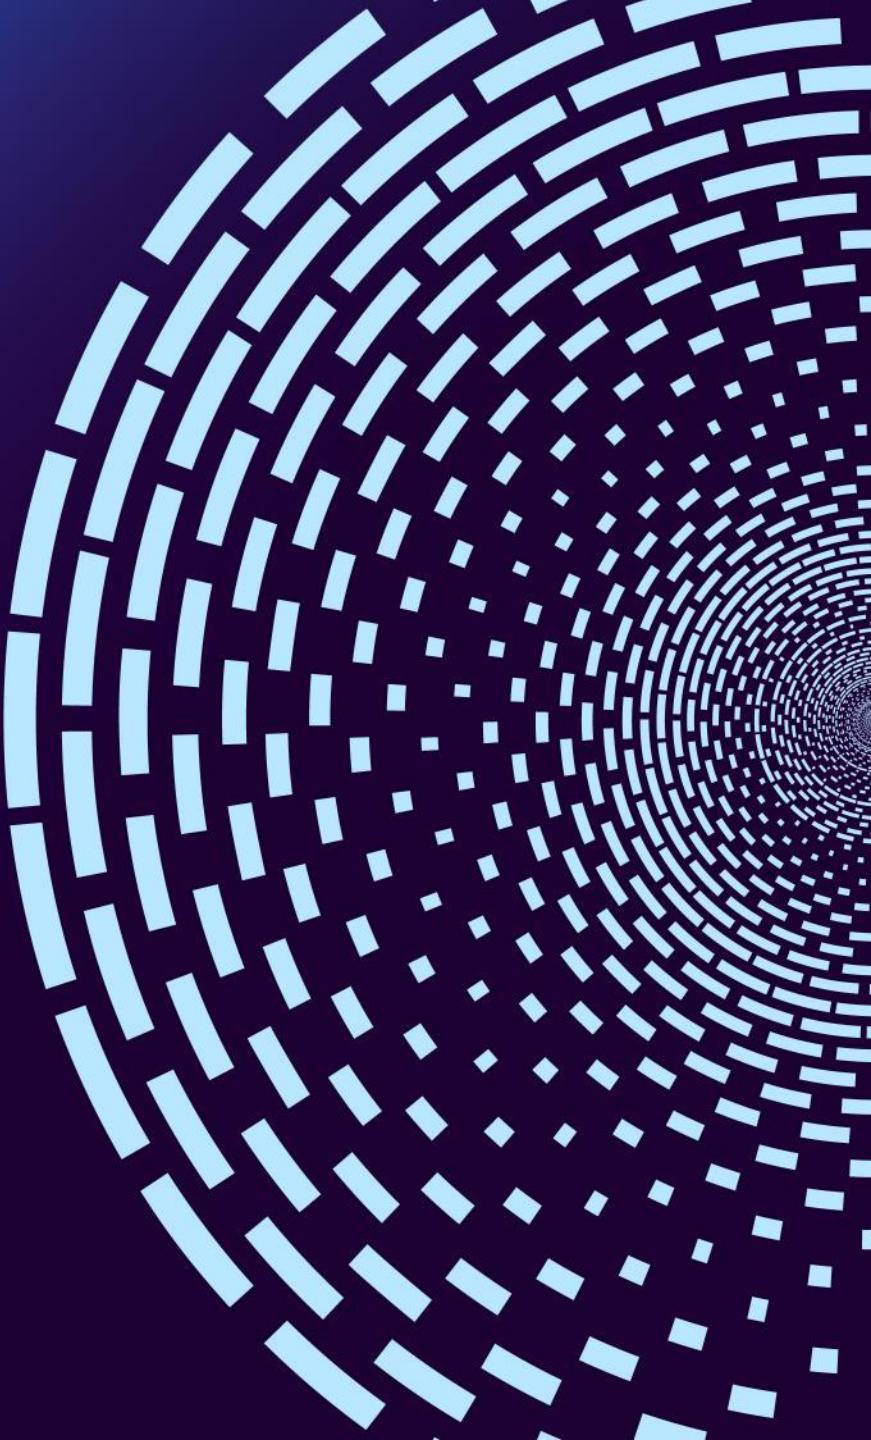




AI Conclave

Online



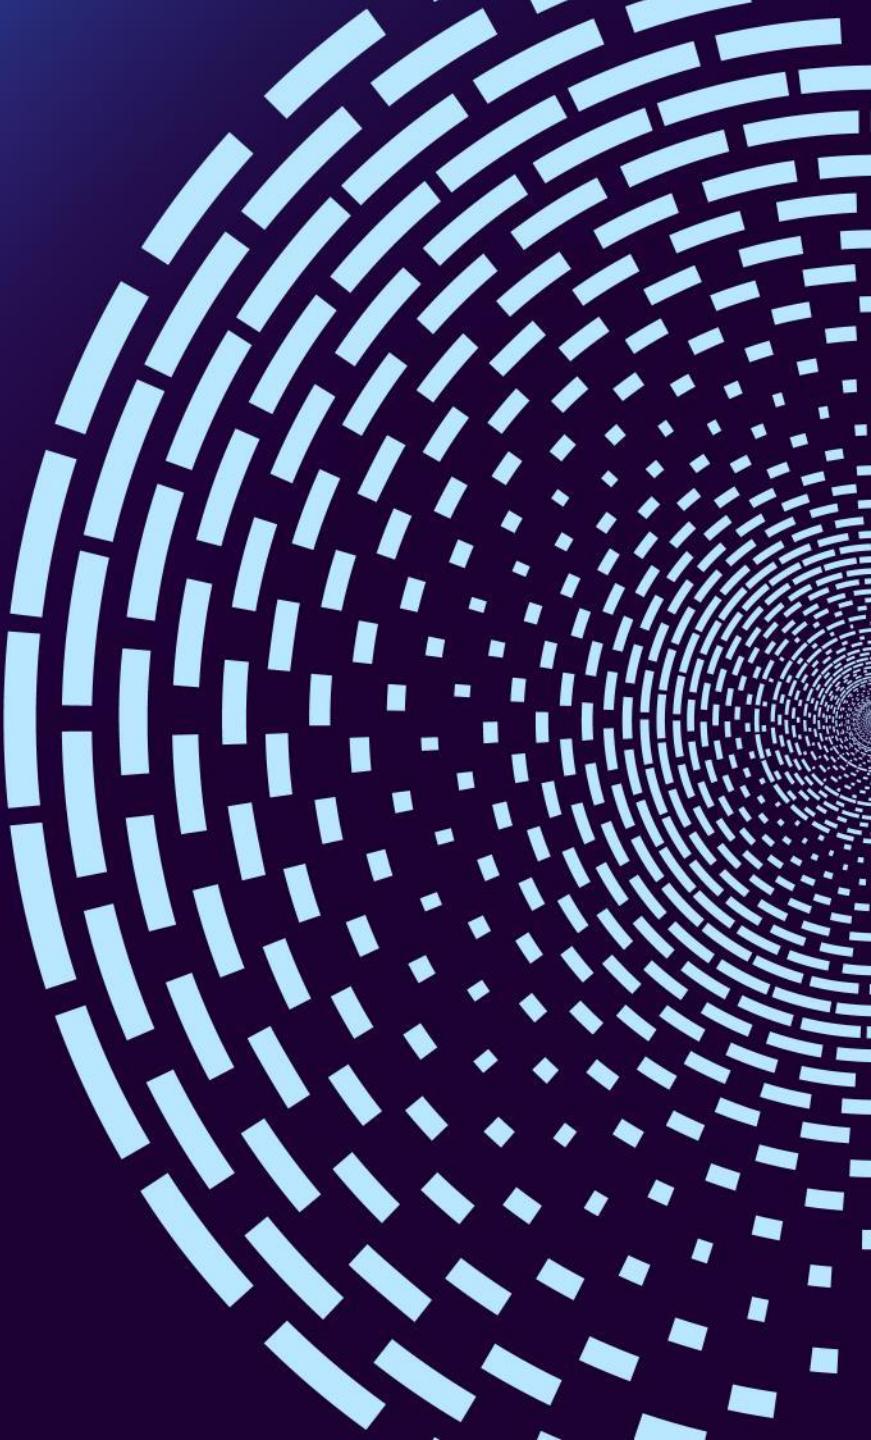


KEY 001

Welcome note

Sundar VG

Director, India Specialists
AWS India





WELCOME



5

TRACKS



30+

BREAKOUT
SESSIONS



3

EXPERIENTIAL
ZONES



20

AWS
CUSTOMERS
FEATURED



10

PARTNERS

Thank you to our Partners

Gold Partners



Silver Partners



India: Largest open market for technology



Unprecedented Demand

3rd largest Startup and SaaS ecosystem

7.5K+ listed companies,
75M + SMBs



Unparalleled Supply

5.2M+ Developers

50% India <25 years; ~900M internet users



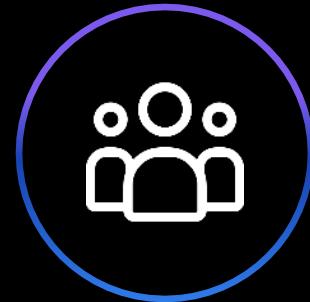
Impact on the world

\$3.9T GDP,
5th largest economy

India stack a template for the world

Indian ISVs, ITESs and GSIs building for the world

The world's largest and most broadly adopted public cloud



Millions

Of active
customers



100k+

Partners in 150+
countries



200+

Fully featured
services



108

Availability Zones



34

Geographic
regions

Our AWS India Mission

**Empowering builders and businesses to build a
better India**

India is building on AWS

83 / Top 100

Indian Enterprises by Market Cap
build on AWS



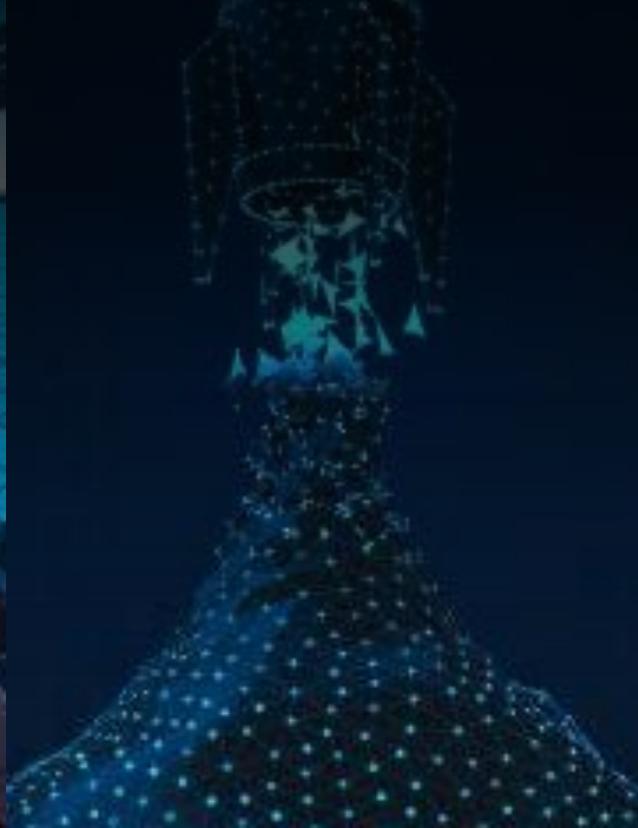
7 / Top 10

Largest Banks in India have adopted
AWS



Over 90% of

India's digital natives
and unicorns



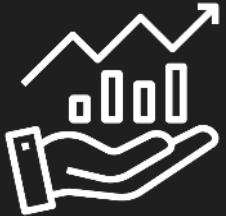
~40K SMBs

Indian SMBs are transforming using
AWS



We have a long-term commitment to India and we're just getting started

\$12.7B



To be invested in the country by 2030

131,700



Jobs supported at external businesses annually by 2030

\$23.3B



Contributed to India's GDP by 2030



Roti



zepto

grofers



Kapda



raymond



NYKAA

JOCKEY®



METRO
SHOES



Makaan





Health



Wealth



Manufacturing, Auto, Retail and Public Sector

Manufacturing



Automotive



Retail



Public Sector



Financial Services

Banking



IndusInd Bank



Insurance



Edelweiss | GENERAL INSURANCE

Capital Mkts



Payments



NBFCs



Startups

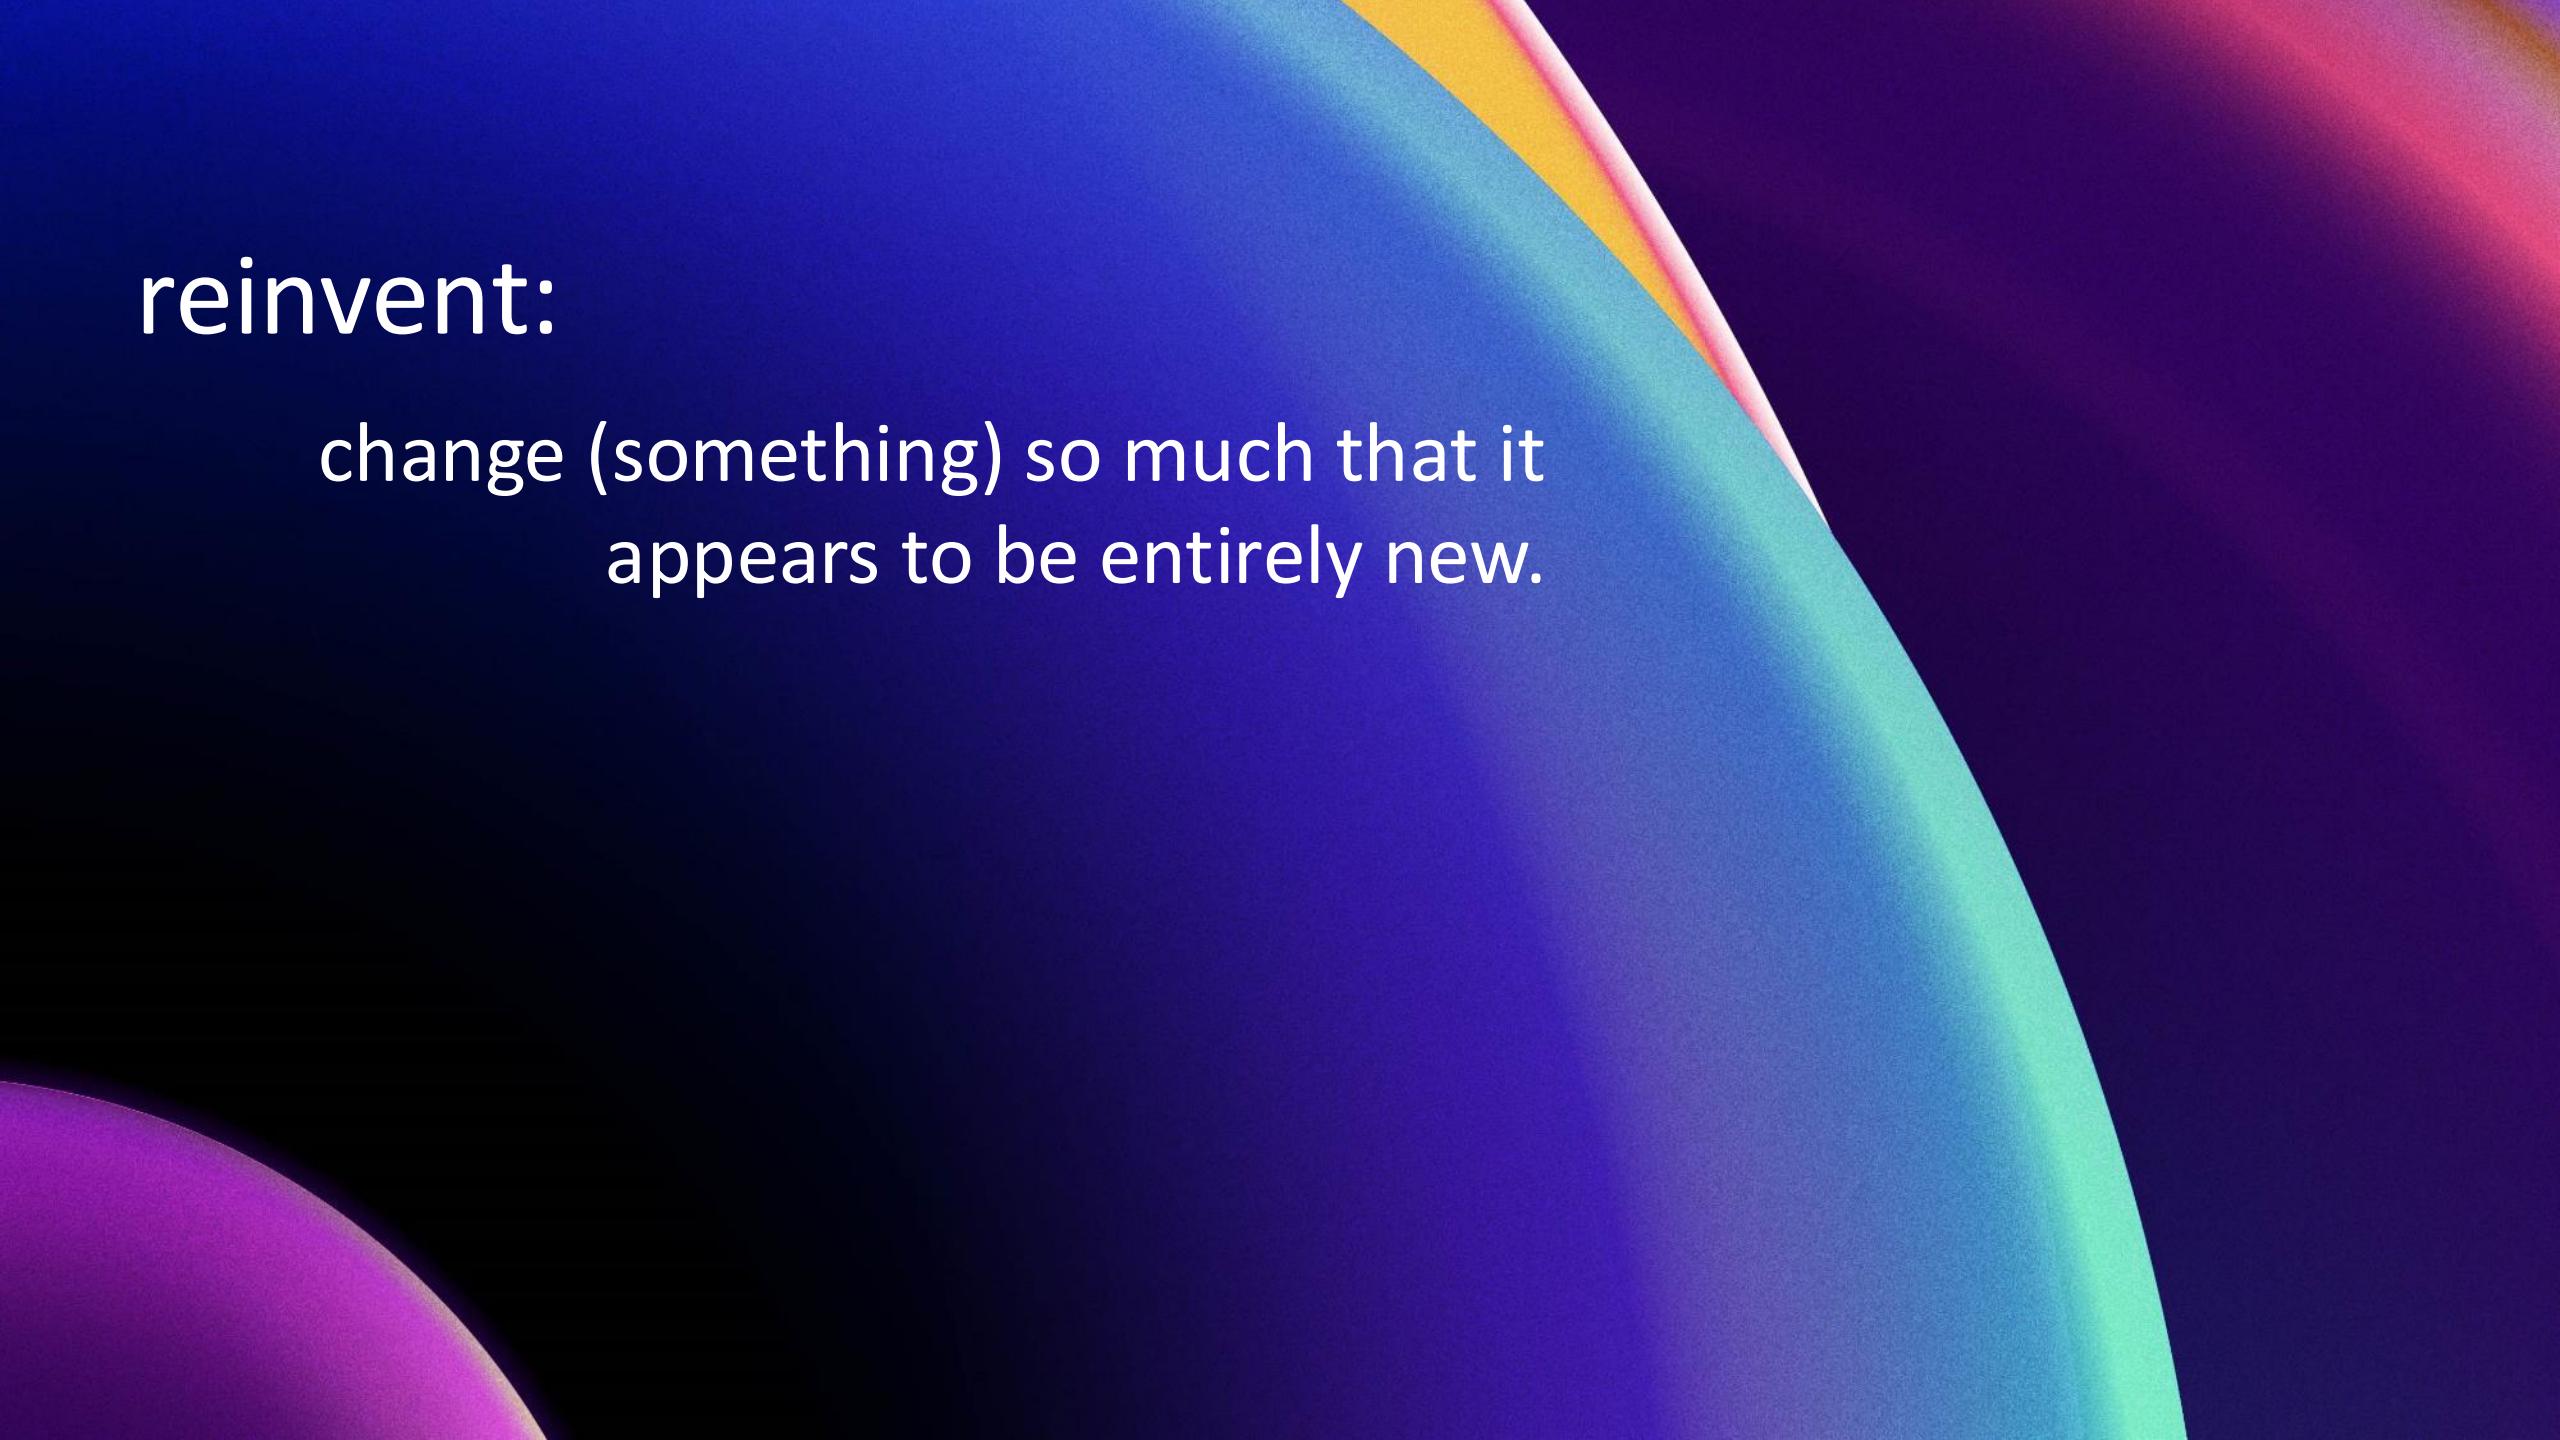


"At Dream11, technology is at the core of everything we do. With over 220 million users, we rely on scalable, innovative systems to deliver seamless and engaging experiences. AWS supports our ability to rapidly prototype and scale new features, helping us stay ahead in a fast-paced market and achieve our vision of making sports better for fans everywhere." – Dream11

Independent Software Vendors



“The range of services offered by AWS is currently unmatched. At least 50 different AWS services, including its entire Media Suite, have been employed by Amagi. Each service has contributed extensively to our business process performance enhancement across the business lifecycle.” – Amagi

The background of the slide features a minimalist design with abstract, overlapping circles in various colors. A large circle on the left is a vibrant purple. To its right and above are smaller circles in shades of blue, green, yellow, and red. The colors transition smoothly into each other where they overlap. The overall effect is clean and modern.

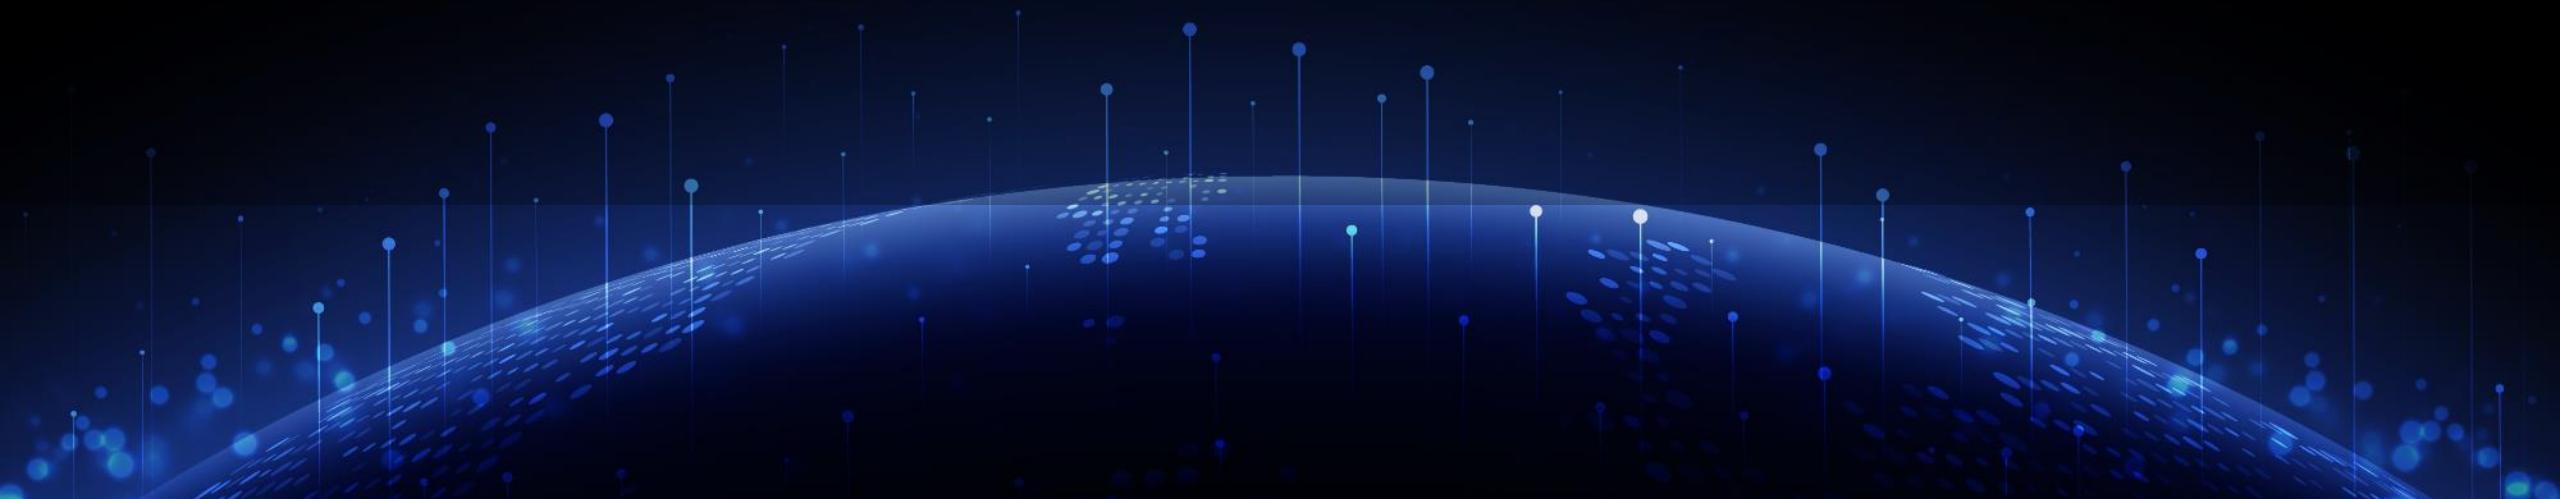
reinvent:

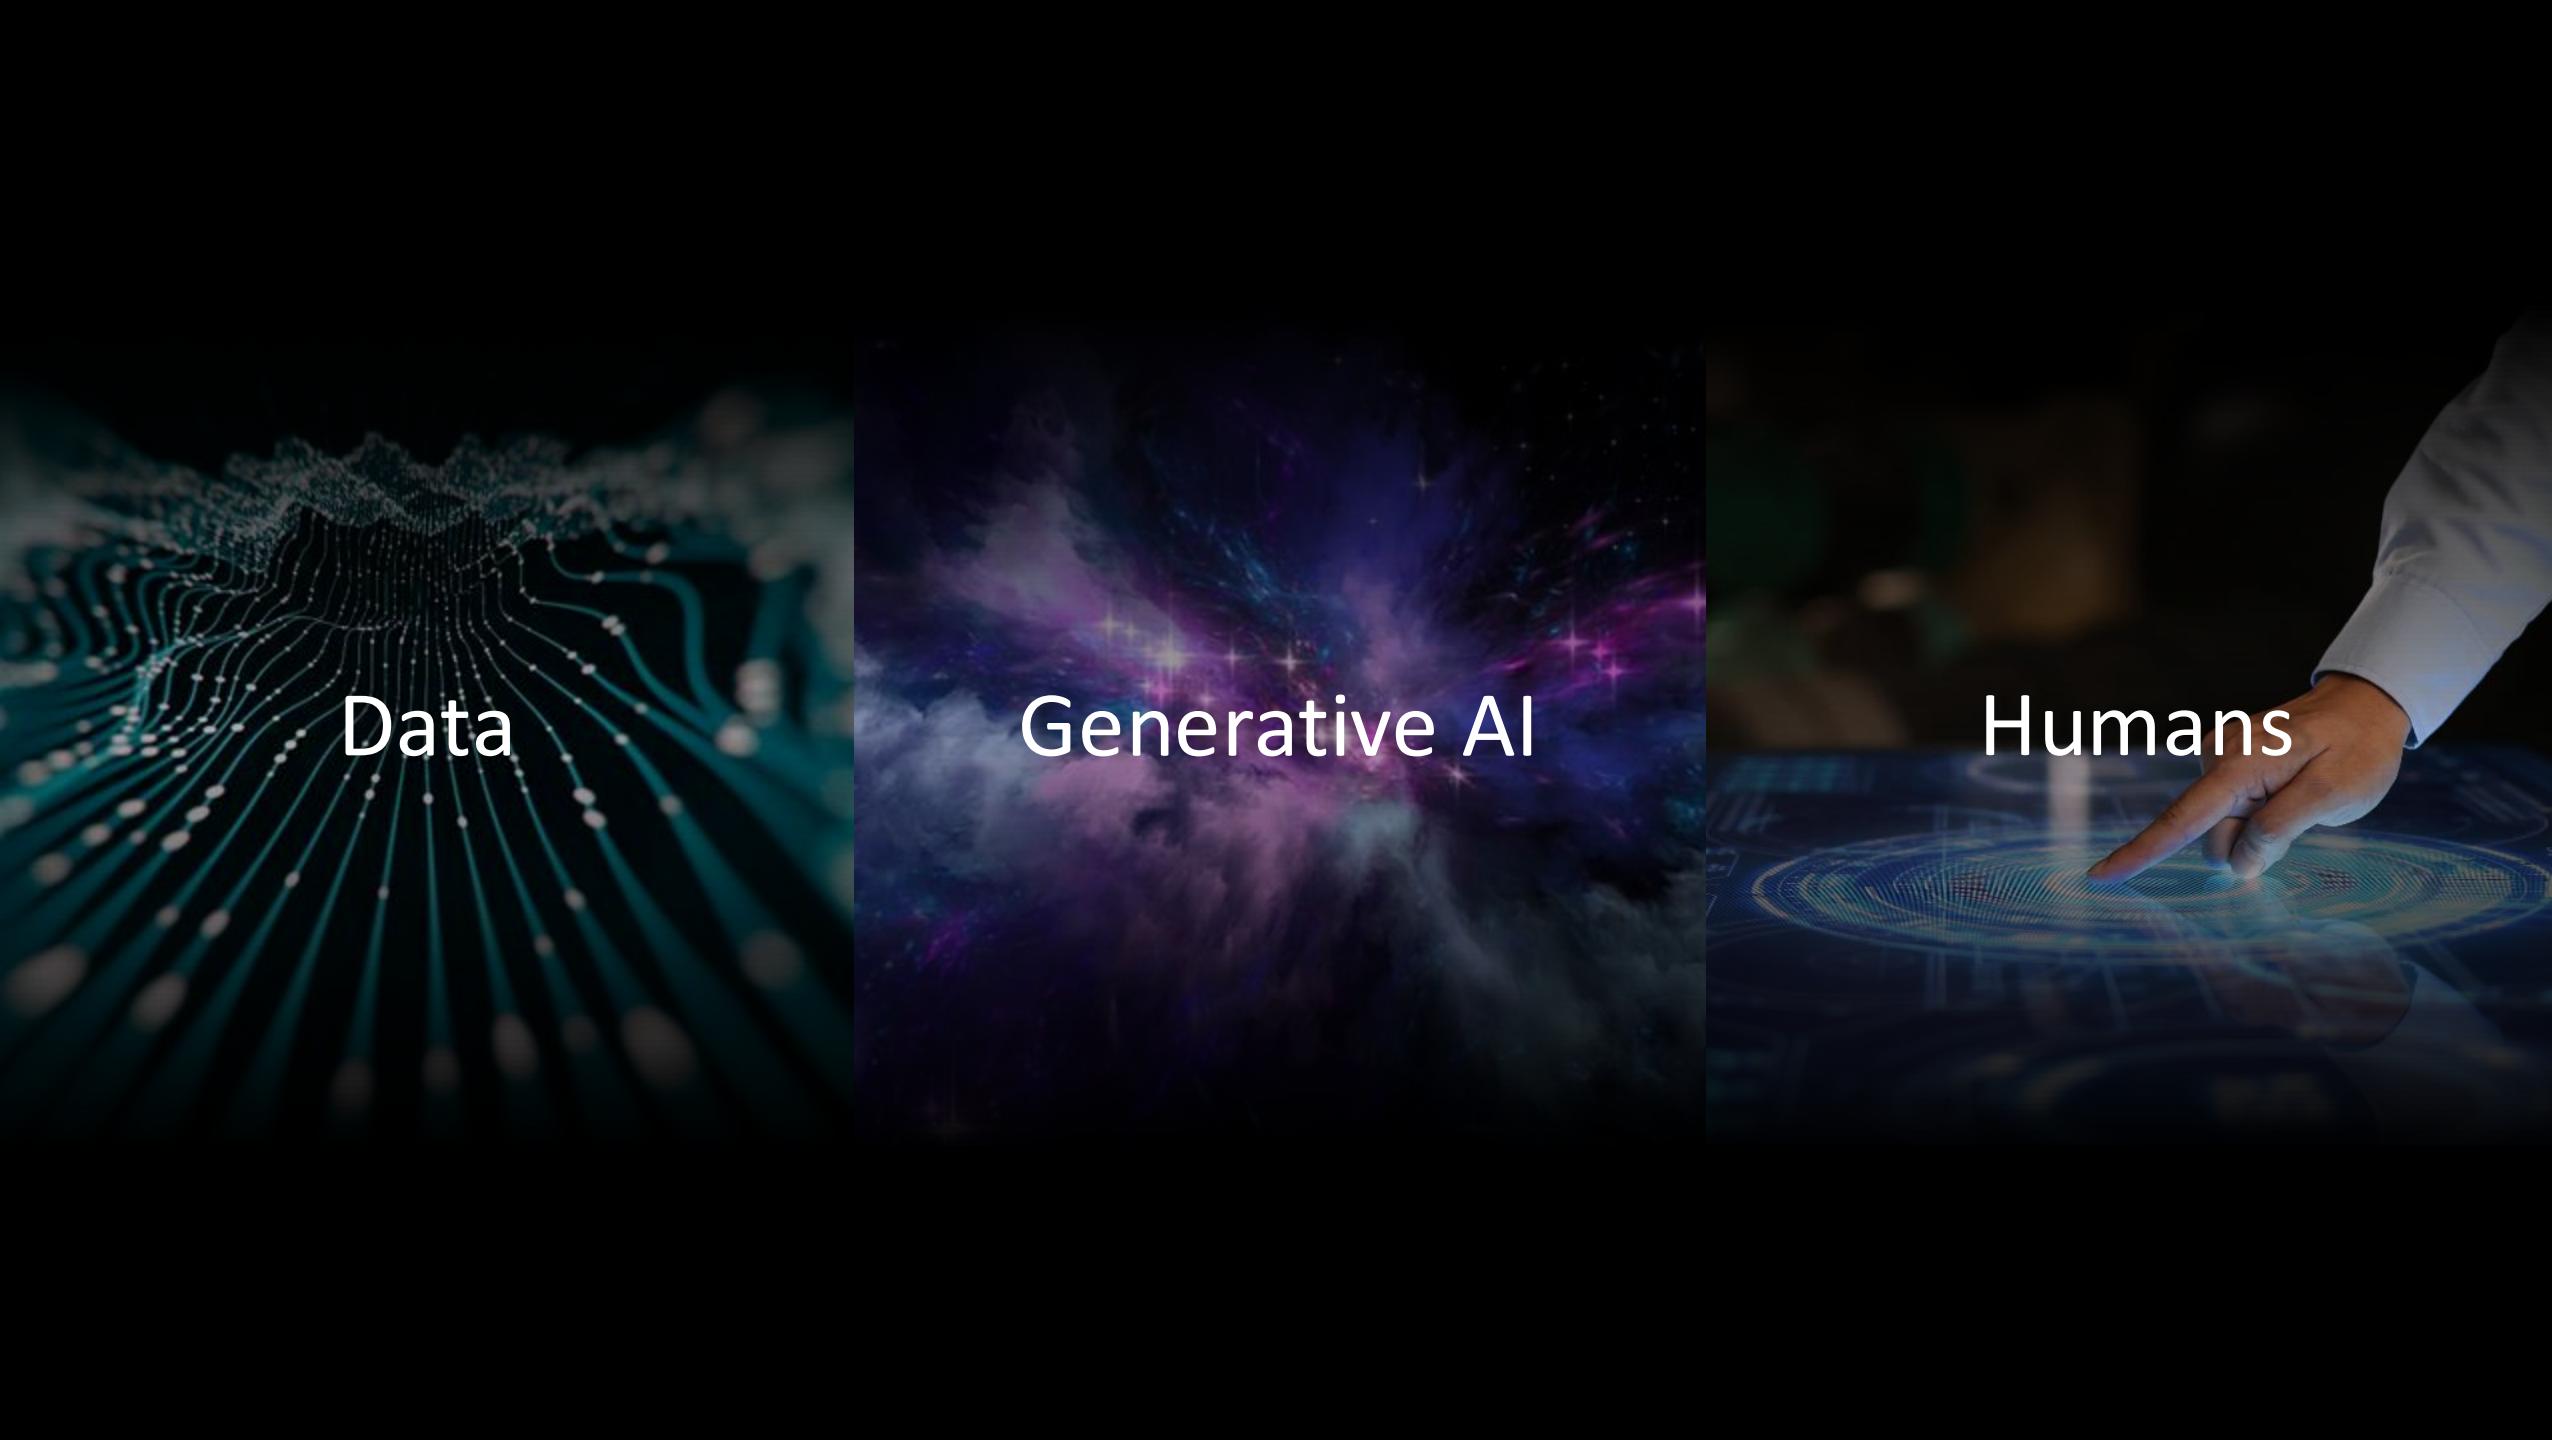
change (something) so much that it appears to be entirely new.



Accelerating new
innovations **everywhere**

Data + Generative AI

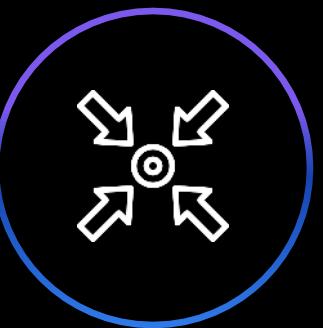




Data

Generative AI

Humans



CONVERGING
PERSONAS



NON-INTEROPERABLE
DATA SILOS



BARRIERS TO AI



TRUST | CONFIDENCE |
COMPLIANCE

UNIFIED

Amazon SageMaker
Unified Studio

Zero ETL

Amazon SageMaker
Lakehouse

Cross Cloud Clean Rooms

Amazon SageMaker
Data and AI Governance

Amazon Bedrock

EFFORTLESS

Amazon S3 Metadata

Serverless Databases

Generative AI for Natural
Language Query

Generative AI powered DMS

Amazon Bedrock
Knowledge Bases

Amazon Bedrock
Data Automation

Amazon Bedrock
Guardrails

Amazon S3 Tables

Amazon Bedrock
Multi Agent Collaboration

Amazon Bedrock
Model Evaluation

[Price] PERFORMANCE

Multi Region Databases

Amazon Nova

Kafka Express Brokers

Purpose built chipset
for LLM

Price Performant Databases

Amazon Q Developer

Amazon SageMaker
HyperPod

Amazon Bedrock Model
Distillation

Next Generation Glue

TYRANNY OF THE

OR

GENIUS OF THE

AND

UNIFIED



Amazon SageMaker Unified Studio



Zero ETL



Amazon SageMaker
Lakehouse



Cross Cloud Clean Rooms



Amazon SageMaker
Data and AI Governance

EFFORTLESS



Amazon S3 Metadata



Generative AI for Natural
Language Query



Generative AI powered DMS



Amazon Bedrock
Knowledge Bases



Amazon Bedrock
Data Automation



Amazon Bedrock
Guardrails



Amazon S3 Tables



Amazon Bedrock
Multi Agent Collaboration



Amazon Bedrock
Model Evaluation



Amazon Bedrock

[Price] PERFORMANCE



Multi Region Databases



Serverless Databases



Amazon Nova



Managed Kafka Express



Purpose built chipset
for LLM



Price Performant Databases



Amazon Q Developer



Amazon SageMaker
HyperPod



Amazon Bedrock
Model Distillation



Next Generation Glue

Featured Speakers



Rick Sears
Director, Data Processing &
Experience, AWS



Luke Anderson
Managing Director, Data & AI,
APJ, AWS



Vasant Rao
Managing Director,
IDC India



Ganesh Kumar Gella
Director, Software Development Engineers,
Bedrock Dev Tooling, AWS



Harsha Bhogle
Indian cricket commentator
and journalist



Ravichandran Ashwin
Indian cricketer

Customer and Partner Speakers



Rajaneesh Kini
Chief Technology Officer,
Cyient



Anuj Variava
Lead, Digital Initiatives,
Raymond Limited



Mridul Gupta
Head of Engineering,
Zolve



Krishna Kumar Chakkirala
Vice President, Data
and AI Practice,
Quinnox



Sachin Dutta
COO,
Canara HSBC Life Insurance



Nishchay Shah
Group CTO & EVP,
Products & AI,
Cactus Communications



Vishnu Vasanth
Co-founder & CEO,
e6data



Aditya Jain
Head of Digital & Analytics,
CEAT Tyres



Vaibhav Dobriyal
Co-founder and
Chief Product Officer,
LUMIQ

Customer and Partner Speakers



Dr Pankaj Dikshit
PhD, Chief Technology Officer,
Government e-Marketplace,
Ministry of Commerce



Amit Prakash
VP, Enterprise Architecture,
Airtel Payments Bank



Poonam Thapar
Engineering Manager,
Zomato



Vaidyanathan S Iyer
Chief Business Officer,
mPokket



Prajakt Deshpande
SVP Engineering,
Kore.ai



Sreedhar Gade
Vice President, Engineering,
Freshworks



Nilav Ghosh
Director of Data Science,
Innovaccer



Mathangi Sri Ramachandran
Chief Data Officer,
Yubi



Jithendra Vepa
CTO,
Observe.ai



Go Create
with AWS



KEY 001

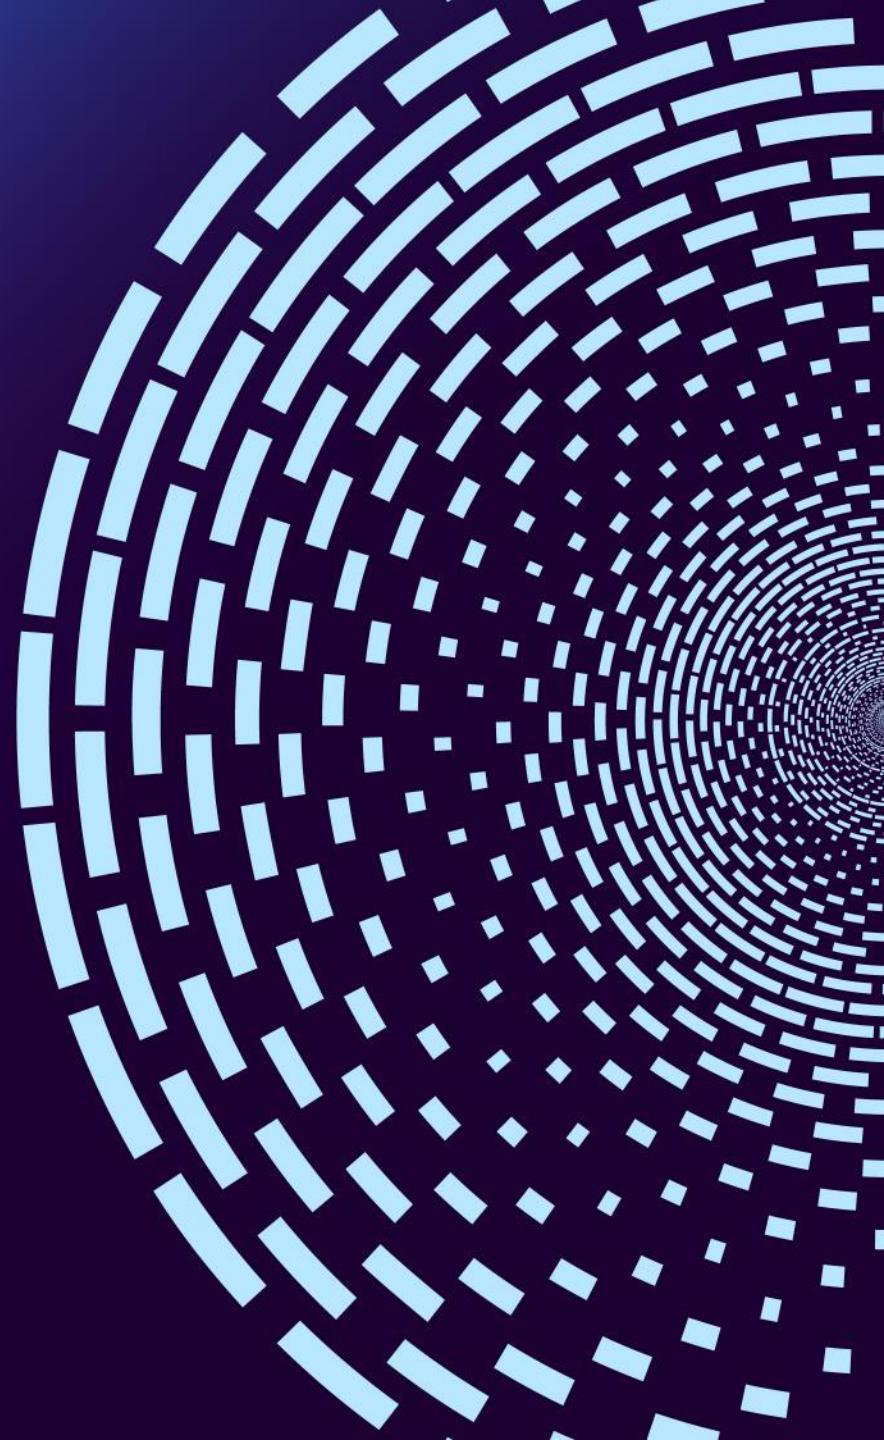
Generative AI & data keynote

Luke Anderson

Managing Director,
Data & AI, APJ
AWS

Rick Sears

Director, Data Processing
& Experience
AWS



366,467 GB

366,527 GB

DATA

ANALYTICS

AI

Effortless

Unified

Performant

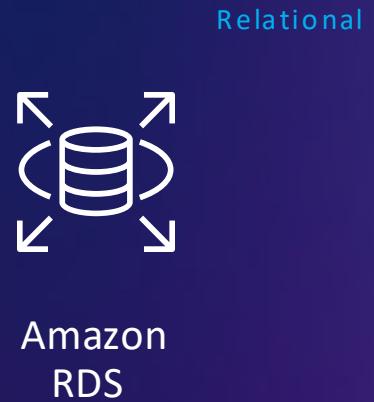
DATA

Effortless



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Comprehensive set of relational & purpose-built databases



Amazon
DynamoDB



Amazon
DocumentDB



Amazon ElastiCache



Amazon
Neptune



Amazon
Timestream



Amazon
QLDB



Amazon
Keyspaces

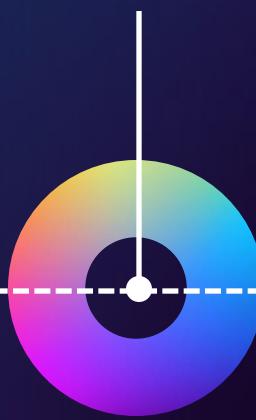


Amazon
MemoryDB

FULLY
MANAGED



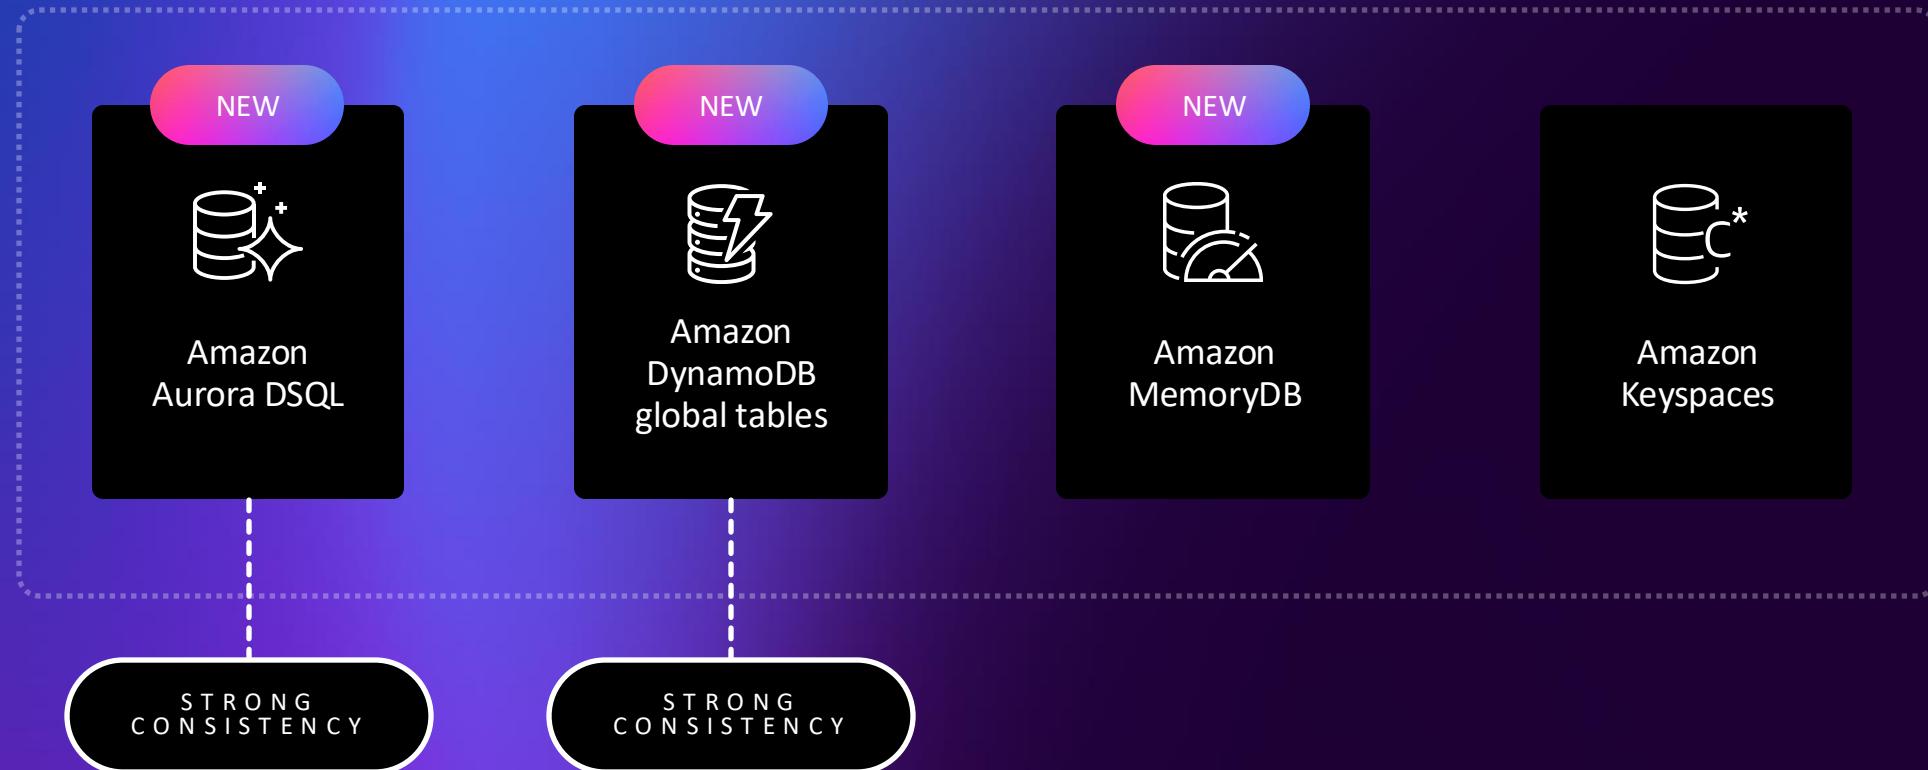
SERVERLESS



?



Active-active multi-Region databases



Aurora PostgreSQL Limitless Database

MANAGED HORIZONTAL SCALE-OUT BEYOND THE LIMITS OF A SINGLE INSTANCE

AUTOMATICALLY SCALE

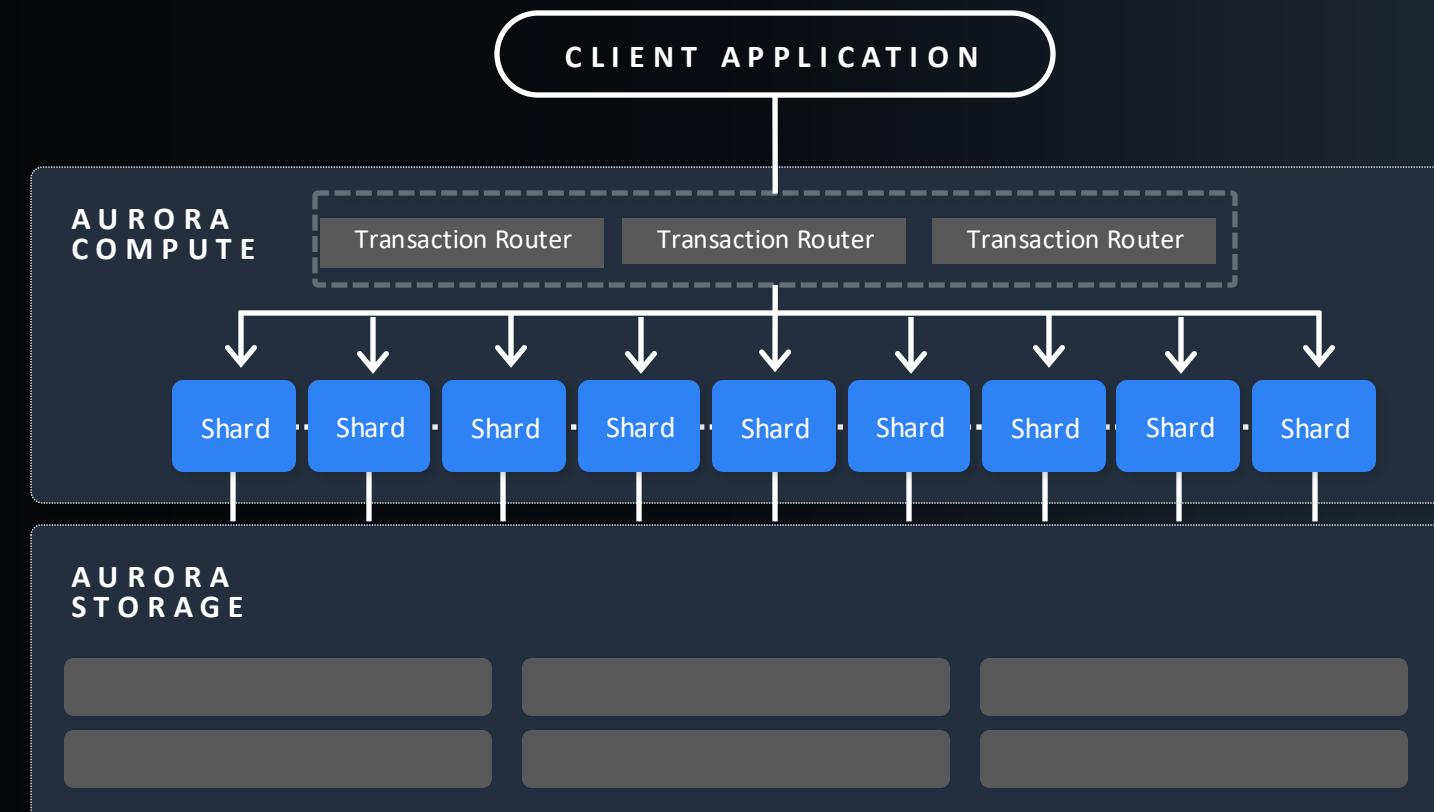
to millions of write transactions per second well beyond current limits of a single PostgreSQL instance

MANAGE PETABYTES

of data in a single database

OPERATIONAL SIMPLICITY

and fast scaling



Oracle Database@AWS

AVAILABLE IN LIMITED PREVIEW

NEW

ORACLE DATABASES

AMAZON EC2 FOR APPLICATIONS



Unify your data across Oracle and AWS to innovate

Easily and quickly migrate Oracle Database workloads to AWS

Simplify management and operations

NEW

AWS DMS Schema Conversion with generative AI

GENERALLY AVAILABLE

Reduce manual time to convert complex database objects

Fully integrated with DMS Schema Conversion

Supports multiple engines to PostgreSQL conversion



Commitment to Open Source

0,000

accepted contributions back to the open source community across PostgreSQL, MySQL, MariaDB, and Valkey



Valkey 8.0

I/O threading architecture

230%

HIGHER
THROUGHPUT

70%

BETTER
LATENCY

20%

LOWER
MEMORY OVERHEAD



New pricing reductions for customers

ELASTICACHE
SERVERLESS
FOR VALKEY

33% ↓

DYNAMO DB
GLOBAL TABLES
ON-DEMAND

50% ↓

KEYSPACES
ON-DEMAND

up to 65% ↓

NODE-BASED
ELASTICACHE
FOR VALKEY

20% ↓

DYNAMO DB
GLOBAL TABLES
WRITES

up to 67% ↓

KEYSPACES
PROVISIONED

up to 20% ↓



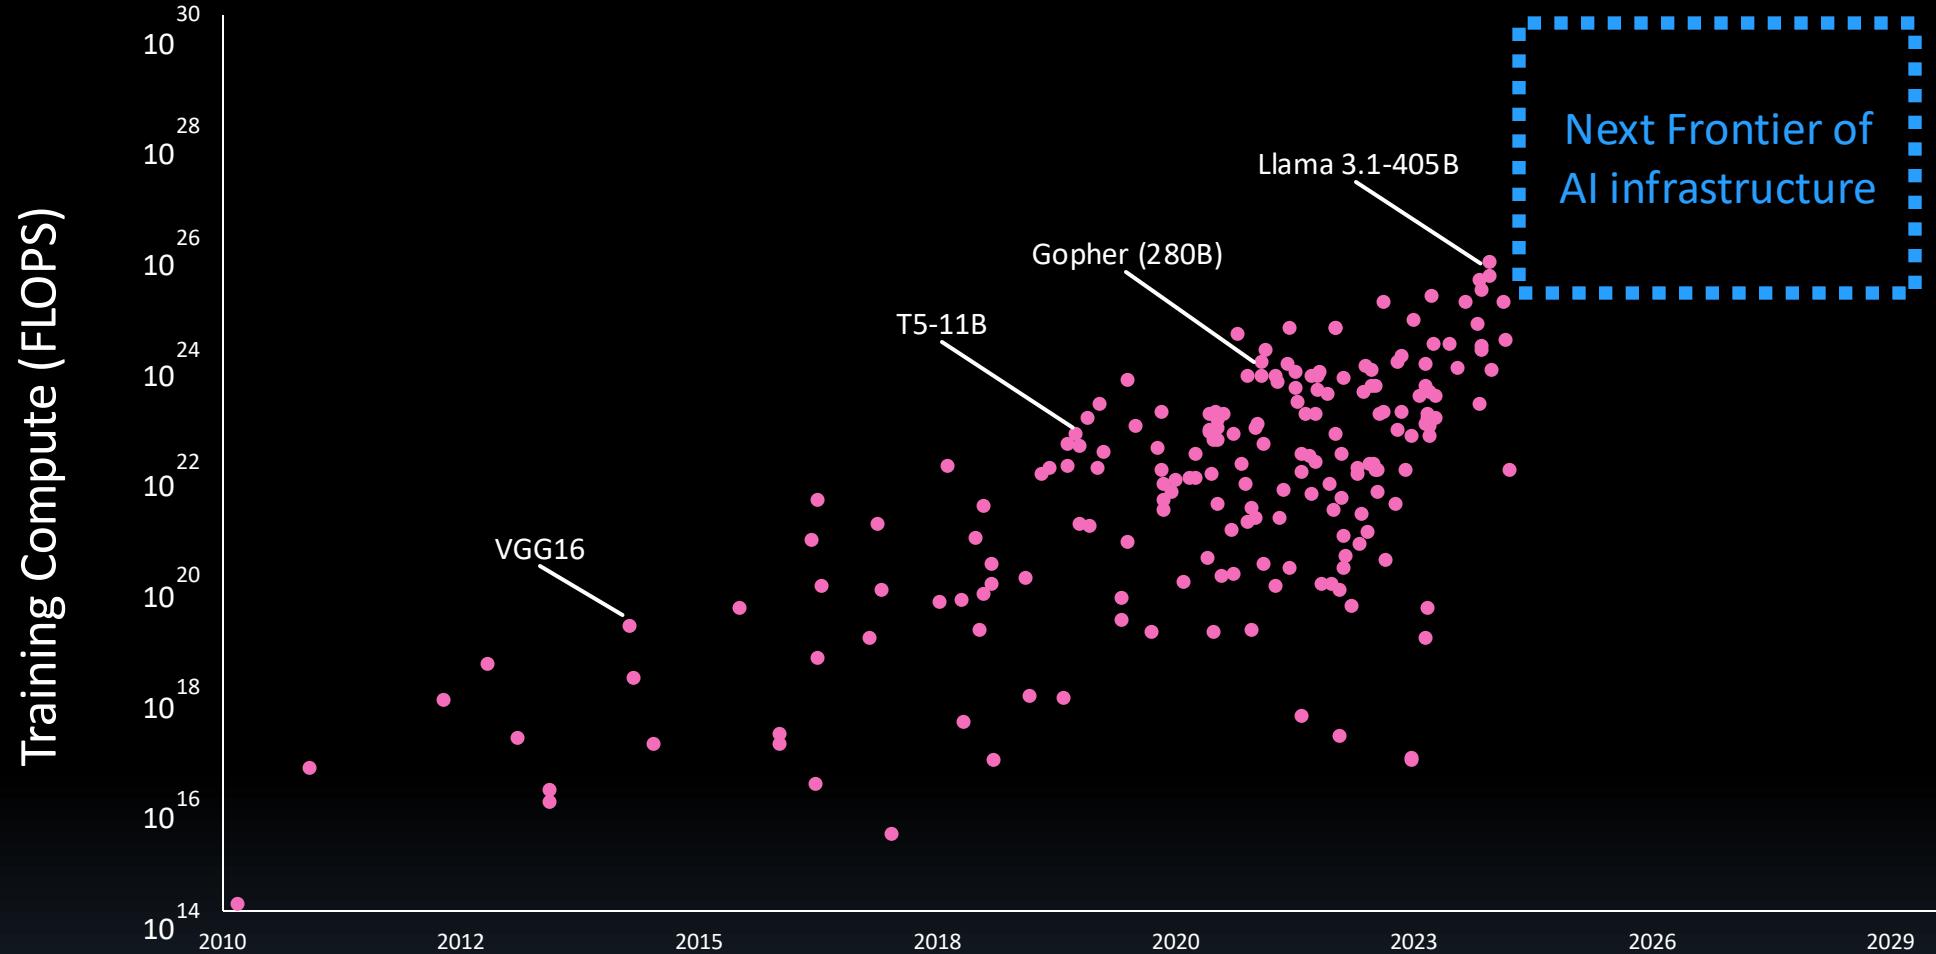
DATA

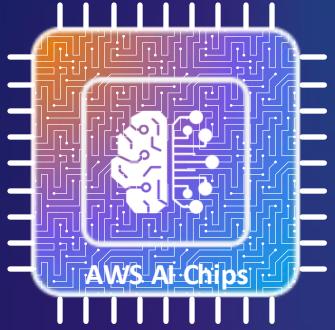
Effortless



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Scaling compute requires next-generation AI infrastructure





AWS AI Chips for Generative AI

AWS Inferentia

Lowest Cost

AWS Inferentia2

Best price performance

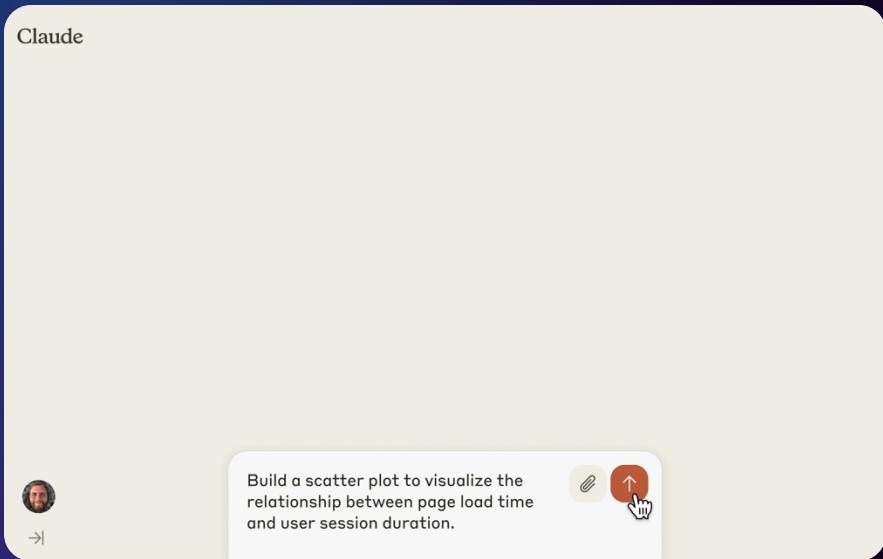
AWS Trainium

Lowest cost to train models up to 70b

AWS Trainium2

Most powerful compute for frontier models

Anthropic models run 60% faster with Trainium2



At Anthropic, our mission is to ensure transformative AI helps people and society flourish. Speed is essential for our enterprise users, that's why we announced the availability of Claude Haiku 3.5, which matches the performance of Claude 3 Opus, at industry leading latency," said Dario Amodei. "We are excited to introduce Claude 3.5 Haiku running on Trn2 for latency-optimized inference, now available via Amazon Bedrock. It provides the fastest response time for our customers, 60% faster than any other instance we use. We are excited to continue migrating more Claude models to Trainium2, so our customers enjoy the performance boost and productivity increase."

Efficiently managing compute resources
across multiple tasks is complex





© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

NEW

Amazon SageMaker AI

Amazon SageMaker
HyperPod flexible training
plans

Save weeks of training time and
help meet timelines and budget

Amazon SageMaker
HyperPod task governance

Maximize accelerator utilization for model
training, fine-tuning, and inference

AI apps from AWS
partners now available in
Amazon SageMaker

Find, deploy, and use AI apps from AWS
partners within Amazon SageMaker



Inference needs are rapidly evolving

Amazon Bedrock

The easiest and fastest way to build and scale generative AI applications

Choose the
best model

Broadest choice of industry leading models

Custom model import

Model evaluation

Optimize for cost,
latency, and accuracy

Latency-optimized inference

Model distillation

Customize
with your data

Fine-tuning

Knowledge Bases

Apply safety and
responsible AI checks

Guardrails

Automated reasoning checks

Build and
orchestrate agents

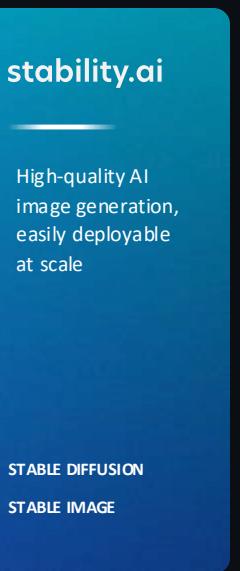
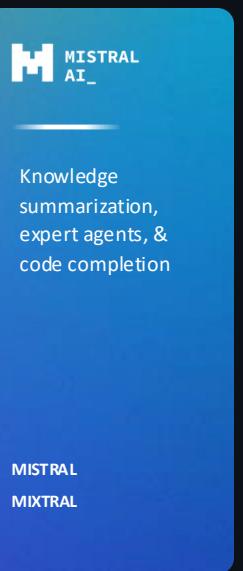
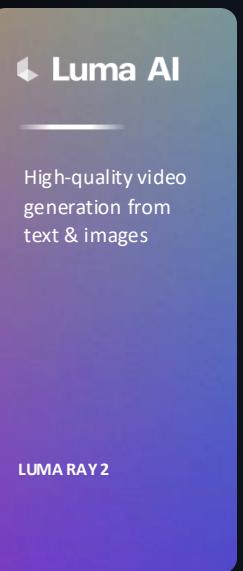
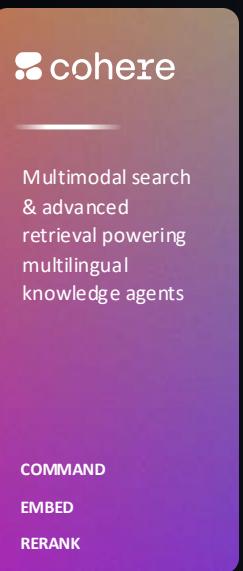
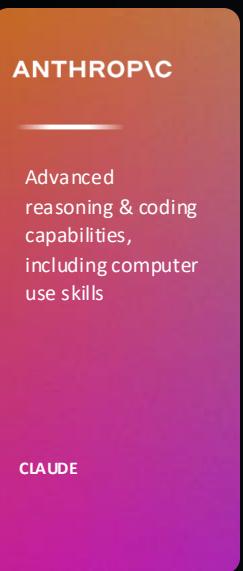
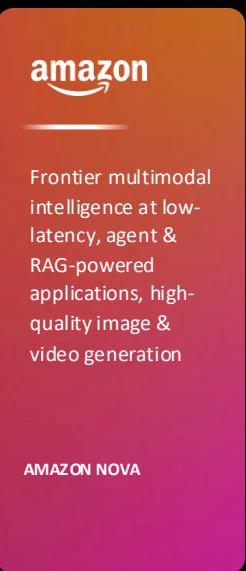
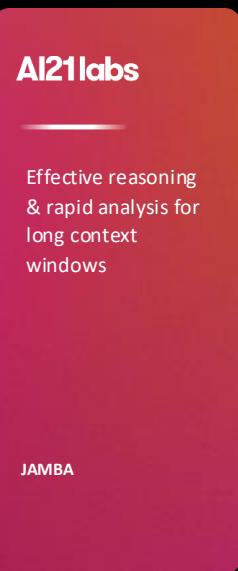
Agents

Multi-agent collaboration

AWS SECURITY, PRIVACY, AND RELIABILITY BUILT-IN



Amazon Bedrock



Amazon Nova

State-of-the art foundation models that deliver frontier intelligence and industry leading price performance



Micro

Text only model that delivers the lowest latency responses at very low cost

Lite

Very low-cost multimodal model that is lightning-fast for processing image, video, and text inputs



Canvas

State-of-the-art image generation model

Pro

Highly capable multimodal model with the best combination of accuracy, speed, and cost for a wide range of tasks

Premier

Most capable multimodal model for complex reasoning tasks and for use as the best teacher for distilling custom models

Reel

State-of-the-art video generation model



Optimize for cost, latency, and accuracy

Latency
optimized inference

Model
distillation

Most enterprise data is unstructured,
multimodal content that isn't ready for generative AI

NEW

Amazon Bedrock Data Automation

Transform unstructured multimodal data for generative AI applications and analytics

AVAILABLE IN PREVIEW

Extract, transform, and generate structured data from multimodal content

Generate customized outputs based on your requirements and business rules

Streamline application workflows with a fully managed, single API experience



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Amazon Bedrock Knowledge Bases

Transform unstructured multimodal data for
generative AI applications and analytics

- Structured Data retrieval
- Multi-modal data processing
- GraphRAG
- RAG evaluation
- Auto generated query filters
- Re-rank API
- Real-time sync for custom data sources
- Streaming responses



Evolving generative AI use cases require evolving safeguards

Amazon Bedrock Guardrails

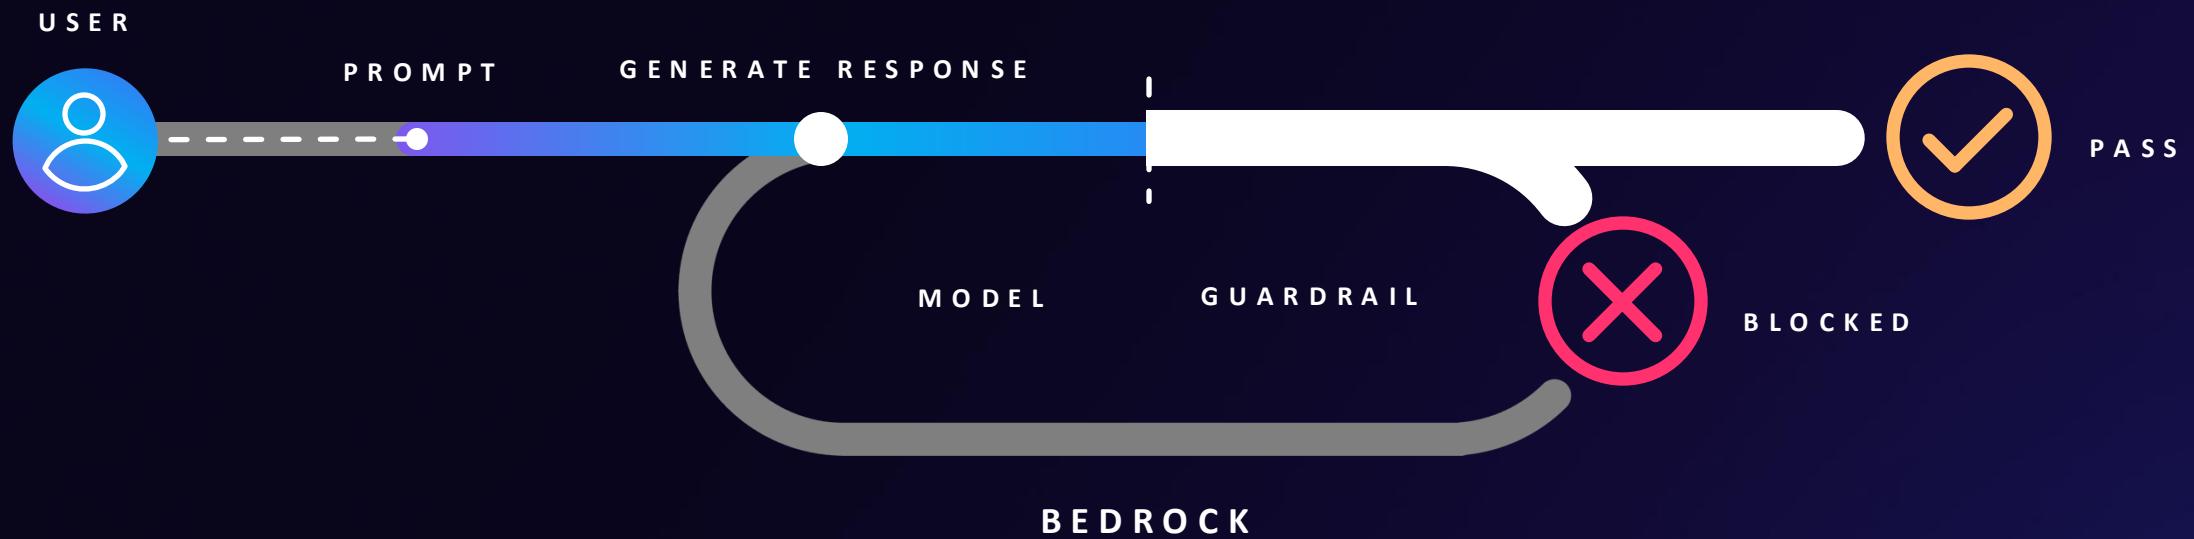
Safeguard your generative AI applications with
your responsible AI policies

Configurable safeguards

Contextual grounding
checks

Automated reasoning
checks





NEW

Amazon Bedrock Guardrails Multimodal toxicity detection

Configurable safeguards for image content

AVAILABLE IN PREVIEW

Enhance security of multimodal generative AI applications

Enable consistent policy control

Available for all foundation models in Amazon Bedrock with image support



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Amazon Bedrock Agents

Execute multi-step tasks across company systems and data sources

Enables generative AI applications to take action in just a few clicks

Breaks down and orchestrates tasks and executes API calls on your behalf

Securely accesses and retrieves company data

Multi-agent collaboration to handle complex, multi-step tasks



Amazon Q accelerates
productivity across organizations



Amazon Q Developer

Most capable generative AI-powered assistant for software development

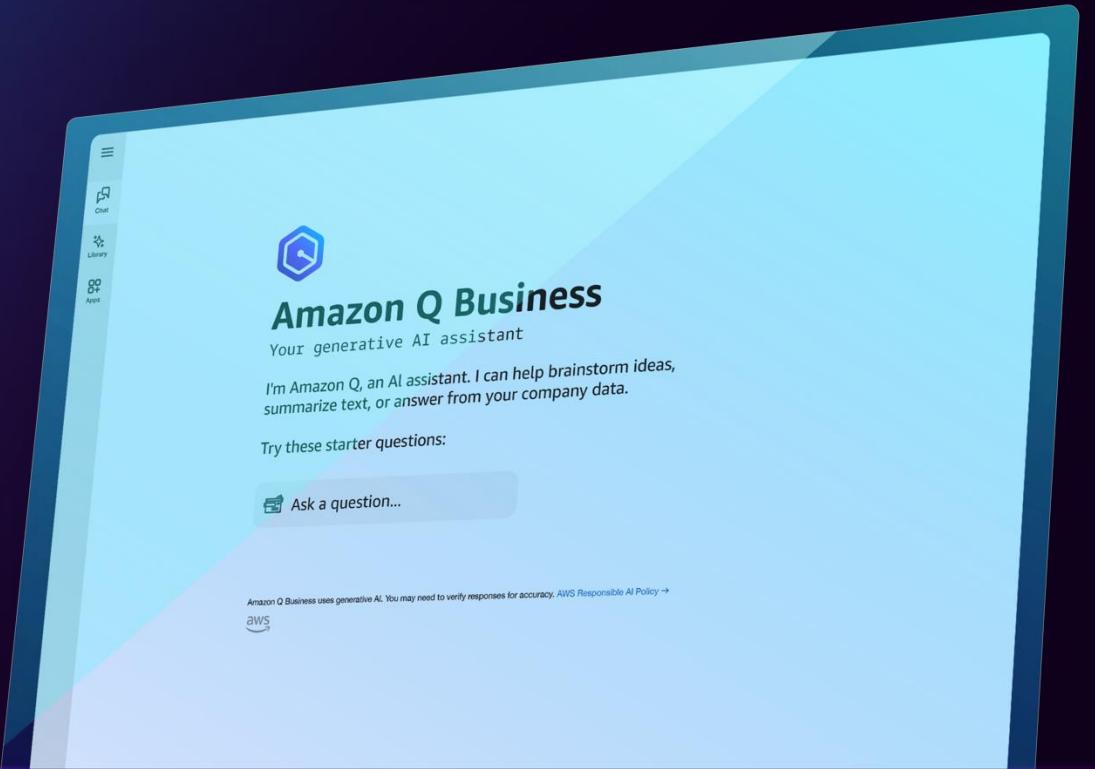
The screenshot shows the AWS Lambda function editor with the following details:

- Function Name:** AWSLambda Chat
- Runtime:** Node.js 14.x
- Handler:** index.handler
- Code Editor:** The code is written in Node.js. It includes logic for handling various API requests related to VM management (create, update, delete, pause, resume, snapshot, restore) and monitoring metrics.
- Annotations:** Numerous annotations are present throughout the code:
 - Comments:** Explain the purpose of specific sections of code.
 - Checkmarks:** Indicate successful validation or completion of certain steps.
 - Links:** Provide external references to documentation or other parts of the system.
- Suggested Fixes:** A sidebar lists suggestions for improving the code:
 - How is the error handled in the handleRequest func.
 - What are the specific performance metrics tracked in the code?
 - How is the asynchronous communication between components?
- Footer:** Includes a note about AWS Lambda's memory limit and a link to the AWS Responsible AI Policy.



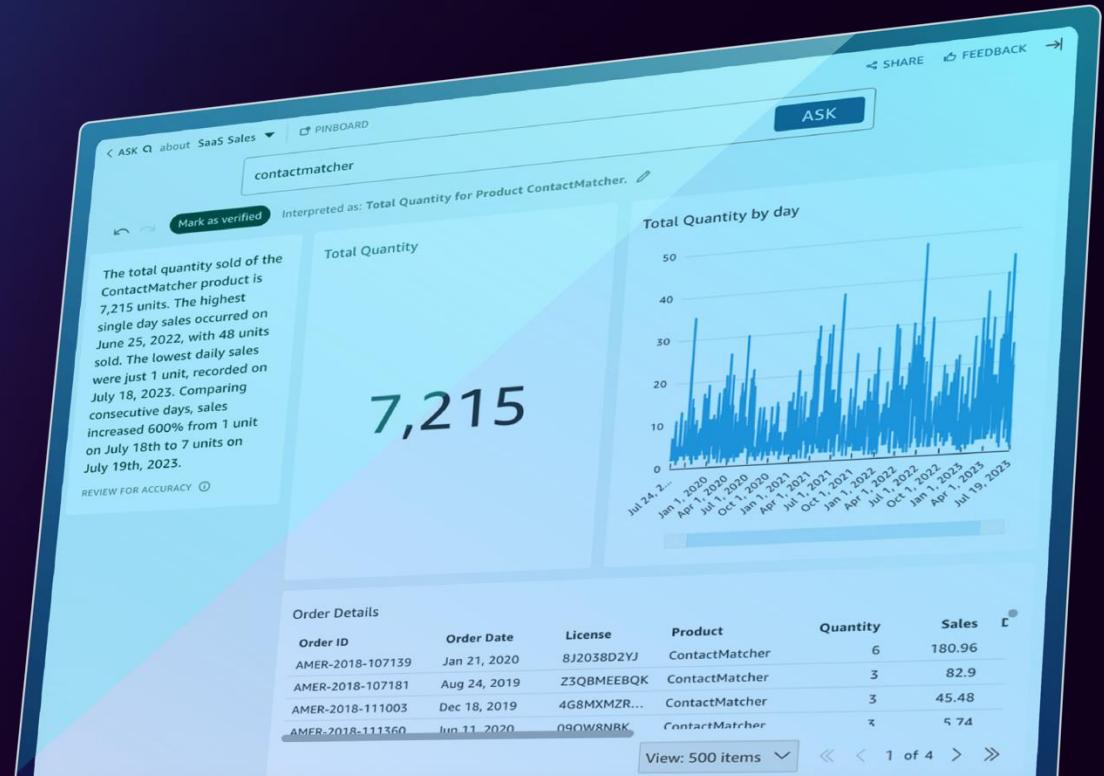
Amazon Q Business

The most capable generative AI-powered assistant for leveraging companies' internal data





Amazon Q in QuickSight



ANALYTICS

AI

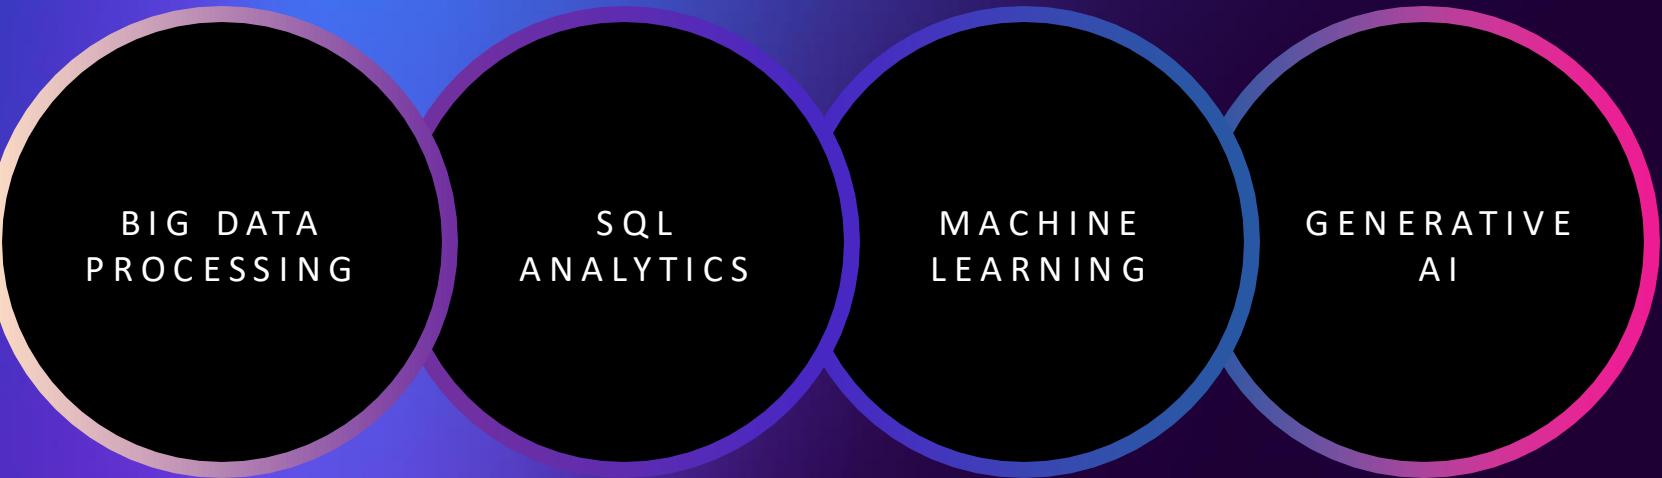
DATA

Unified

Performant

Effortless







© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.
© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Bring together widely adopted and comprehensive ML, AI, and analytics capabilities from AWS

Experience



Act



Govern



Store



NEW

The next generation of
 Amazon SageMaker

The center for all your data, analytics, and AI



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

NEW



The next generation of Amazon SageMaker

Collaborate and build faster

Develop and scale AI use cases

Reduce data silos

Meet your enterprise security needs





Unified Studio Preview

Streaming | Data Processing | SQL Analytics | Model Development | Gen AI App Development | Business Intelligence | Search Analytics

comingsoon

comingsoon

comingsoon

Data & AI Governance

Lakehouse



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Amazon SageMaker

Unified Studio Preview

Streaming | Data Processing | SQL Analytics | Model Development | Gen AI App Development | Business Intelligence | Search Analytics

comingsoon

comingsoon

comingsoon

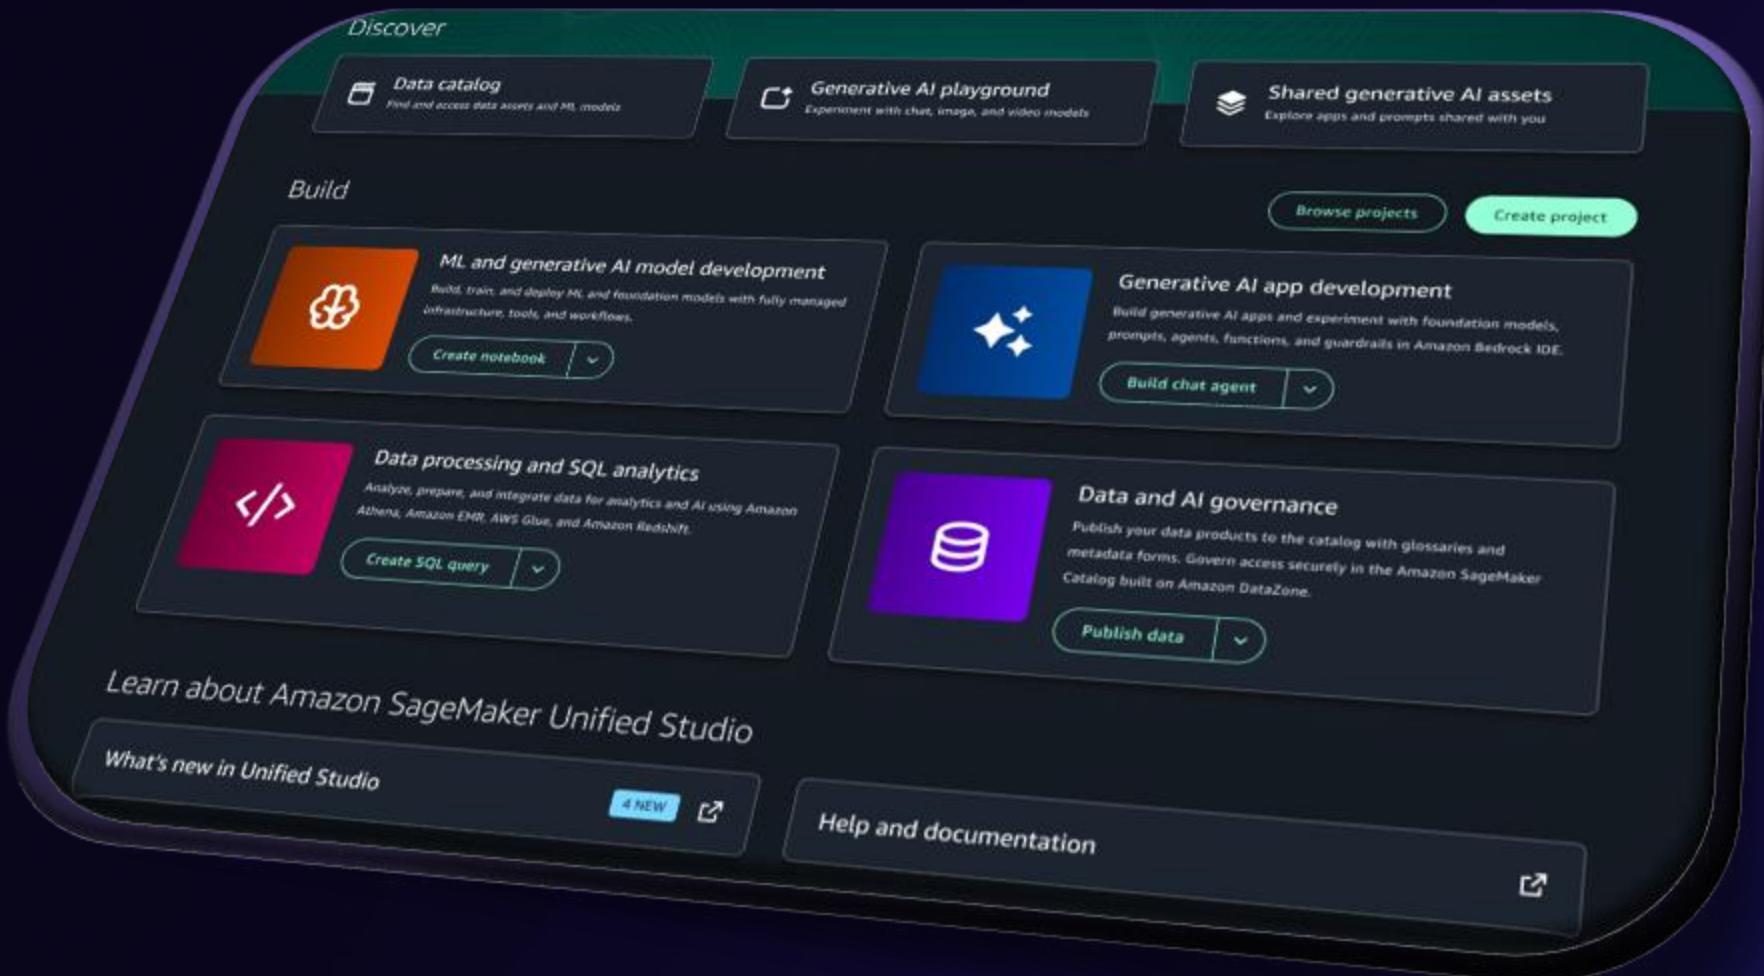
Data & AI Governance

Lakehouse

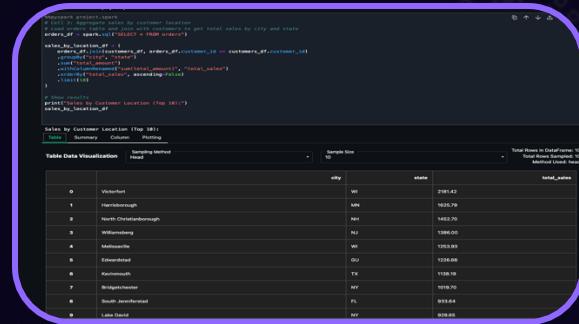
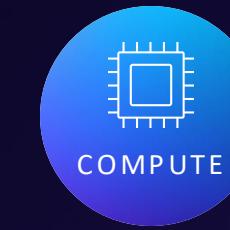


© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

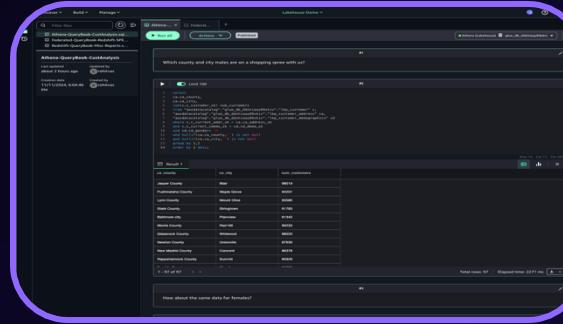
A single data and AI development environment



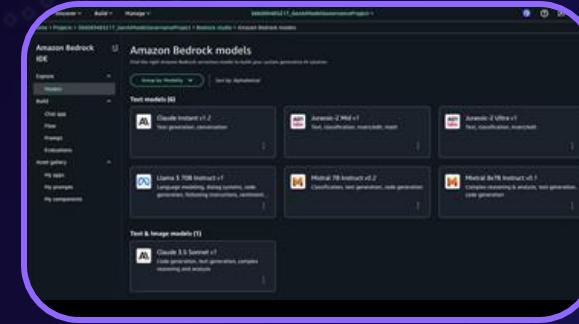
Seamless access to all data, all resources, and all tools



JUPYTER NOTEBOOKS



SQL EDITOR



BEDROCK IDE
(preview)

Accelerated by Amazon Q Developer ...

The screenshot shows the AWS Data Explorer interface with a purple oval highlighting the main workspace. On the left, the Data explorer sidebar lists 'Lakehouse' and 'Redshift' databases. The main workspace displays two queries:

- Query #1:** Analyzes customer demographics. It selects gender, average age, and customer count from the customer_demographics table.
- Query #2:** Analyzes purchase patterns. It joins purchase_history and customer_demographics tables to calculate average purchase amount, unique customers, and total orders per income bracket.

To the right, an 'Amazon Q' panel is open, containing a generated SQL query to find top users by number of orders:

```
1 SELECT
2   cd.customer_id,
3   COUNT(ph.order_id) AS total_
4   FROM
5   sales_churn_ctb6edif88vand.|_
6   JOIN sales_churn_ctb6edif88v|_
7   GROUP BY
8   cd.customer_id
9   ORDER BY
10  total_orders DESC
11 LIMIT
12  100;
```

The Amazon Q panel also includes a text input field for generating more suggestions and a note about using generative AI.





Unified Studio

Streaming | Data Processing | SQL Analytics | Machine Learning | Gen AI Development | Business Intelligence | Search

Data & AI Governance

GENERALLY AVAILABLE

Lakehouse



Amazon SageMaker Data and AI Governance

DISCOVER, GOVERN, AND COLLABORATE ON DATA AND AI SECURELY

SageMaker Catalog



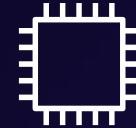
Data



Models



GenAI



Compute

Business Metadata

Sales_2024

Technical name: [Technical Name] • Domain unit: [Domain unit] • Asset type: Glue Table

Description of all the information and artifacts required for creation of Sales dataset for

BUSINESS METADATA SCHEMA DATA QUALITY LINEAGE MY SUBSCRIPTION

README

• Vivamus et nullis id ex, eget tristique lectus.

• Nulla et cursus lacus. Duis tempus posuere justo sed ultricies. Nunc euismod est quis est tempus.

• Quia tacita lenore viseluptat. Nam euismod nulla eu sem varus consempat. Fusca et tortor nec sit interdum laurum et non hinc.

• Pellentesque varius est a collectudine finibus.

SHOW LESS →

GLOSSARY TERMS

[Term] [Term] [Term] [Term]

METADATA FORMS ⓘ

[System defined metadata form]

[sample content: Create a new data source to import technical metadata of assets from AWS Glue data catalog or Amazon Redshift warehouse into the Amazon DataZone inventory.]

Key	Value
Alias	Cell Value

ACTIONS

- Query with Almond
- Open in JupyterLab-notebook
- Edit
- Unpublished
- Add to existing data product
- Create a new data product with this asset
- Delete

V1.2.3

Last modified by: [data source name] or [username]

Last metadata updated: Yesterday

Created by: [data source name] or [username]

Created: Jun 12, 2012

Owner: Performance Central

Subscription approval: Required

What does MaxDome offers and why I should use MaxDome?

How can my data science team ensure data security when sharing curated datasets with external partners?

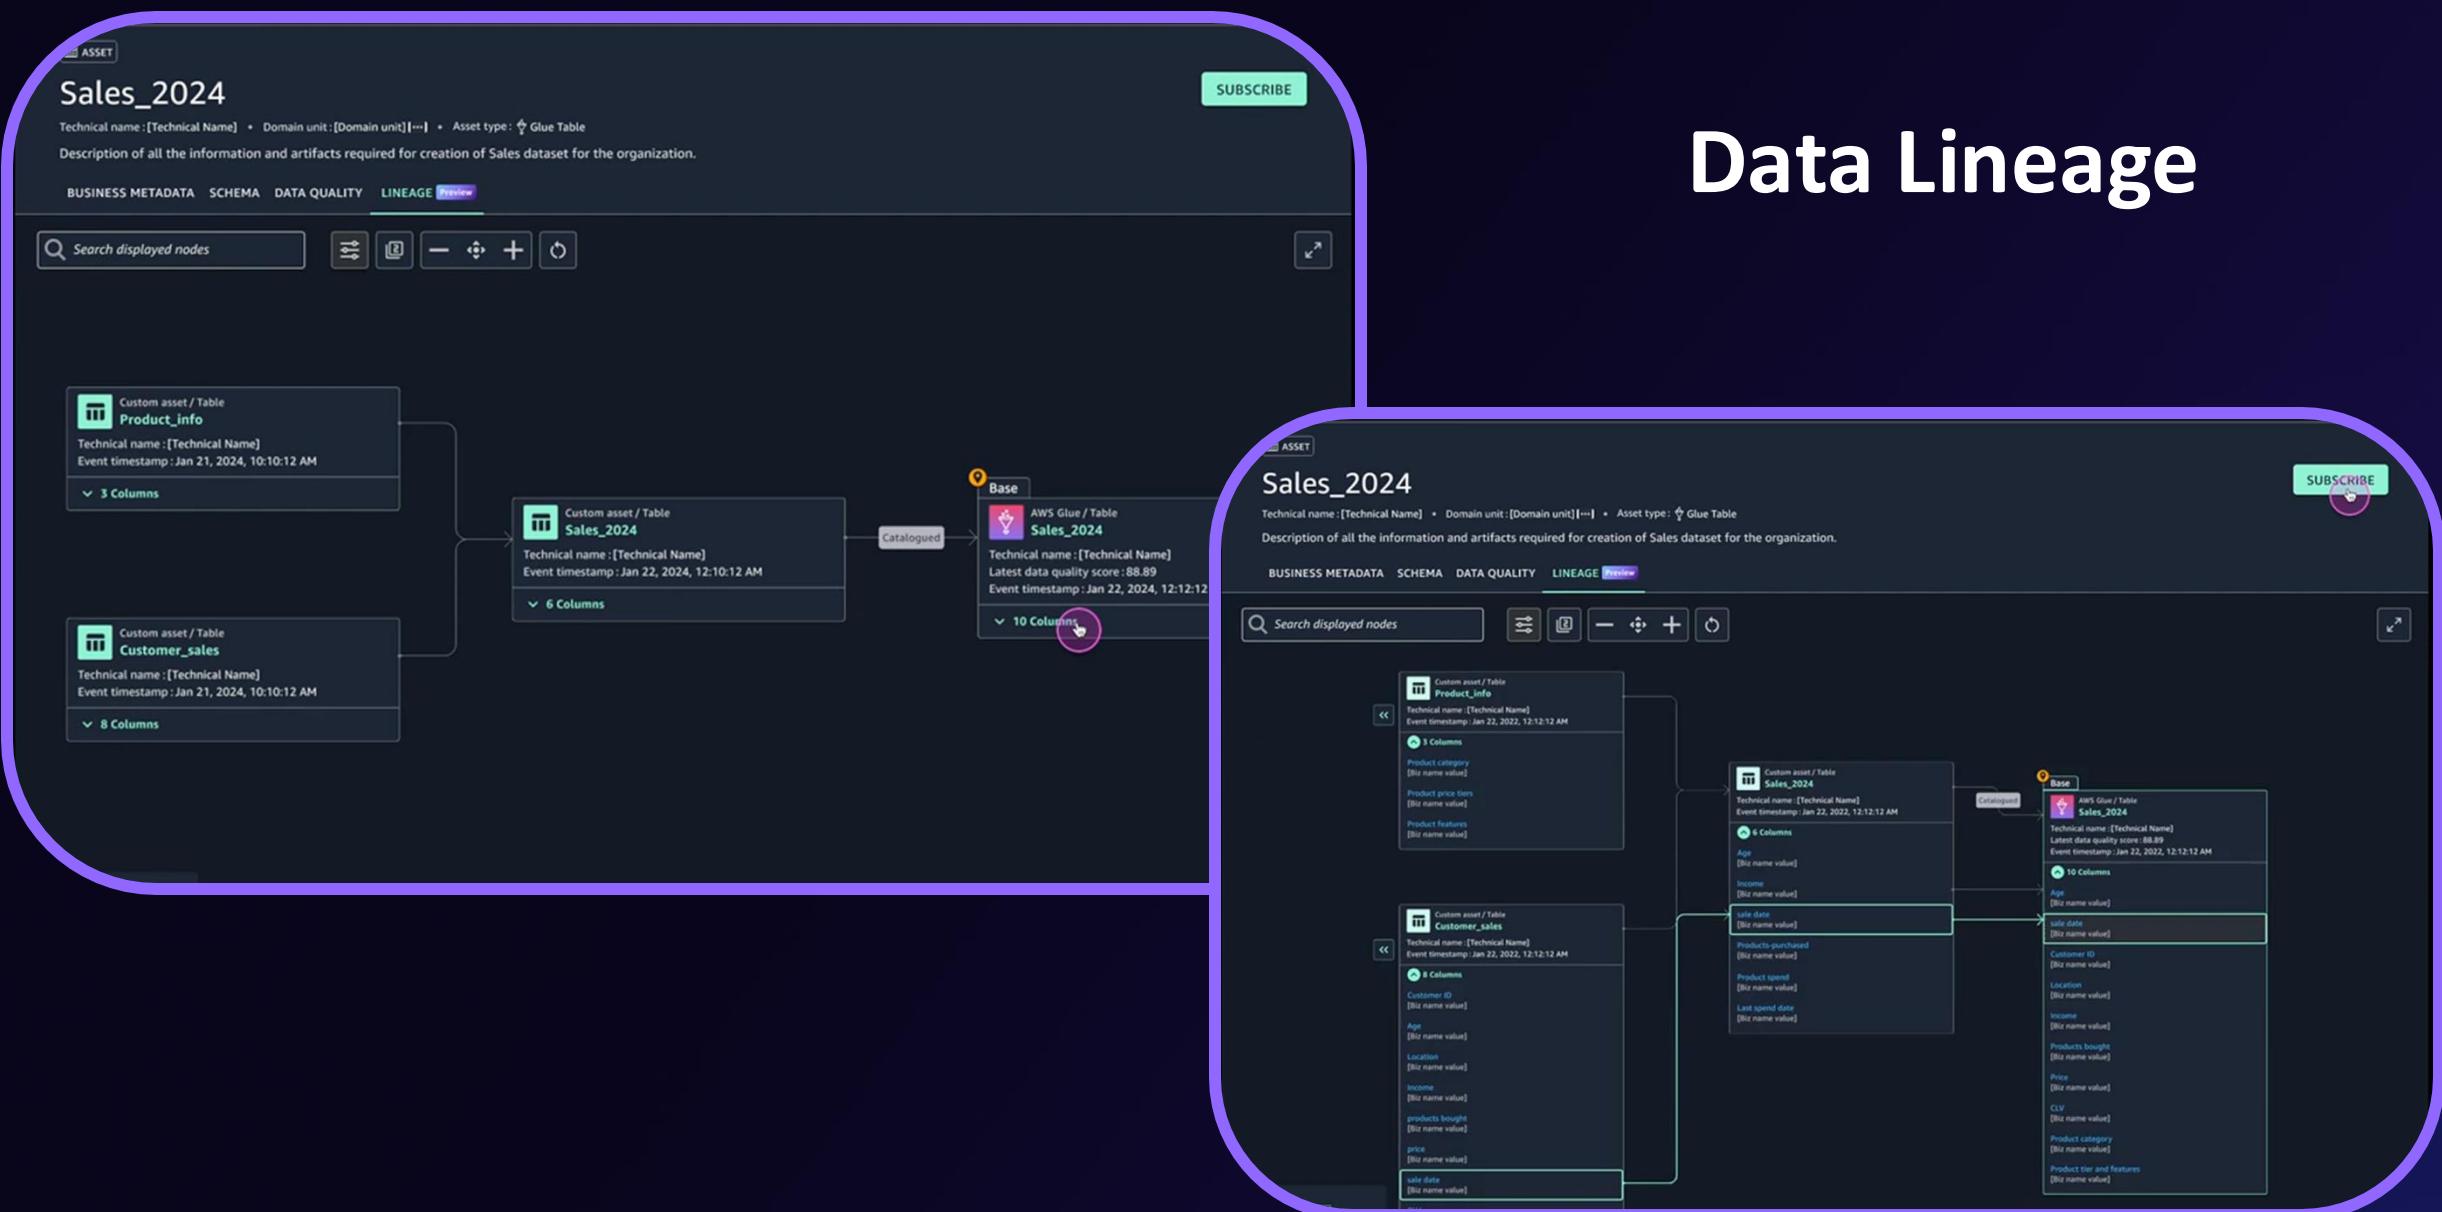
Link to learn more. Please help us improve by leaving feedback.

what data should I use to get a better idea of my product sales?

You have at least 10 Northstar resources in us-east-1 region in this AWS account.

Name	Asset type	Description
sales_2024	Glue Table	description: data about sales in maxce product
real_time_sales_data	Glue Table	description: This is a description of the asset
sales_transactions	Glue Table	description: This is a description of the asset
sales_purchases	Glue Table	description: This is a description of the asset
sales_items	Glue Table	description: This is a description of the asset
sales_data	Glue Table	description: This is a description of the asset
purchase_history	Glue Table	description: This is a description of the asset
asset	Glue Table	description: This is a description of the asset
asset	Glue Table	description: This is a description of the asset

Data Lineage





Unified Studio

Streaming | Data Processing | SQL Analytics | Model Development | Gen AI App Development | Business Intelligence | Search

Data & AI Governance

Lakehouse

GENERALLY AVAILABLE

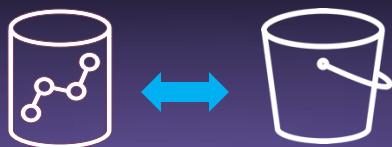


© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.



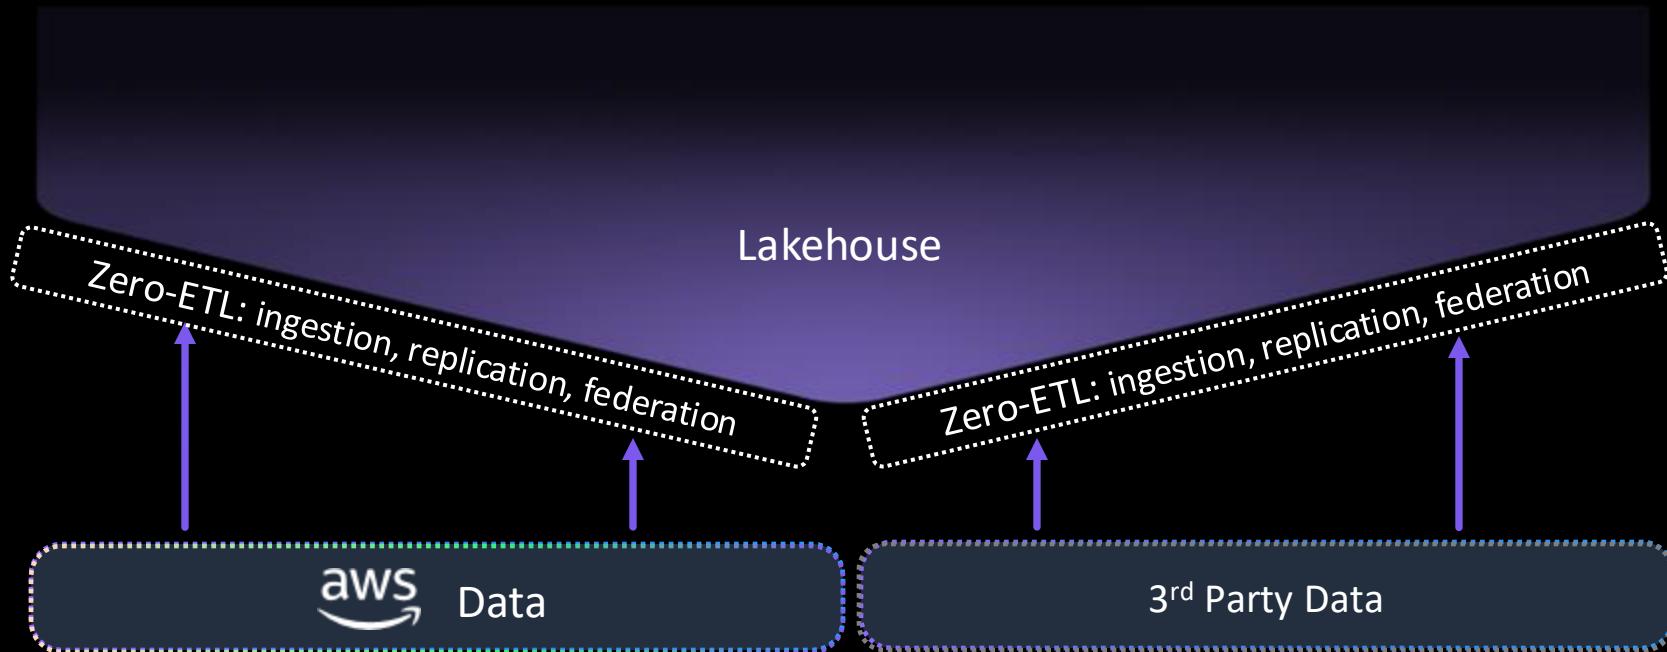
Amazon SageMaker Lakehouse

SageMaker Lakehouse Catalog



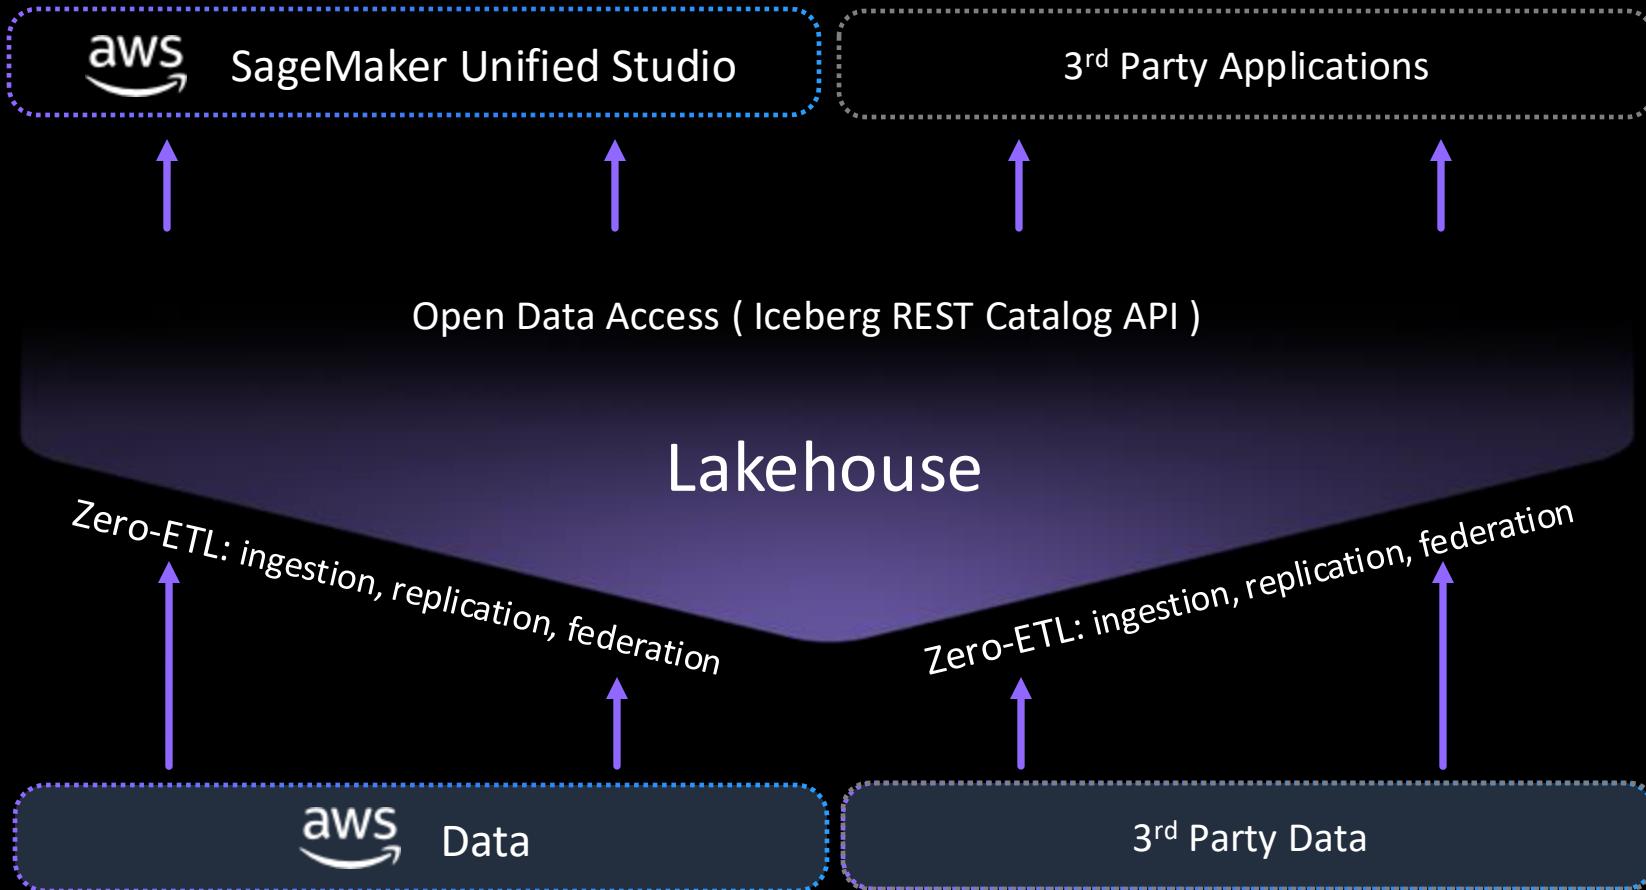


Amazon SageMaker Lakehouse





Amazon SageMaker Lakehouse



A single view of all your data in the Unified Studio





Amazon SageMaker

ZERO-ETL Integrations



Amazon Aurora
MySQL



Amazon Aurora
PostgreSQL



Amazon RDS for
MySQL



Amazon
DynamoDB

Amazon Data Sources

Facebook Ads

Instagram Ads

Salesforce

Salesforce Pardot

SAP

ServiceNow

ZenDesk

Zoho CRM

New

SaaS Applications



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

NEW

Amazon S3 Tables

Up to 3x better query performance and 10x increase in transactions per second for Apache Iceberg tables

SAGEMAKER SUPPORT FOR S3 TABLES IN PREVIEW



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.



Unified Studio Preview

Streaming | Data Processing | SQL Analytics | Model Development | Gen AI App Development | Business Intelligence | Search Analytics

comingsoon

comingsoon

comingsoon

Data & AI Governance

Lakehouse



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Laying the groundwork
for new innovations to take flight



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.
© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

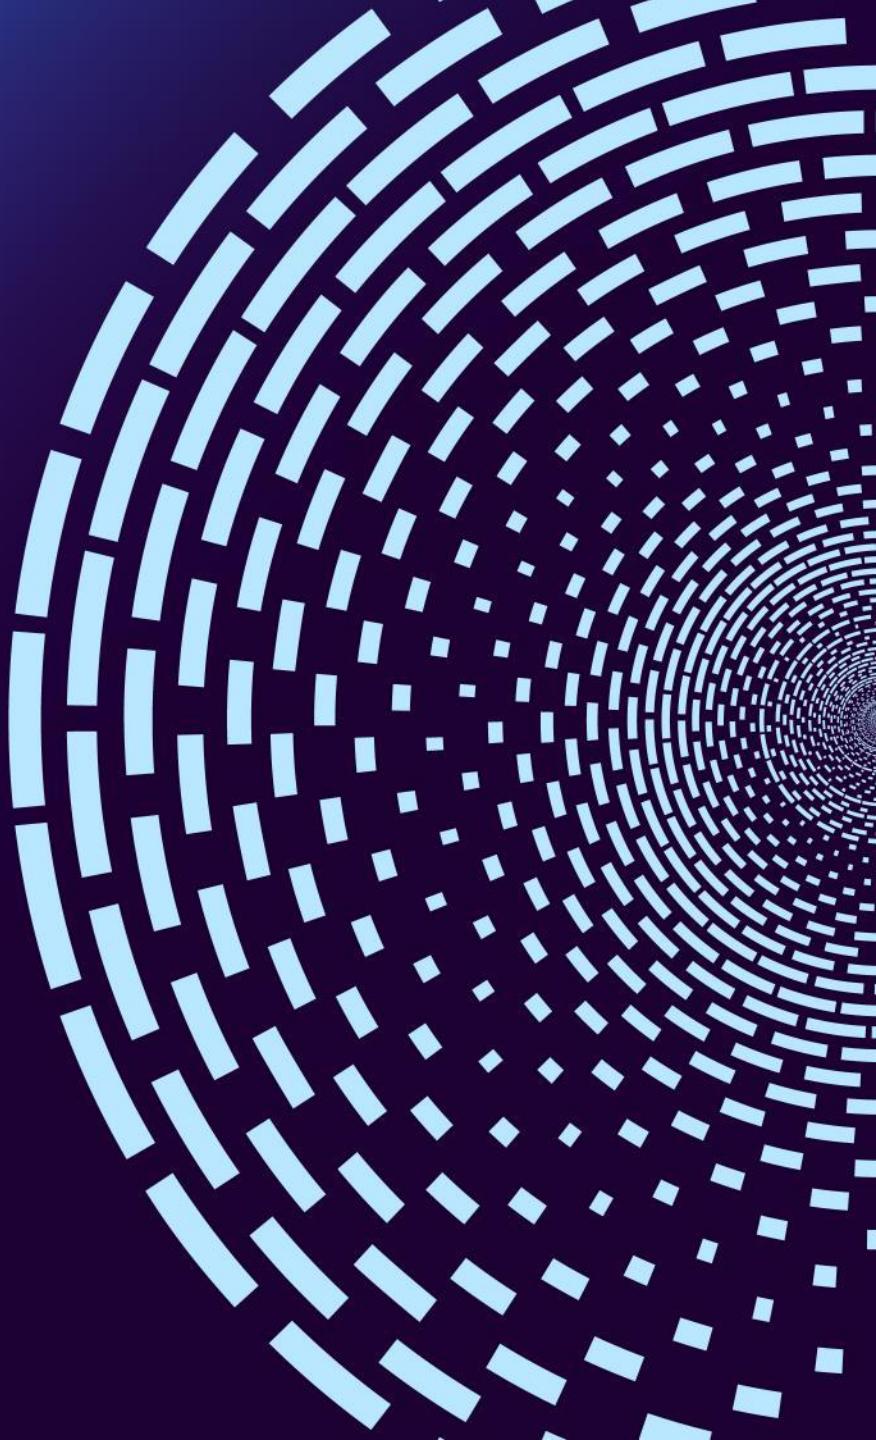


KEY 001

Data and generative AI driven transformation in India

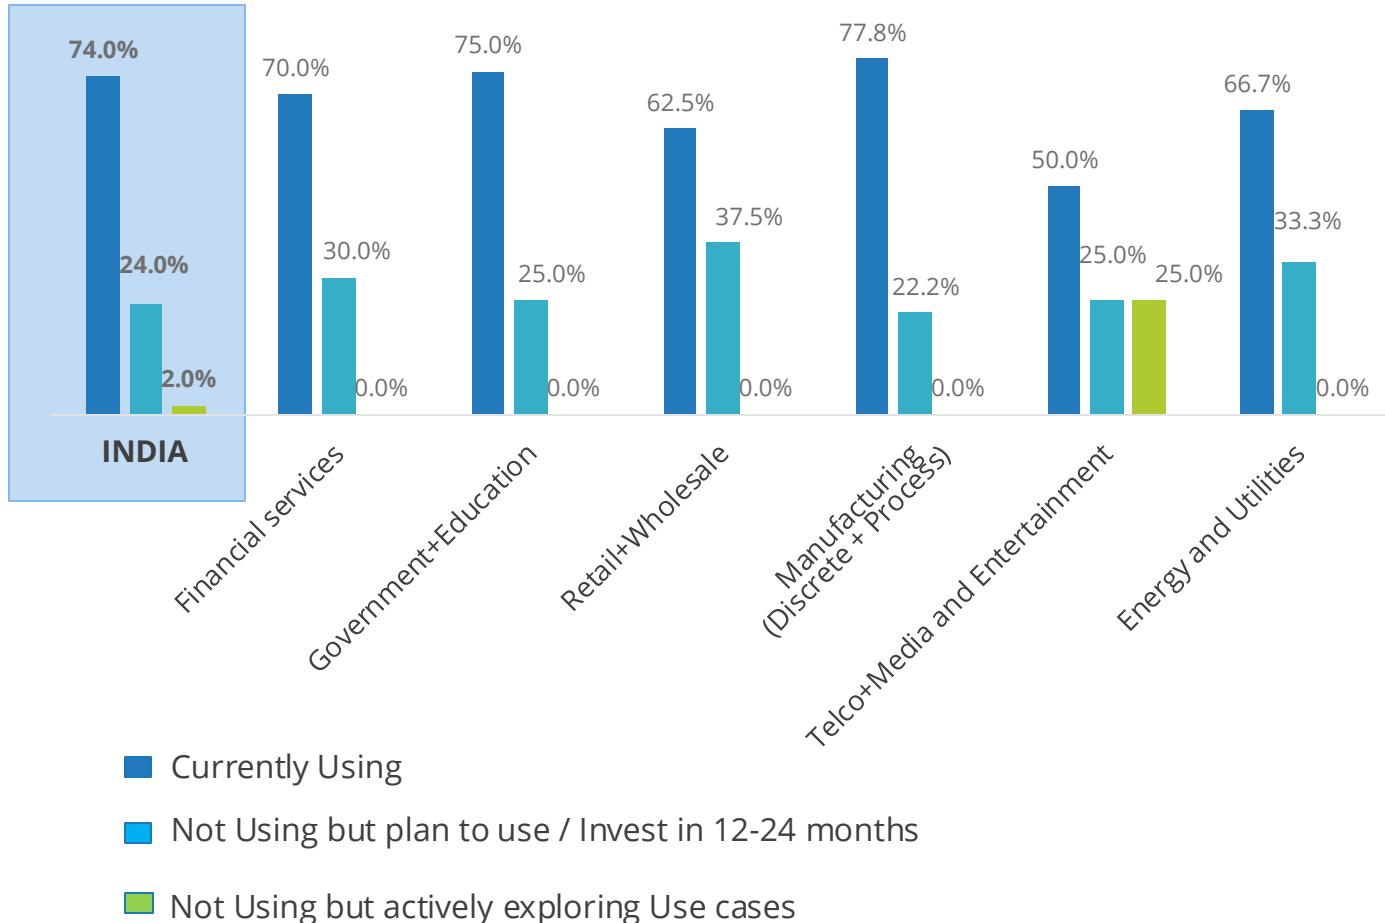
Vasant Rao

Managing Director
IDC India & South Asia

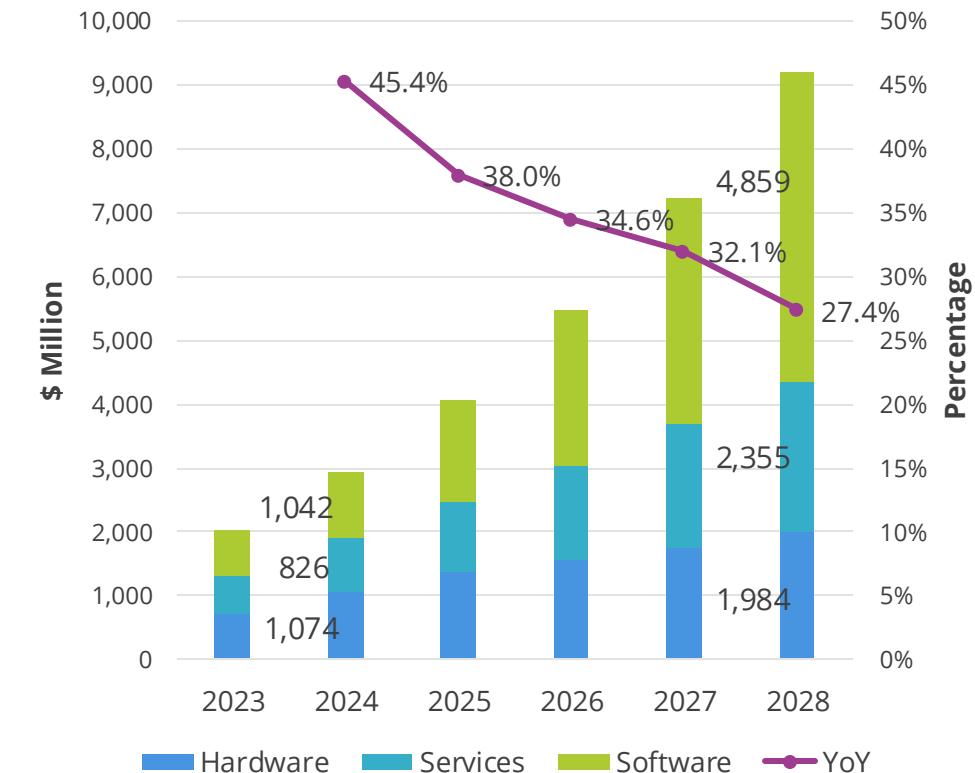


AI is well entrenched in India..

AI adoption in India in 0-24 months



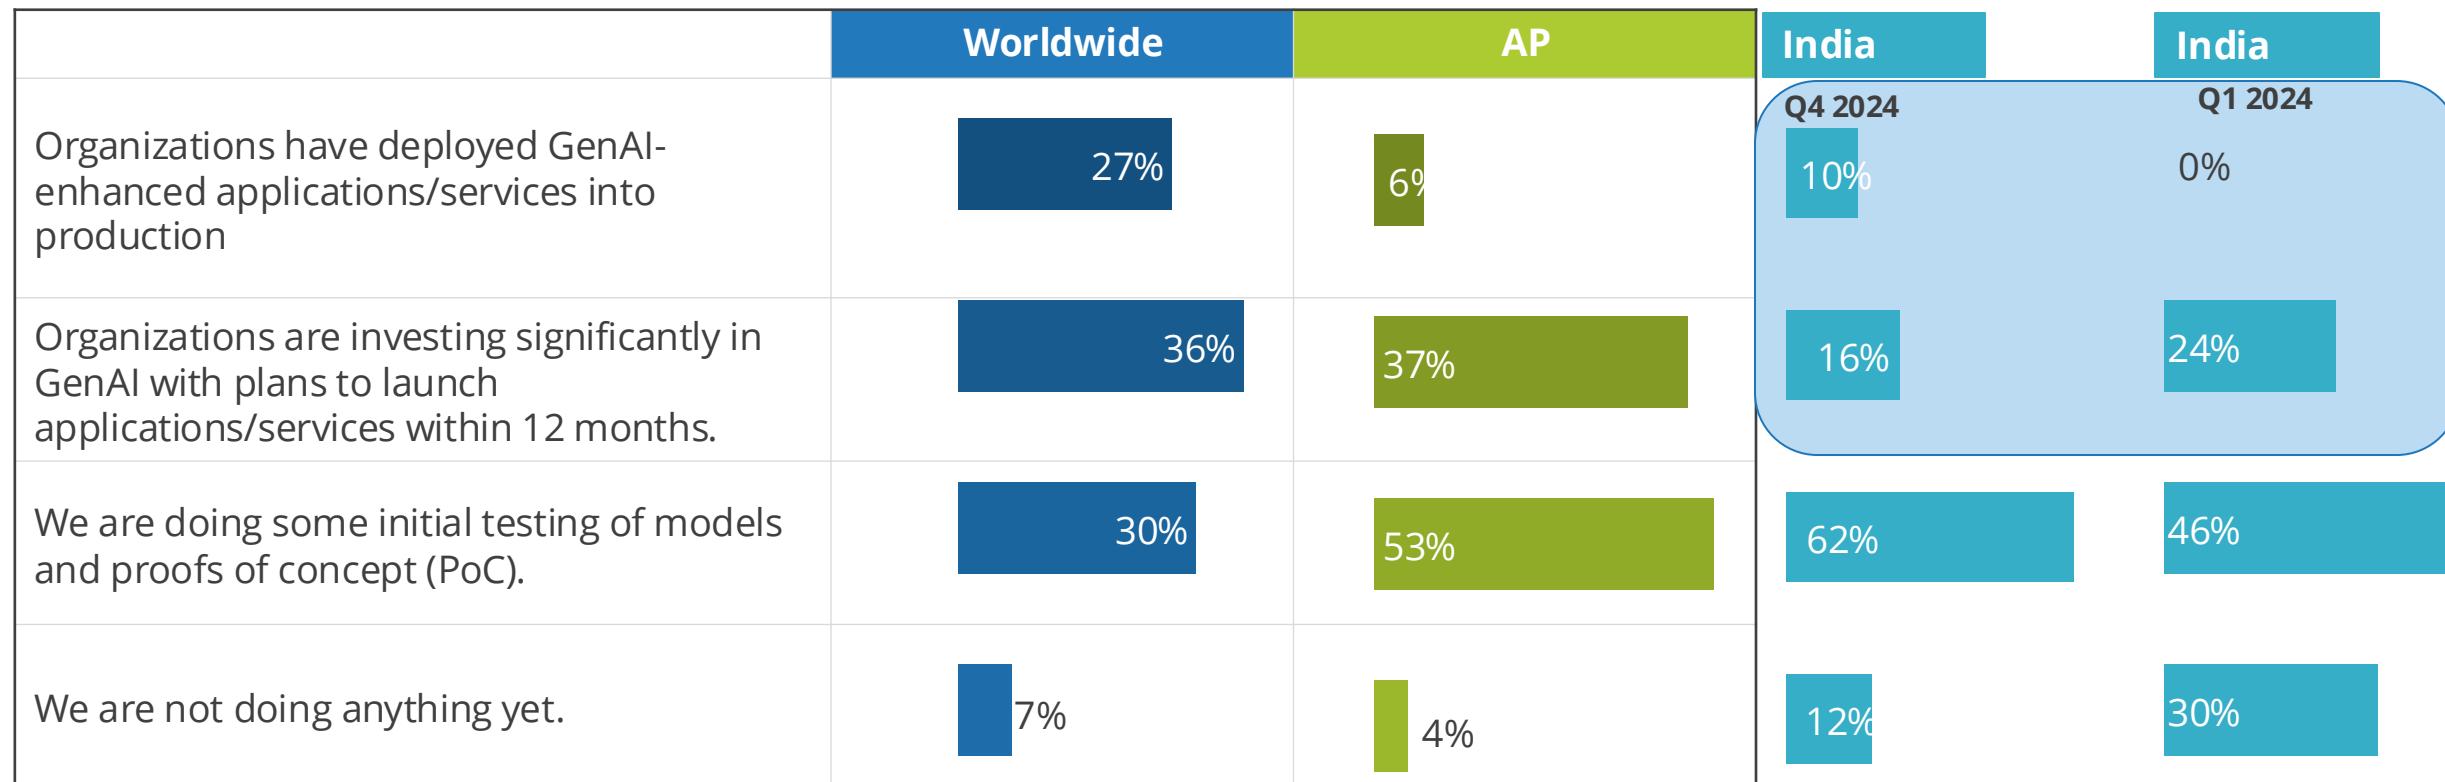
India Market Outlook:
Total AI Spending
CAGR 35.4%, and reach \$9.2Bn by 2028 from \$2.9
Bn in 2024



In 2024, Public Cloud Services in India accounted for about 85% of the AI Hardware Infrastructure market, with a projected CAGR of 24% over the next five years, while On-premise/Others grew at 15%.

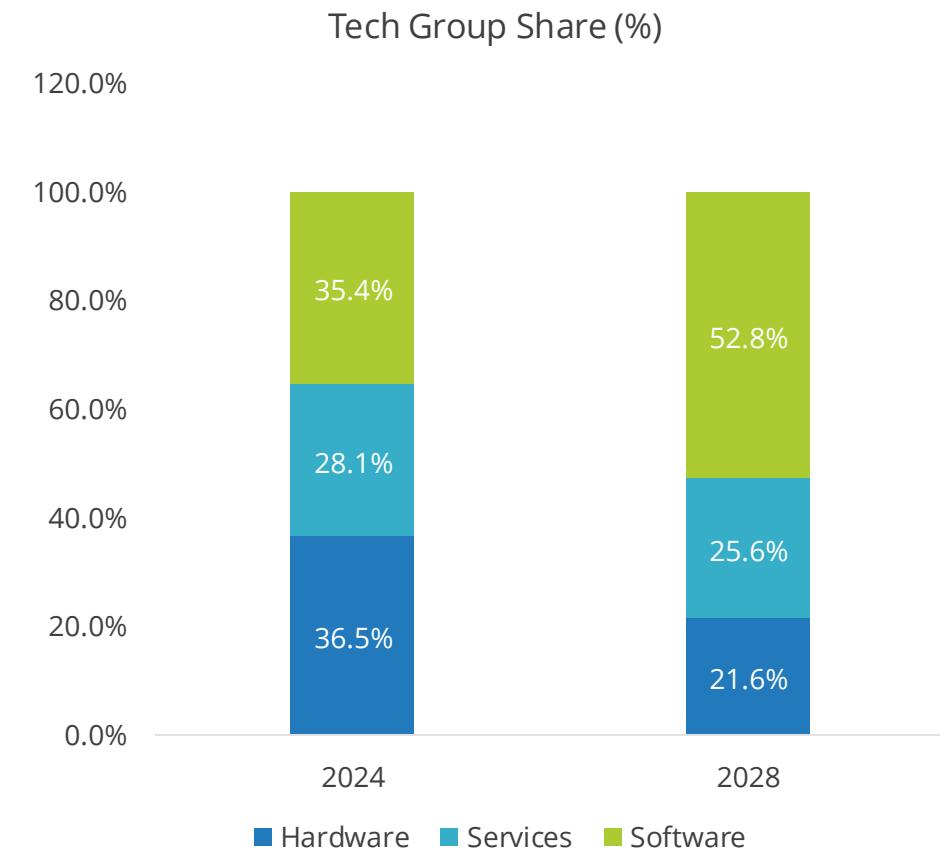
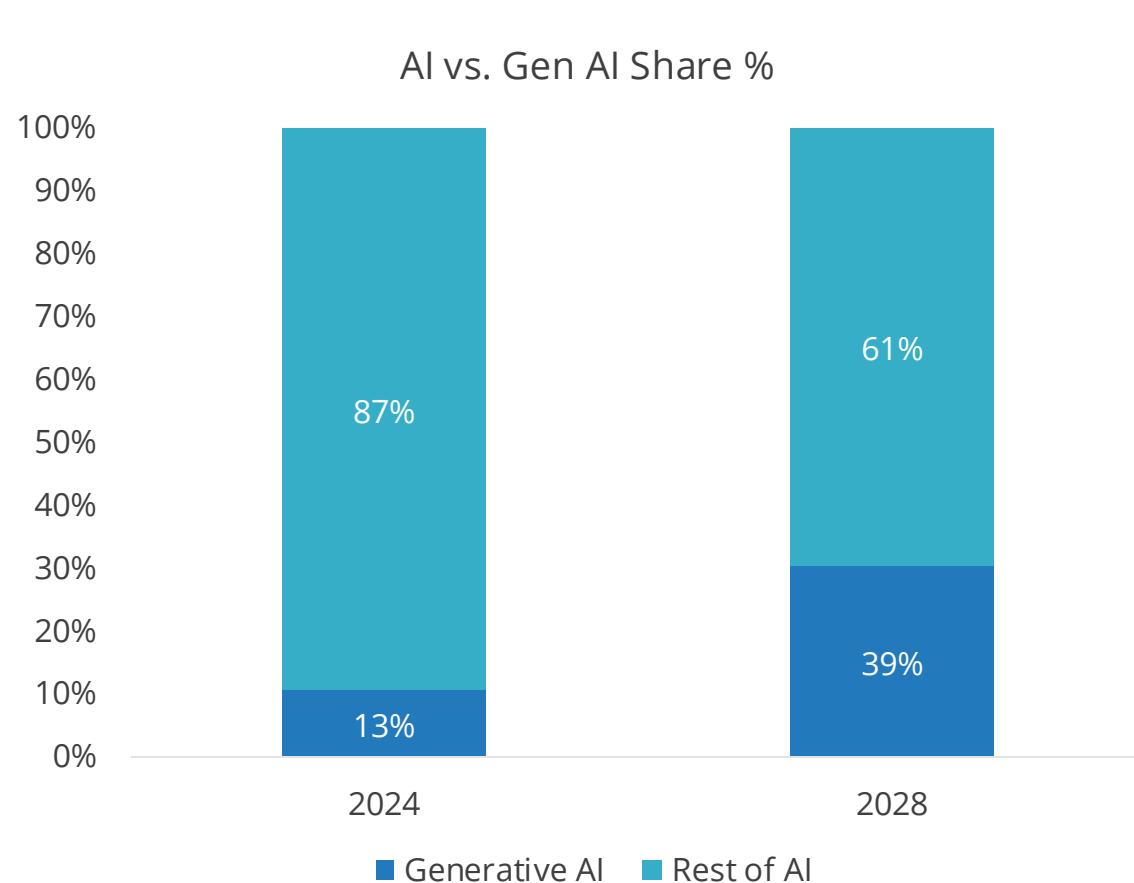
GenAI Adoption is WIP

Q. What is your organization's current state of evaluating or using Generative AI (GenAI)?



There is significant optimism on Gen AI

India GenAI Spend on **Software, Hardware, Services** to grow at 91% CAGR (2023-2028) to Reach \$3.6 billion by 2028



.. But Gen AI realities need to catch up



35%

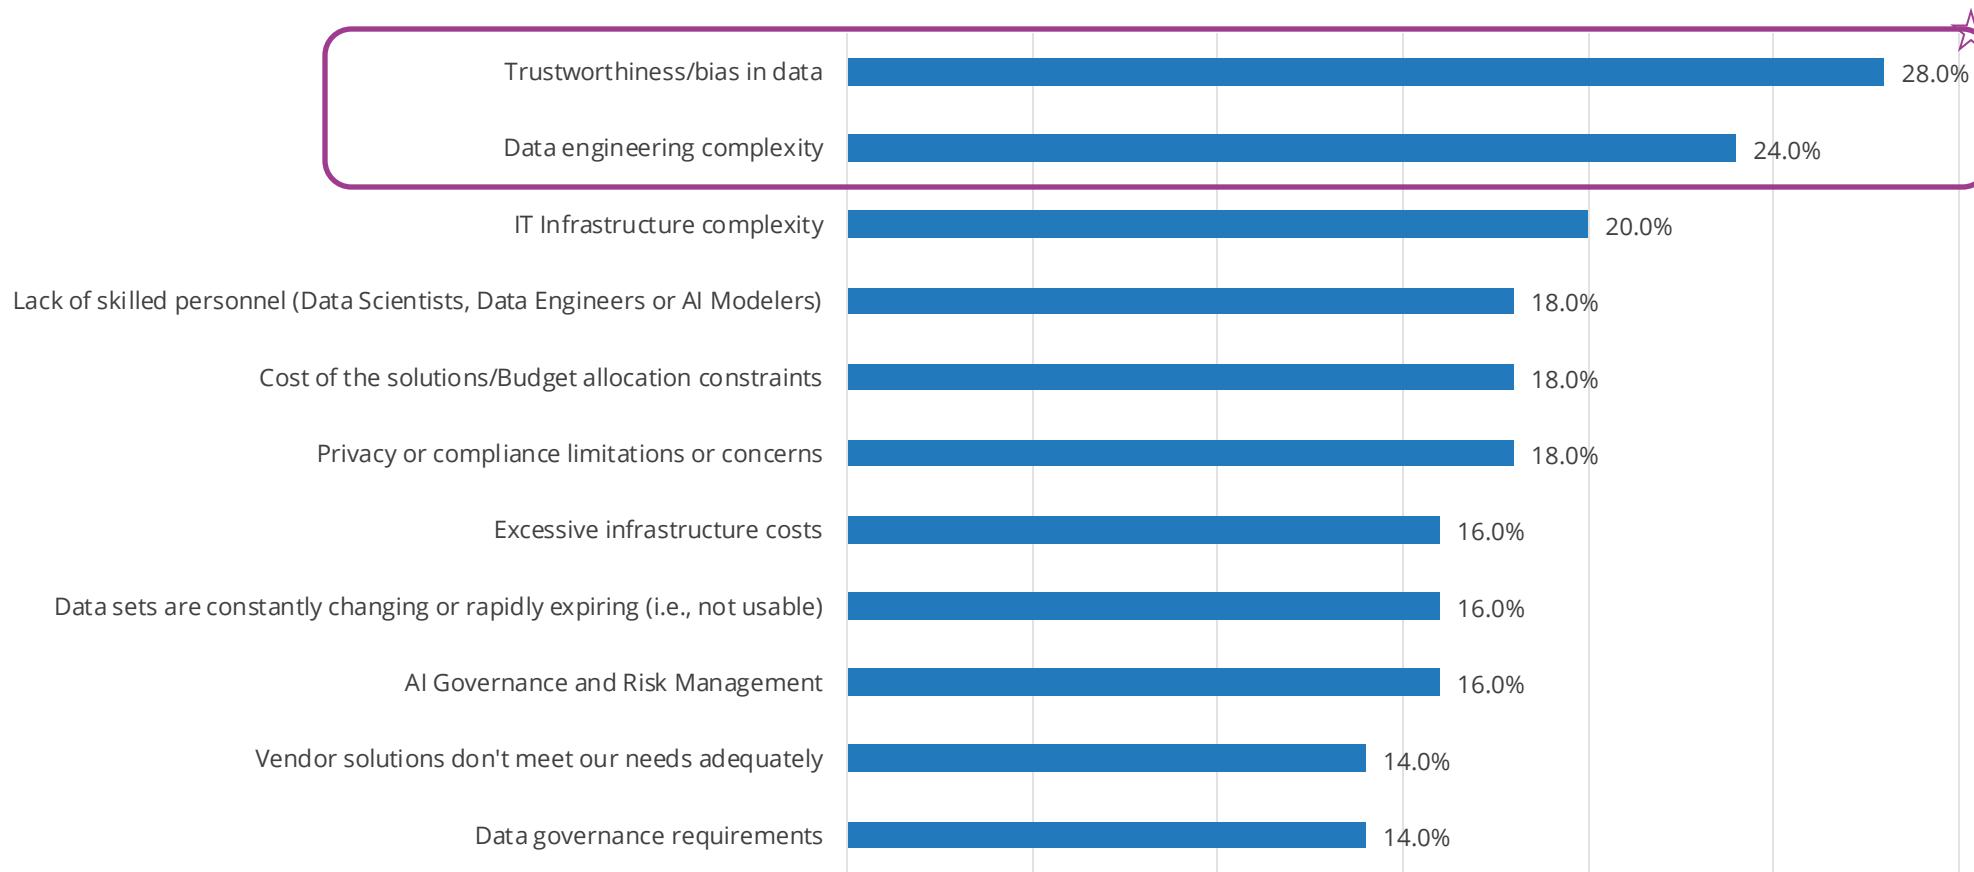
of Gen AI POCs did not succeed due to “poor quality / poorly labeled data sets”.

Lack of skills & Infrastructure, high costs ranked similarly high in reasons for poor success

So what gives?

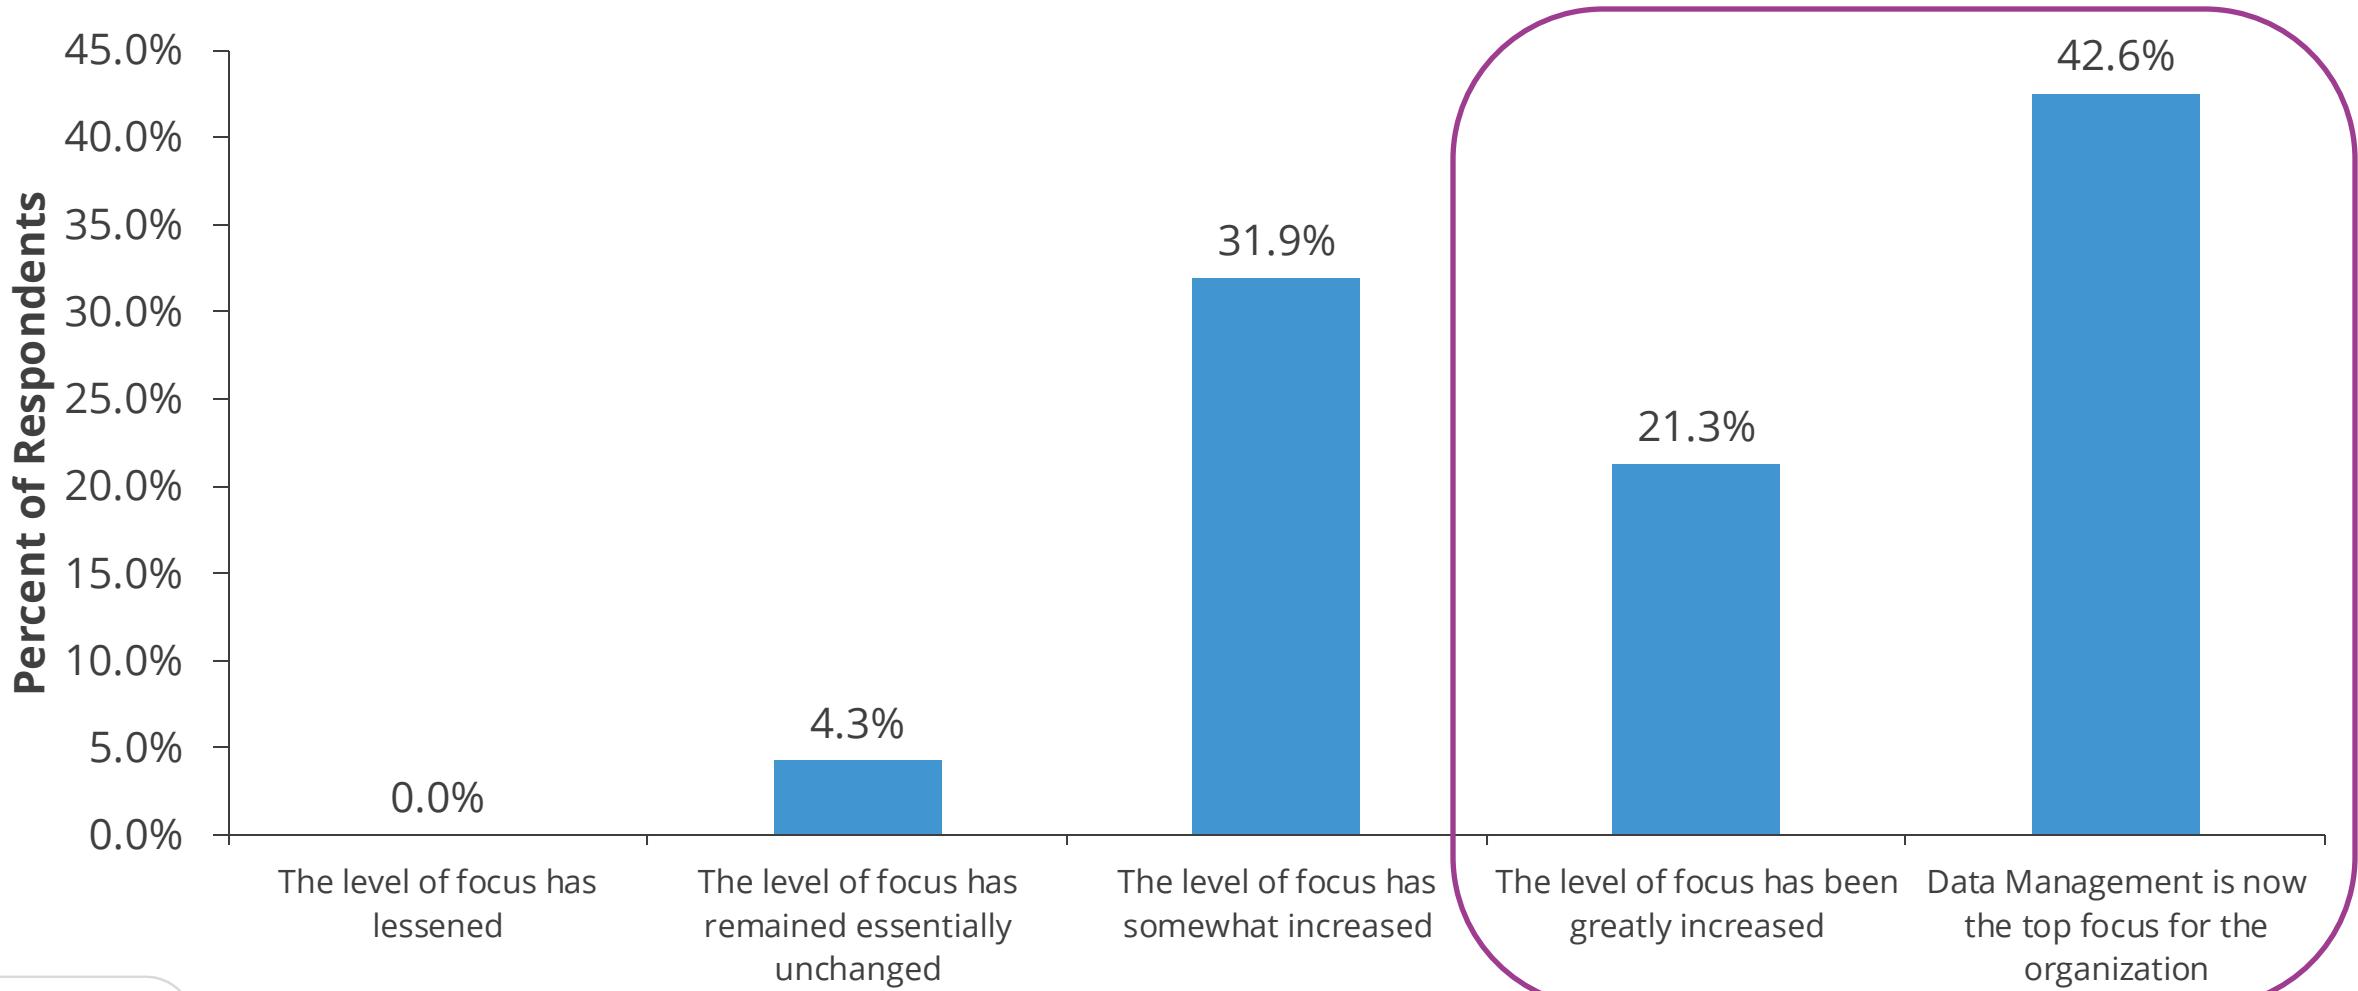
Successful AI Adoption Requires: Addressing data biases to ensure fairness and accuracy. Overcoming data engineering complexity through infrastructure and talent investment.

Q. Overall Concern - Please select the top 3 challenges involved in implementing AI technology in your organization.



Not surprisingly, Data Management is a key Focus Areas for Organizations in India

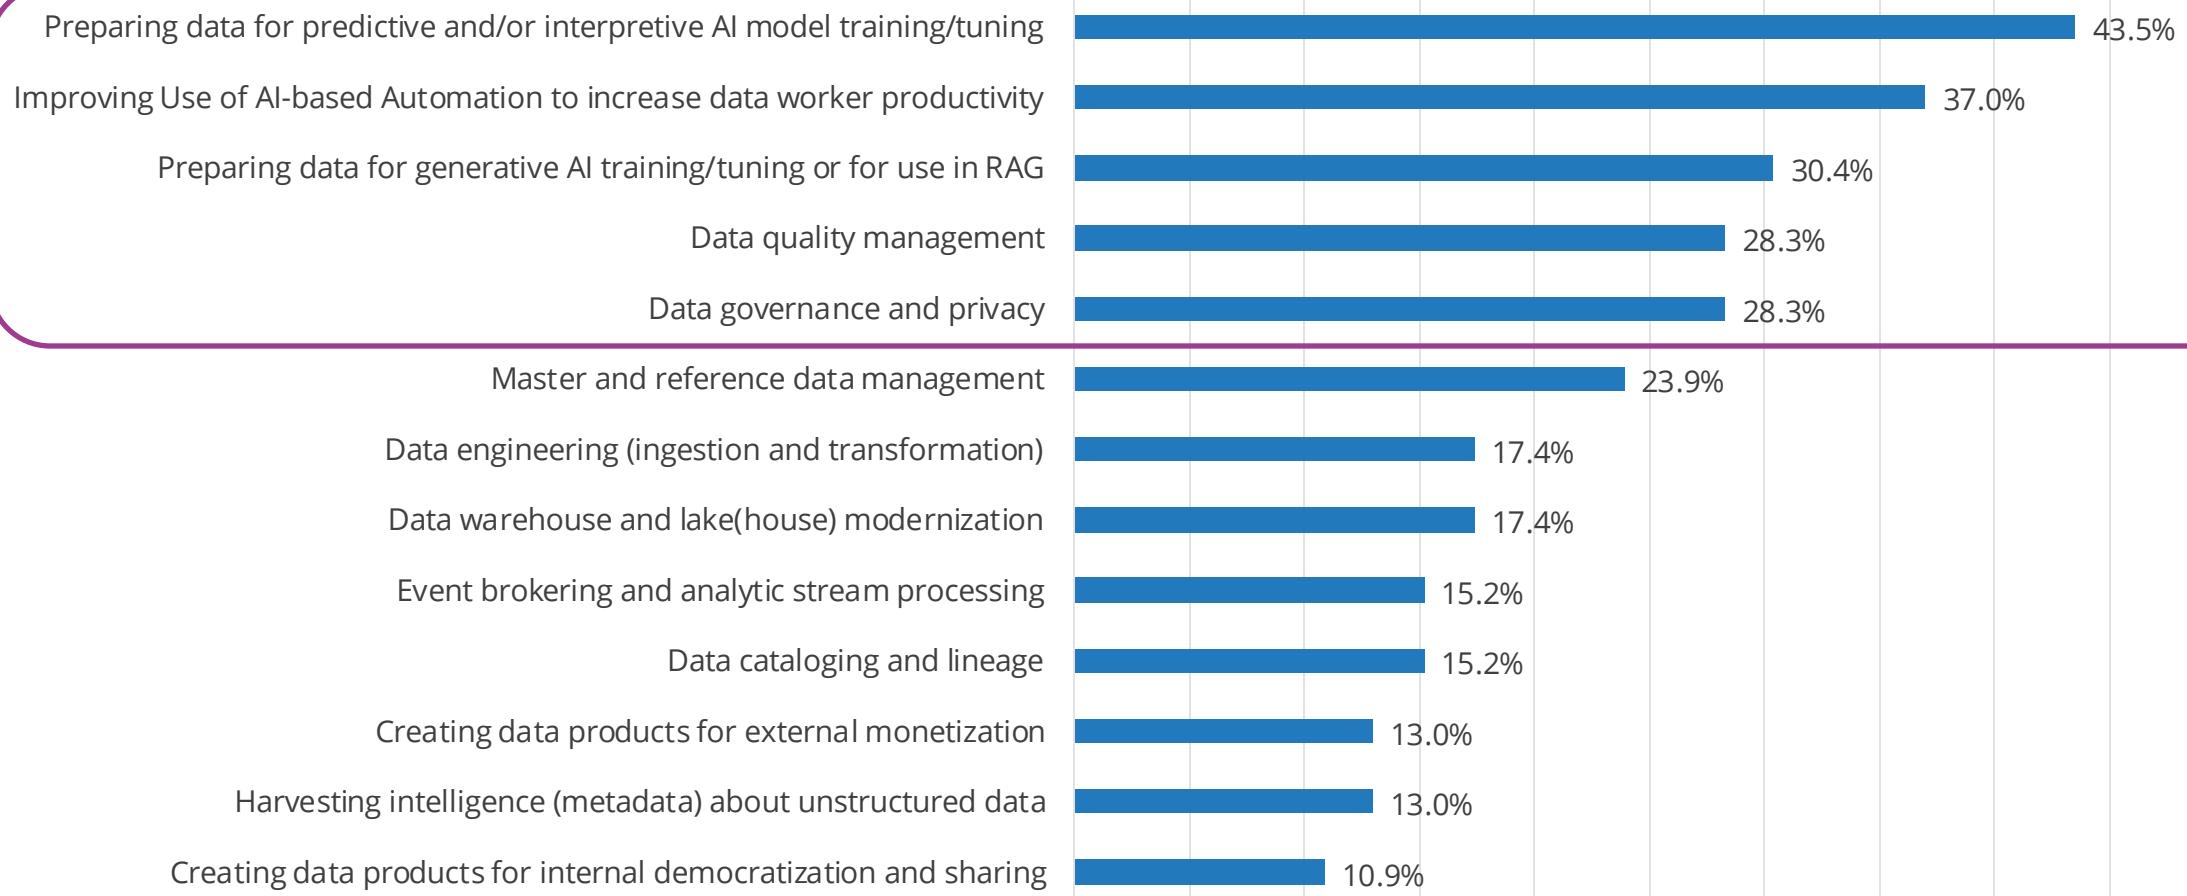
Q: Since the emergence of Generative AI (GenAI), how has the organization's focus on Data Management changed?



In Data, what will Indian Organization's invest in 2025?

Q: What are the overall organization's top 3 strategic investment priorities for data in the next 12-18 months?

Strategic Investment Priorities for Data



.. And what will be the key trends in AI Platforms?



LLMs + Classic AI models



Build once deploy anywhere AI applications



Integration of vendor products



AI tools + Applications



GenAI feature stores



Agentic workflows will be precursor to autonomous AI applications



AI based System Architecture,
Autonomous agentic workflows will necessitate AI system Architects



Model Certification

IDC FutureScape Predictions 2025 and Beyond

By 2027, without strategic and organizational pivots, frustration with low ROI due to poor planning and misallocated spending will cause 30% of A1000 organizations to reduce GenAI investments.

By 2027, the use of data-as-a-product architectures will significantly break down data silos in 50% of large enterprises in APEJ.

By 2028, data collaboration via data exchanges and/or data clean rooms will penetrate 70% of A2000 organizations.

By 2027, over 50% of organizations in APEJ will be ready to use their data with GenAI because of the increased adoption of data platforms for data storage, standardization, and access.

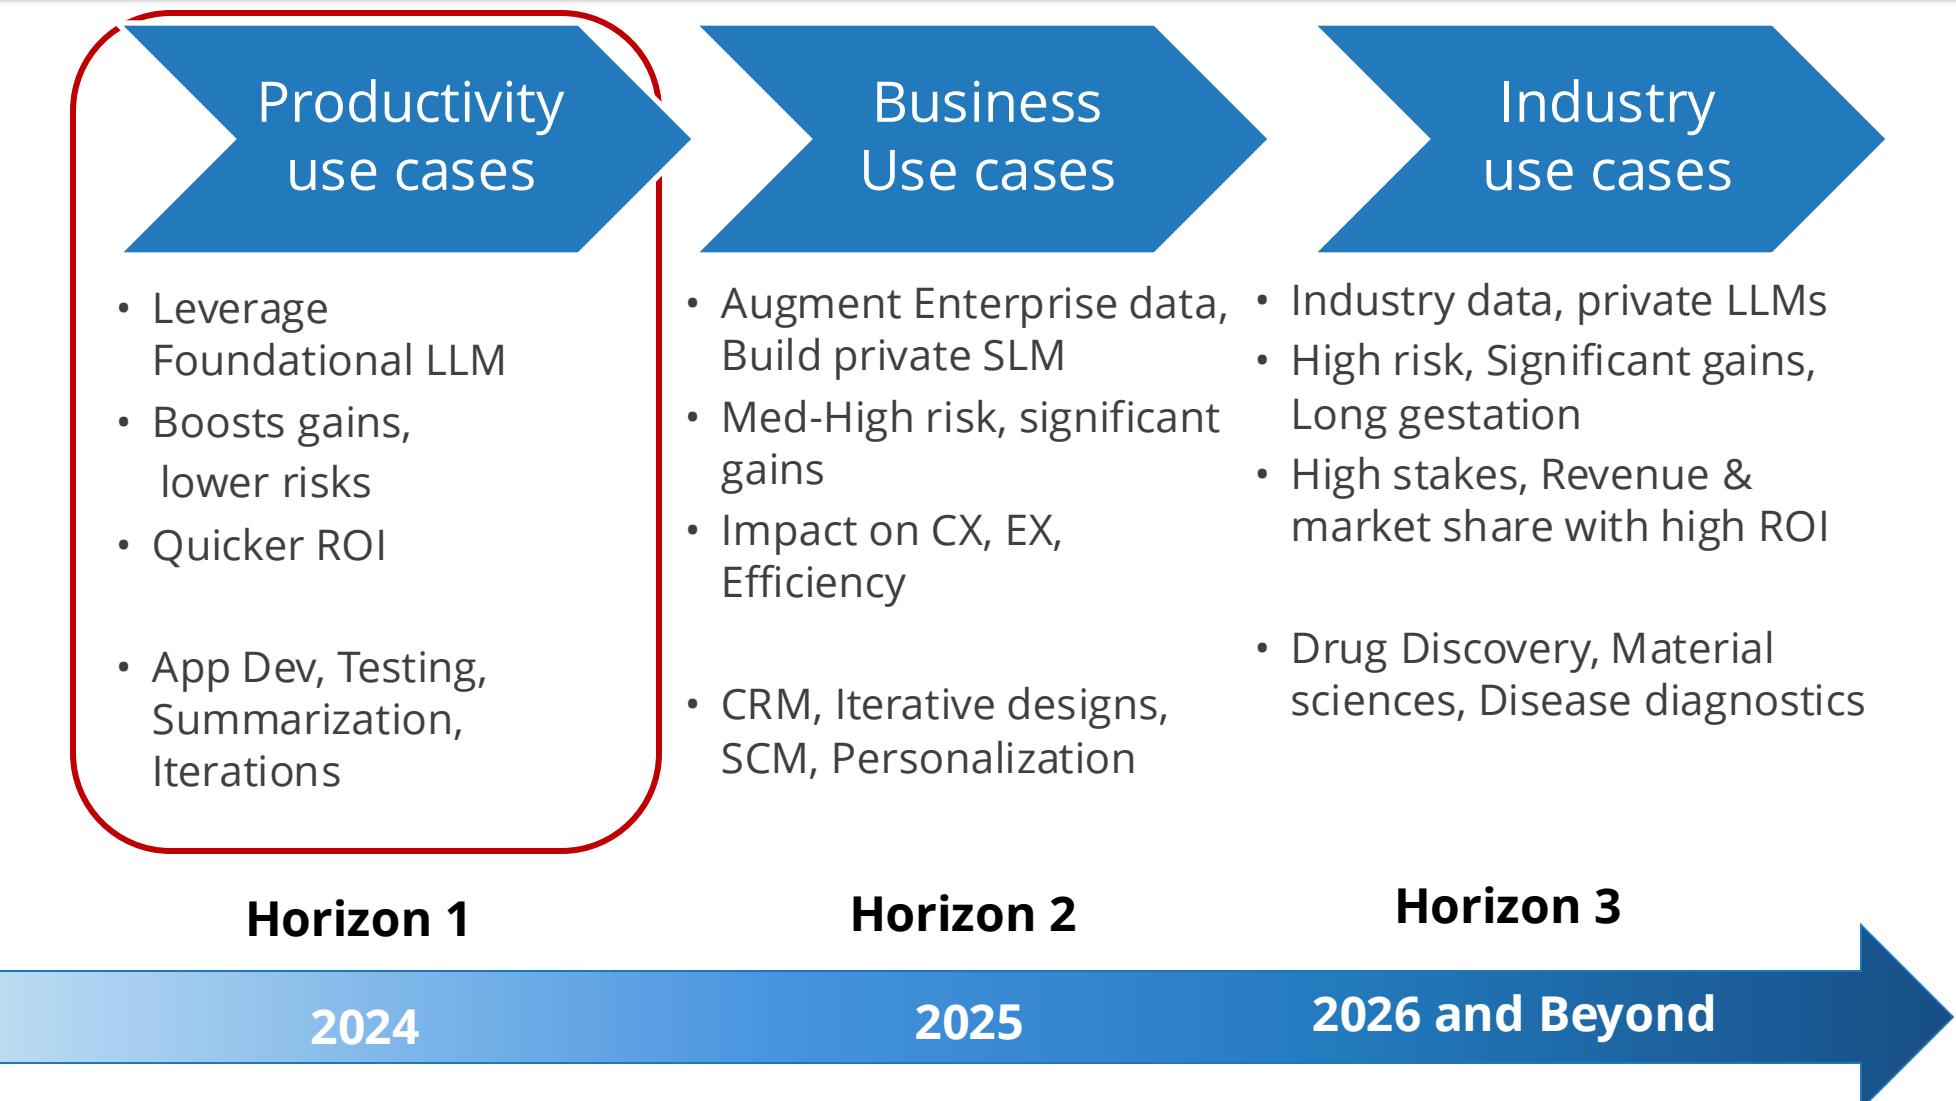
By 2028, 75% of database implementations in APEJ will be a part of broader ecosystems of data technologies rather than standalone implementations.



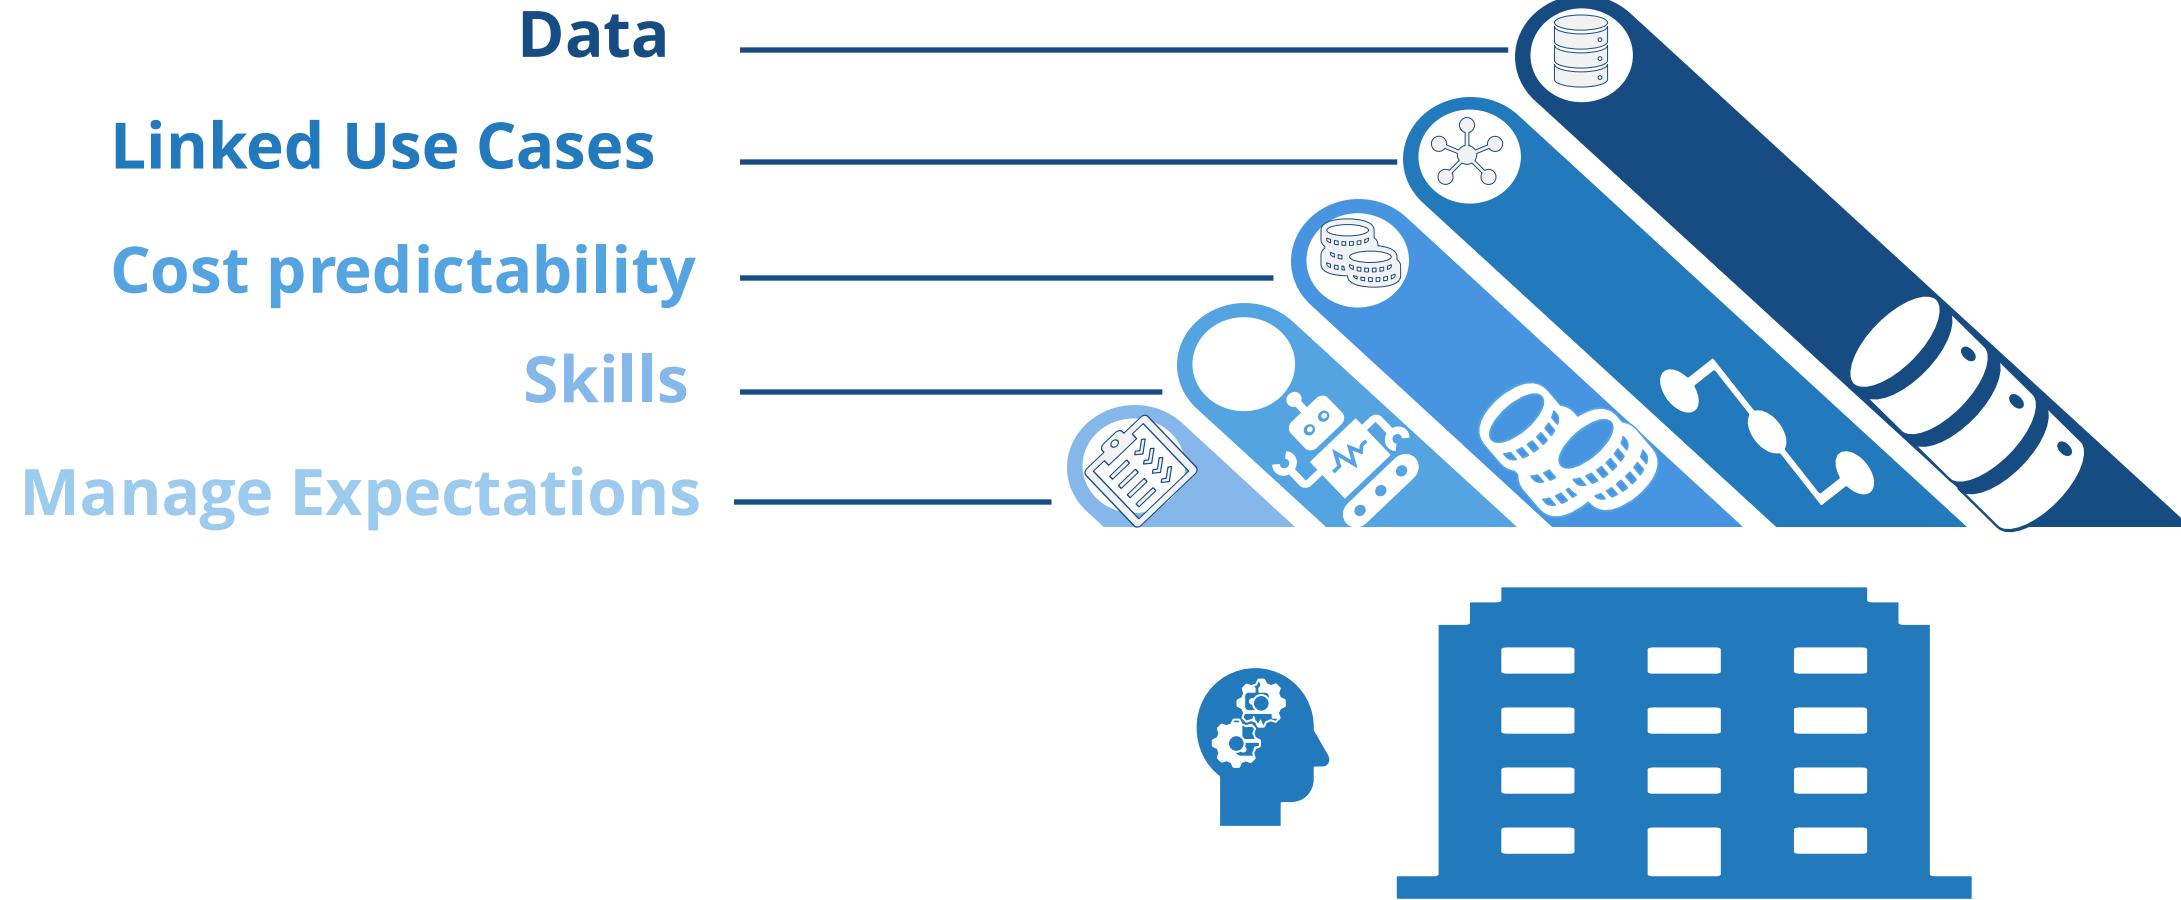
Source: IDC FutureScape: Worldwide Services 2025 Predictions — Asia/Pacific (Excluding Japan) Implications; for Prediction 1

Source: IDC FutureScape: Worldwide Data and Analytics 2025 Predictions — Asia/Pacific (Excluding Japan) Implications

Gen AI – Path to Impact



Recommendations





Thank you!