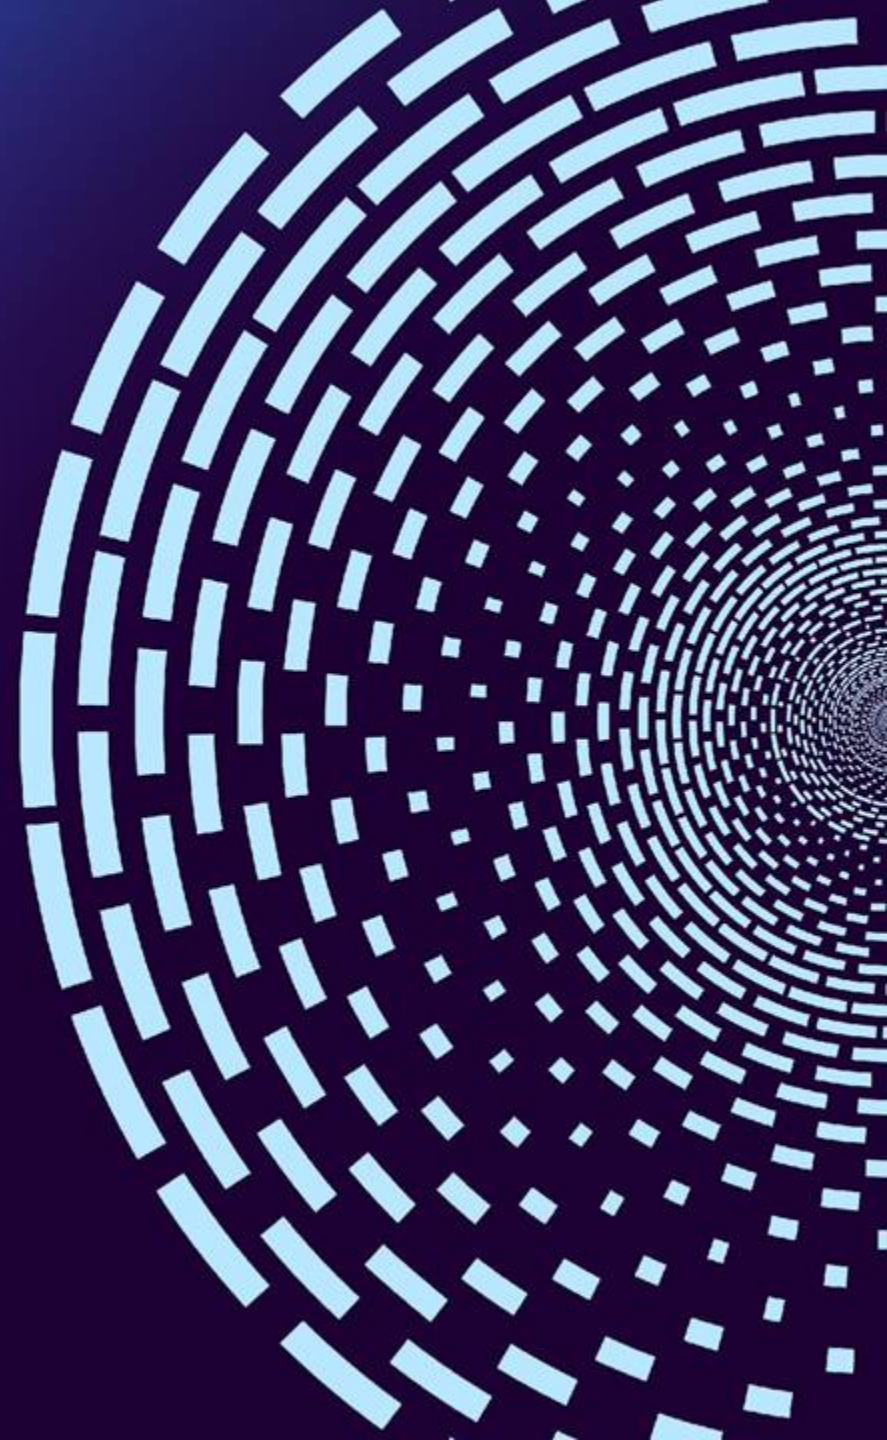




AI Conclave

Online



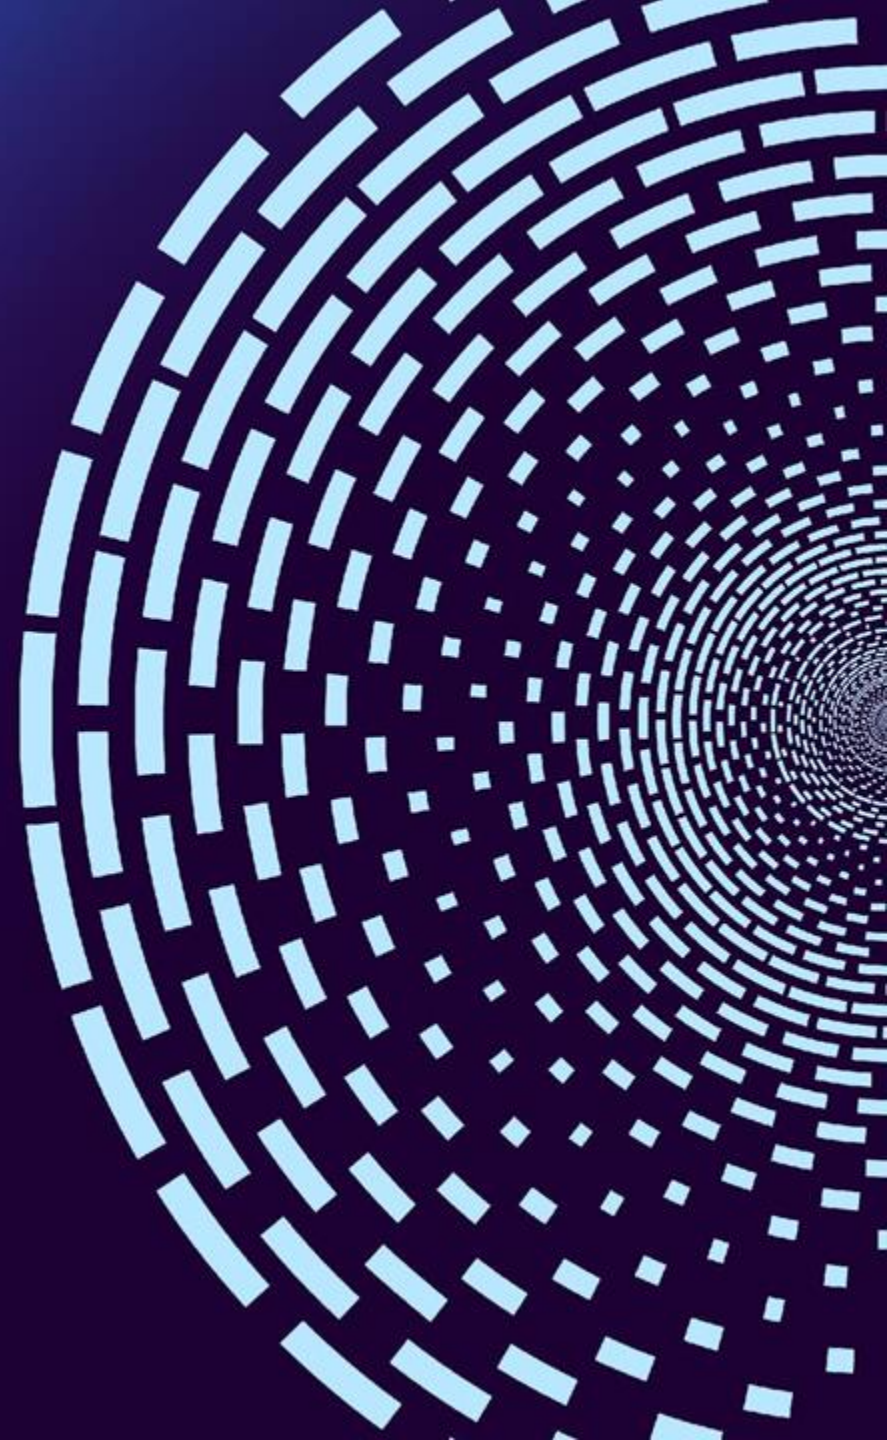


AIOT103

Exploring agents with Amazon Bedrock

Sahil Verma

Senior GTM Specialist Solutions Architect,
Generative AI
AWS India



Agenda

- Amazon Bedrock
- Introduction to agents
- Building agents with Amazon Bedrock
- Deep dive on the agent's capabilities, solution and patterns
- Demo – Build an agent from scratch
- Resources





Amazon Bedrock

The easiest way to build and scale generative AI applications with powerful tools and foundation models

Choice of leading FMs through a single API

Model import, distillation and fine-tuning

Generative AI Tools – Knowledge Bases (RAG), Guardrails, Flows, and Agents

Security, privacy, and data governance

2025 – The year of agents



What is an AI Agent?



Intelligent,
autonomous
systems



Plan,
reason,
and act



Access to
enterprise
data



Ability to
use tools

Momentum behind Amazon Bedrock Agents



Investment &
medical
research



Insurance
claims
processing



Customer
experience



Marketing
assistants



Root cause
analysis


What generative AI customers are asking for



Help me automate
complex workflows



Help me move
faster



Help me find
more robust and
scalable solutions

Amazon Bedrock Agents

Agentic building blocks




Choice of foundation models



Memory, Knowledge Bases,
Guardrails



Tools and action groups



Trace, debug, and observability

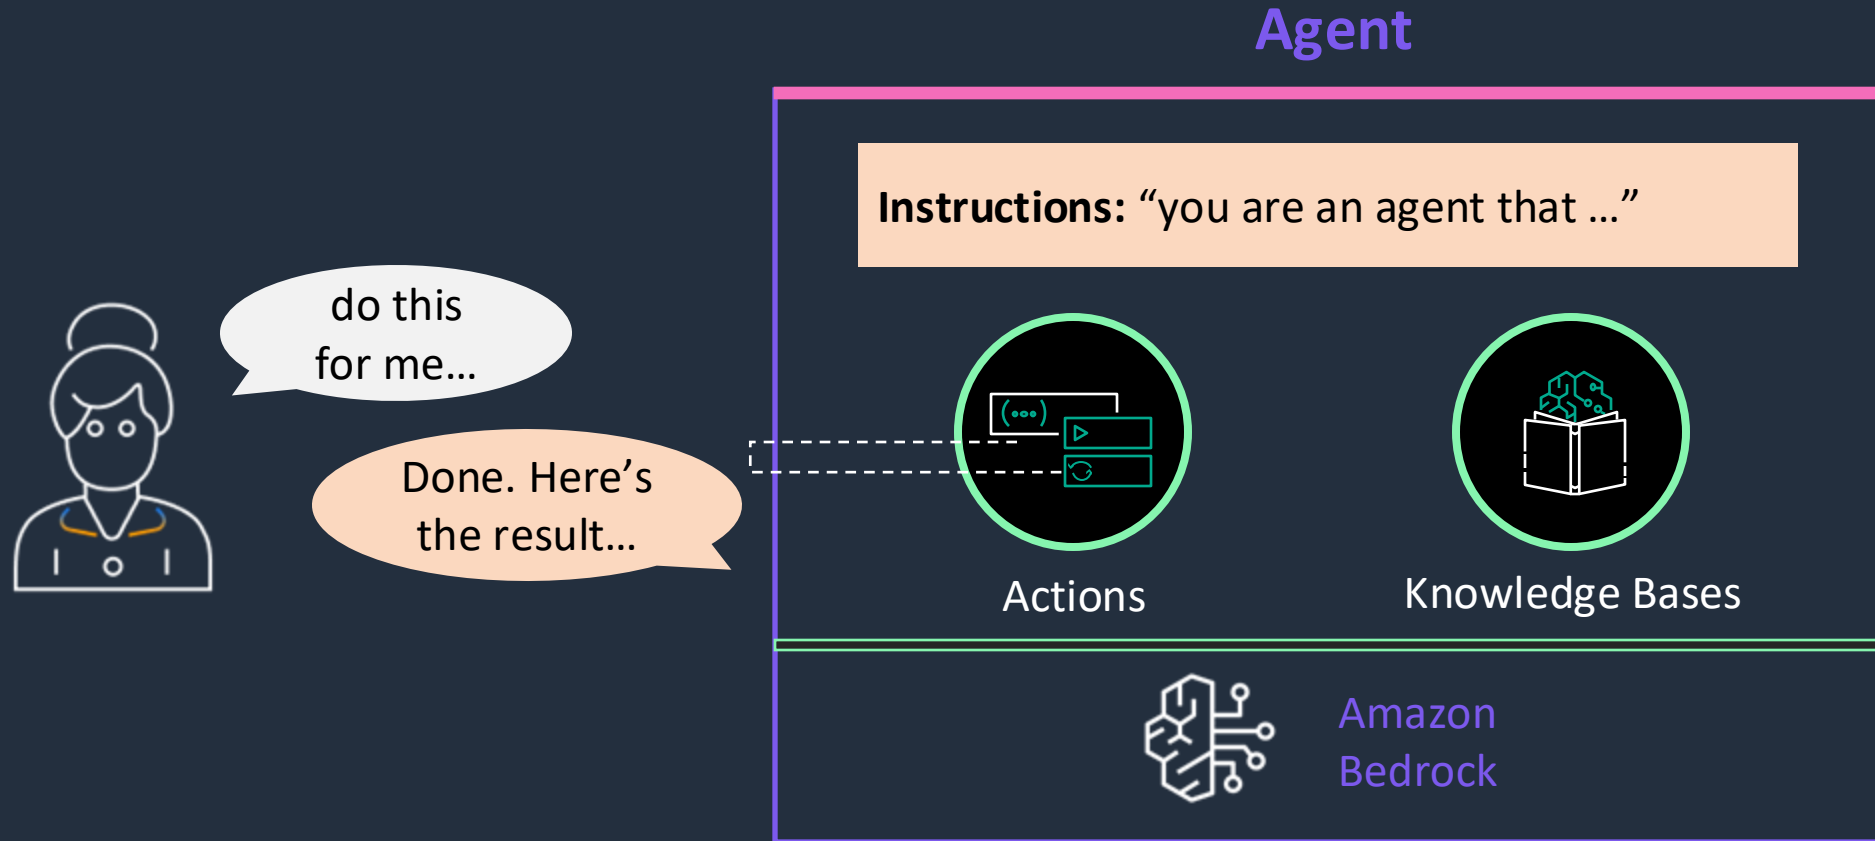


Multi-agent collaboration

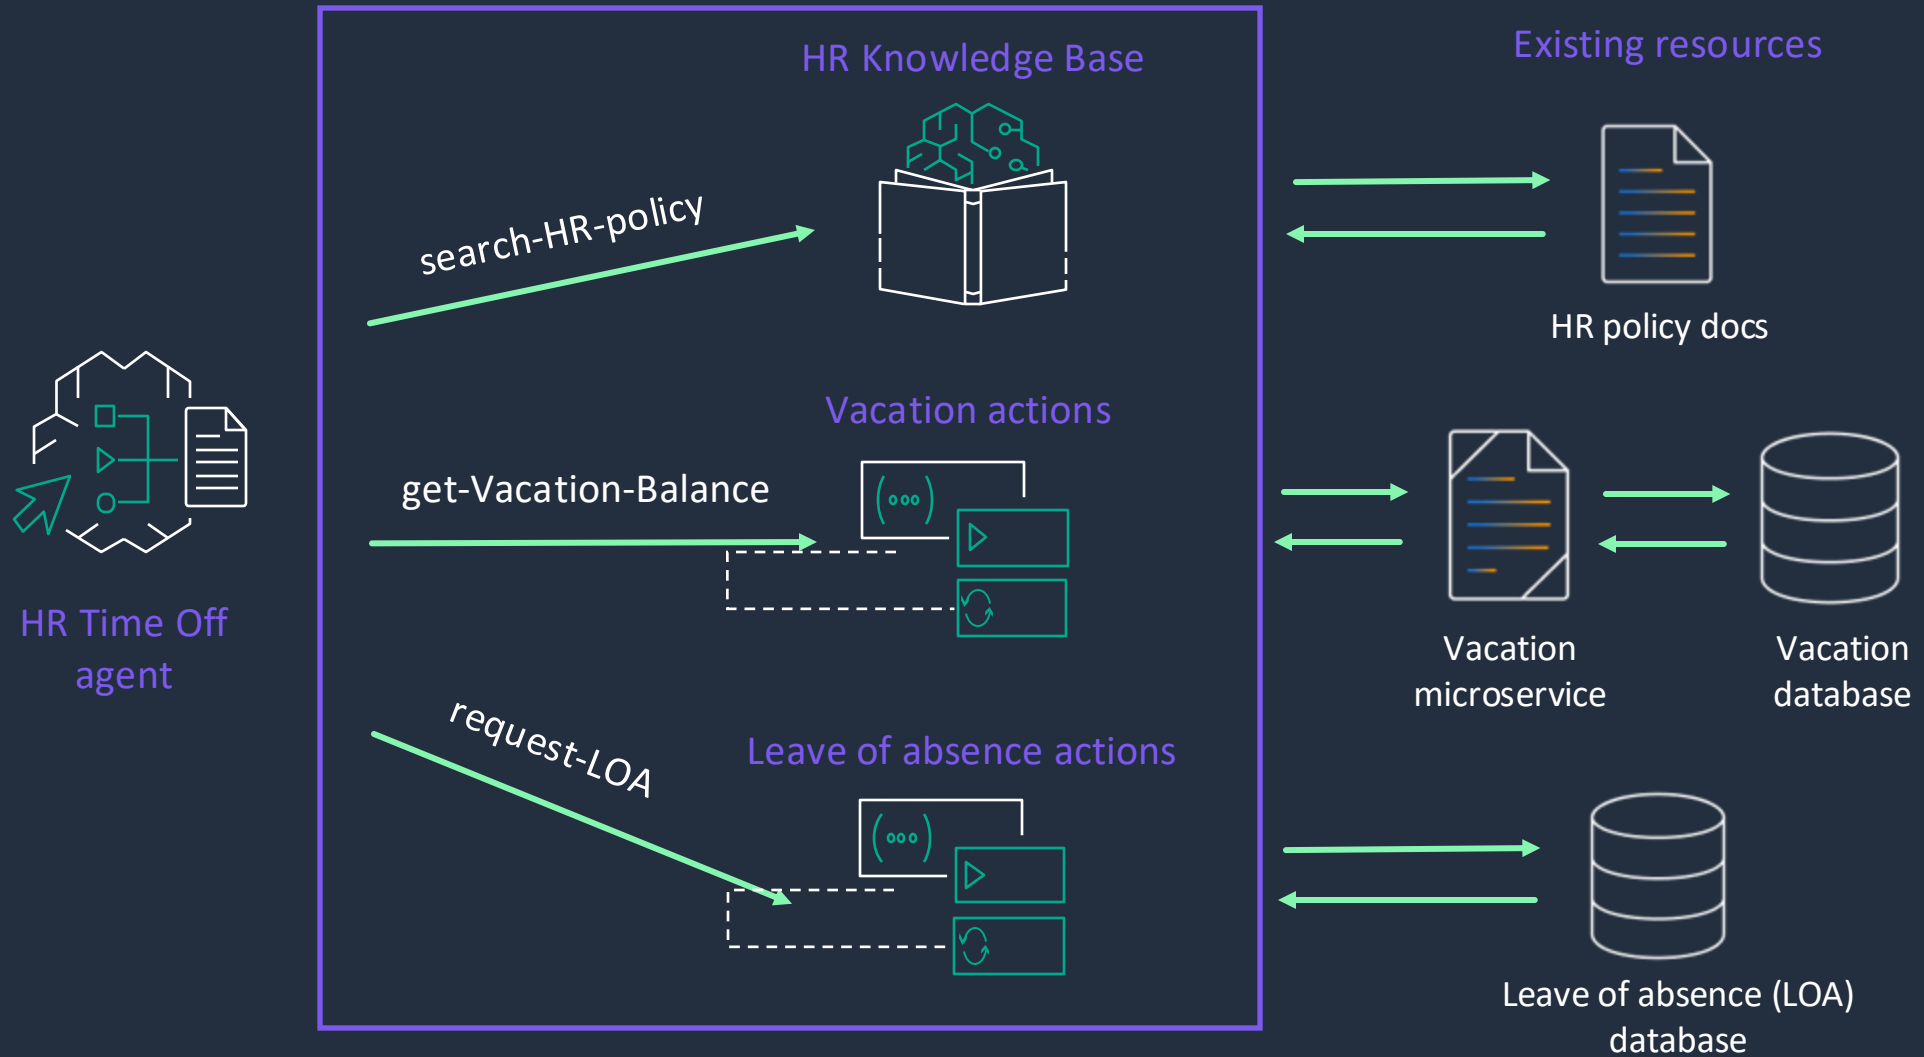
NEW



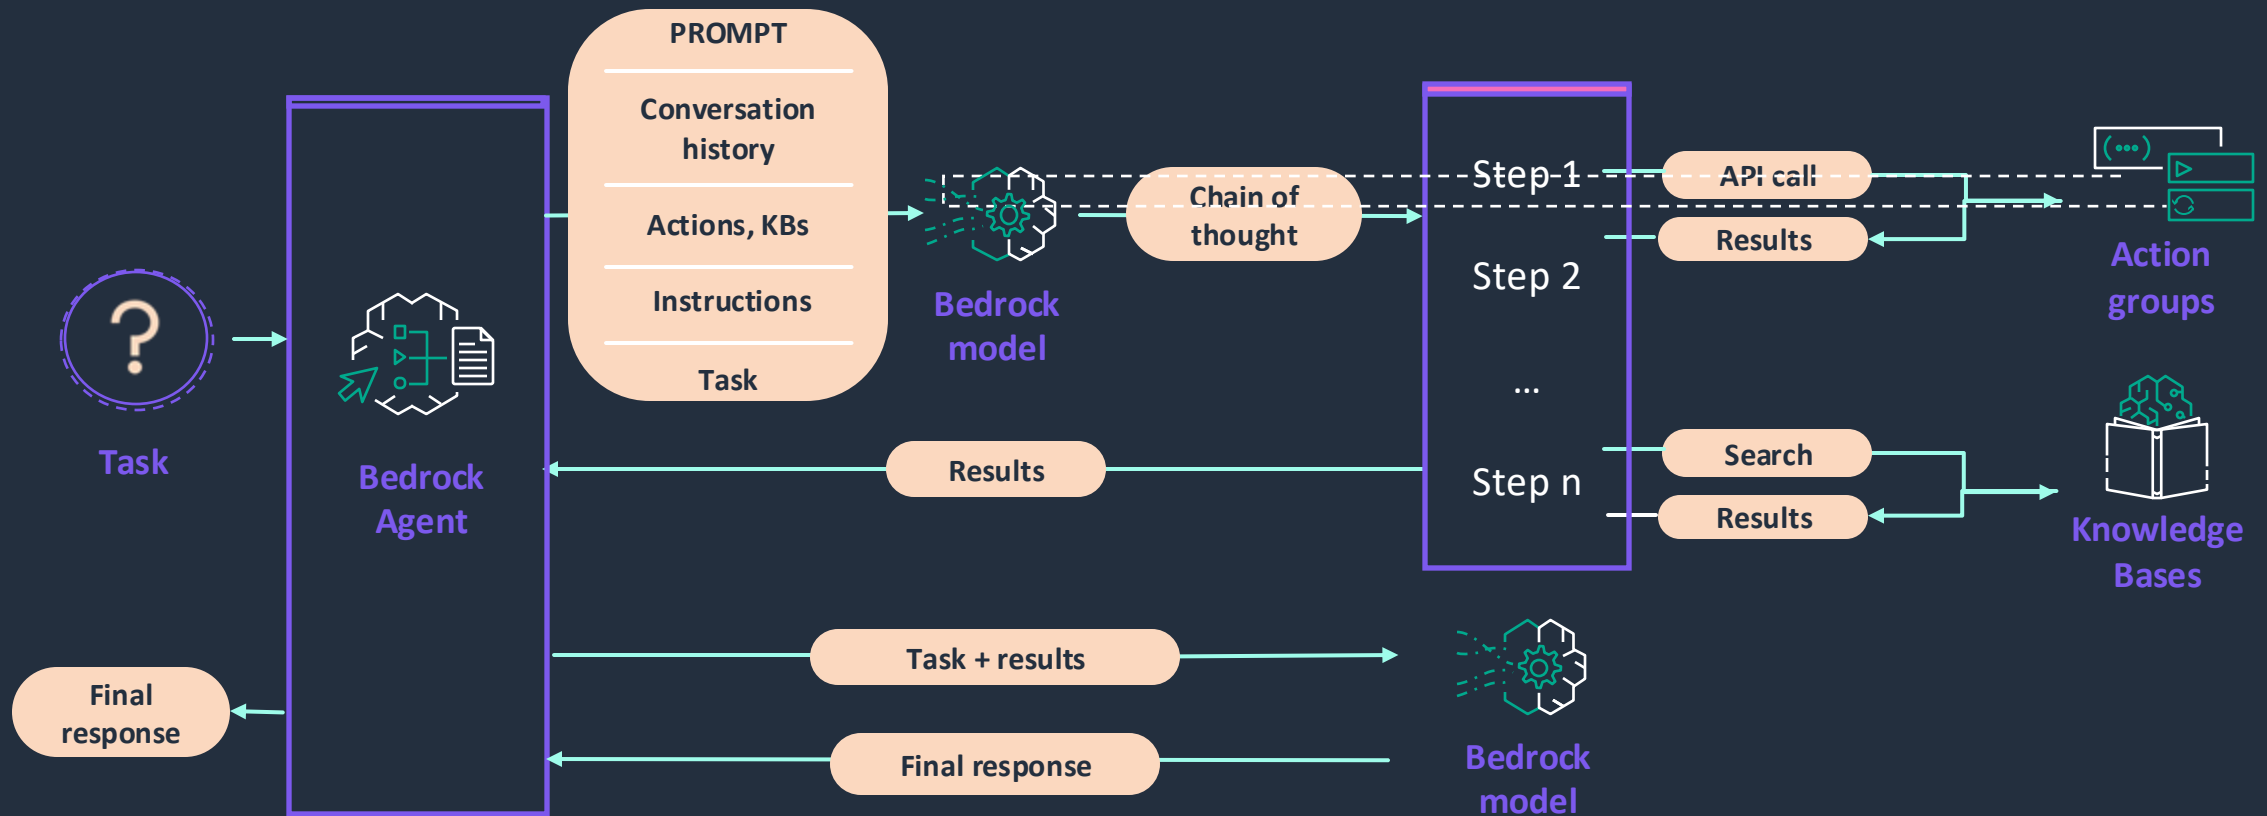
Agent basics



Example - HR Time Off Agent



Agent orchestration - ReAct



Agent breaks task into subtasks, determines the right sequence, and executes actions and knowledge searches on the fly

Custom orchestration

GRANULAR CONTROL OVER TASK PLANNING, COMPLETION, AND VERIFICATION



Full control over
orchestration strategy



Real-time adjustments



Reusable across
use cases

Other orchestration strategies : Plan and Solve, Reason without Observation, Tree of Thought, and Standard Operating Procedures (SOP)

Agent orchestration is transparent – Trace

The screenshot displays the AWS Agent console interface. On the left, a chat window shows a user prompt "send reminder for claim-857" and a successful response: "Reminder sent successfully for claim-857. The send reminder status is InProgress and the tracking ID is 50e8400-e29b-41d4-a716-446655440000." Below the response is a "< Hide trace" link. On the right, the "Trace (4)" panel is open, showing the "Orchestration and knowledge base" step. It includes a "Copy all" button and a description: "After running a prompt in the test window, this orchestration trace allows you to explore the trace steps to understand the linear chain of thought used by the agent's orchestration prompt component. If a knowledge base was invoked, the trace also allows you to see how the results from the knowledge base were summarized to generate an observation that is used for for orchestrating the next step or generating the response." The trace for "Step 1" is shown in a code block, detailing the model invocation input and the rationale for the next steps.

send reminder for claim-857

Reminder sent successfully for claim-857. The send reminder status is InProgress and the tracking ID is 50e8400-e29b-41d4-a716-446655440000.
< Hide trace

Trace (4)
Request ID: dcd6938c-73b2-414d-b589-0ea18d6af40e

Pre-processing | **Orchestration and knowledge base** | Post-processing

Orchestration and knowledge base

After running a prompt in the test window, this orchestration trace allows you to explore the trace steps to understand the linear chain of thought used by the agent's orchestration prompt component. If a knowledge base was invoked, the trace also allows you to see how the results from the knowledge base were summarized to generate an observation that is used for for orchestrating the next step or generating the response.

▼ Step 1 [Show config](#)

Trace

```
1 {  
2   "modelInvocationInput": {  
18   "rationale": {  
19     "text": "To answer this question, I will:\n\n1. Call  
GET::claims-actions::getOutstandingPaperwork  
function to get the list of pending documents for  
claim-857. \n\n2. Check if I have the pending  
documents to send a reminder. \n\n3. If I have the  
pending documents, I will call POST::claims  
-actions::sendReminders function to send the  
reminder for claim-857.\n\n4. Return a success  
message to the user.\n\nI have double checked and  
made sure that I have been provided the GET  
::claims-actions::getOutstandingPaperwork and POST
```

Detailed orchestration trace in the console and from the SDK



Each action group has 3 key elements



Action group description

Overview of actions provided – helps agent know when this action group is relevant

API schema



- Rich definition of each action
- Operation name, input parameters, data types, response details
- Helps agents know **when to use it, how to call it, and how to use results**
- Language agnostic API definition using industry-standard schema



Lambda function

- Implementation of each action
- Contains either business logic or wraps microservices, databases, or tools
- Serverless, scalable, secure
- Choice of programming language (Python, C#, JavaScript, Java, ...)

Action group example

UtilityActionGroup

Description: "This action group provides a set of commonly used actions. Use these actions for things like sending emails and getting team member lists."

API schema

```
{ "openapi": "3.0.0",  
  "info": { "title": "Utility Actions",  
            "description": "..."},  
  "paths": {  
    "/sendEmail": {  
      "post": {  
        "description": "This operation ..."  
        "operationId": "sendEmail",  
        "requestBody": { ... },  
        "responses": {  
          "200": {  
            "text/plain": {...}  
          }  
        }  
      }  
    }  
  }  
}
```

Lambda function

```
def lambda_handler(event, context):  
    if event['apiPath'] == '/sendEmail':  
        result = sendEmail(event)  
    elif event['apiPath'] == '/getTeam':  
        result = getTeam(event)  
  
    response_body = {'application/json':  
                     {'body': result}}  
    action_response = { ...  
                       'responseBody': response_body  
    }  
  
    return {'messageVersion': '1.0',  
           'response': action_response}
```



Code interpretation for Amazon Bedrock Agents

Allow agents to generate and execute code to perform complex analysis

Enable agents to generate and execute code to answer questions and solve problems

Automatically generate charts and analysis using generated code

Analyze files automatically with generated code, including CSV, XLS, JSON, DOC, HTML, TXT, and PDF

Run code in an isolated environment, with built-in guardrails enabled



Memory retention for Amazon Bedrock Agents

Enable agents to keep and use summaries of prior interactions over time

Build agents that learn from previous interactions for more seamless conversations over time

Enable more personalized experiences

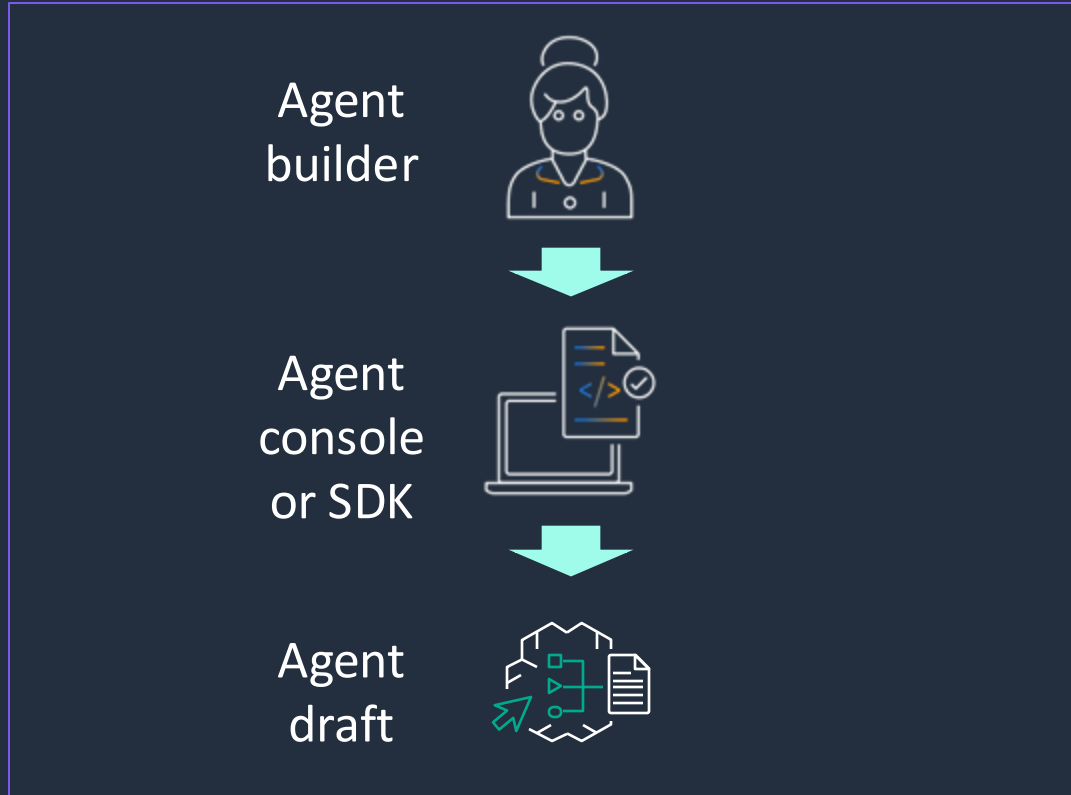
Use unique memory IDs for separation between users

Retain the summaries for up to 30 days

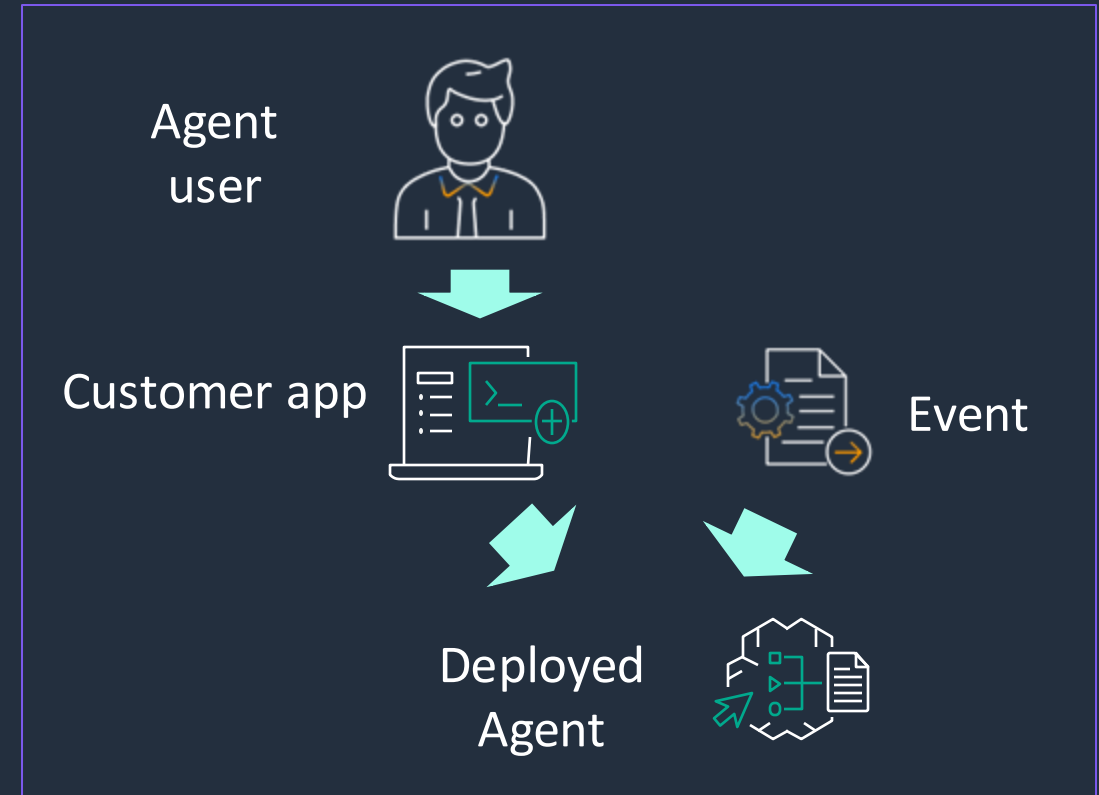


Agents can be deployed and invoked from any app

Building and testing agents



Using production agents



To **deploy** an agent, you create a new **alias**, and optionally a new **version**

Invoking an Agent via the SDK

Invoke agent

```
response = client.invoke_agent(  
    inputText='<user request>',  
    agentId=agent_id,  
    agentAliasId=agent_alias_id,  
    sessionId=str(uuid.uuid1()),  
    enableTrace=True )
```

Process the response stream

```
for event in response['completion']:  
    if 'chunk' in event:  
        data = event['chunk']['bytes']  
        answer = data.decode('utf8')  
        print(f"Answer:\n{answer}")  
    elif 'trace' in event:  
        print(json.dumps(event['trace'],  
                        indent=2))
```


Inline agents

DYNAMICALLY CONFIGURE AGENTS AT RUNTIME




Dynamic agent configuration

Modify agent's instructions, action groups, knowledge bases and other parameters on the fly



Flexible integration

Easily incorporate external APIs and other tooling as needed for each interaction



Contextual adaptation

Adjust the agent's responses based on the user role, preferences or specific scenarios

Amazon Bedrock Guardrails

Apply safeguards customized to your gen AI application requirements and responsible AI policies

Configure thresholds to filter undesirable and potentially harmful text

Define and disallow denied topics with short natural language descriptions

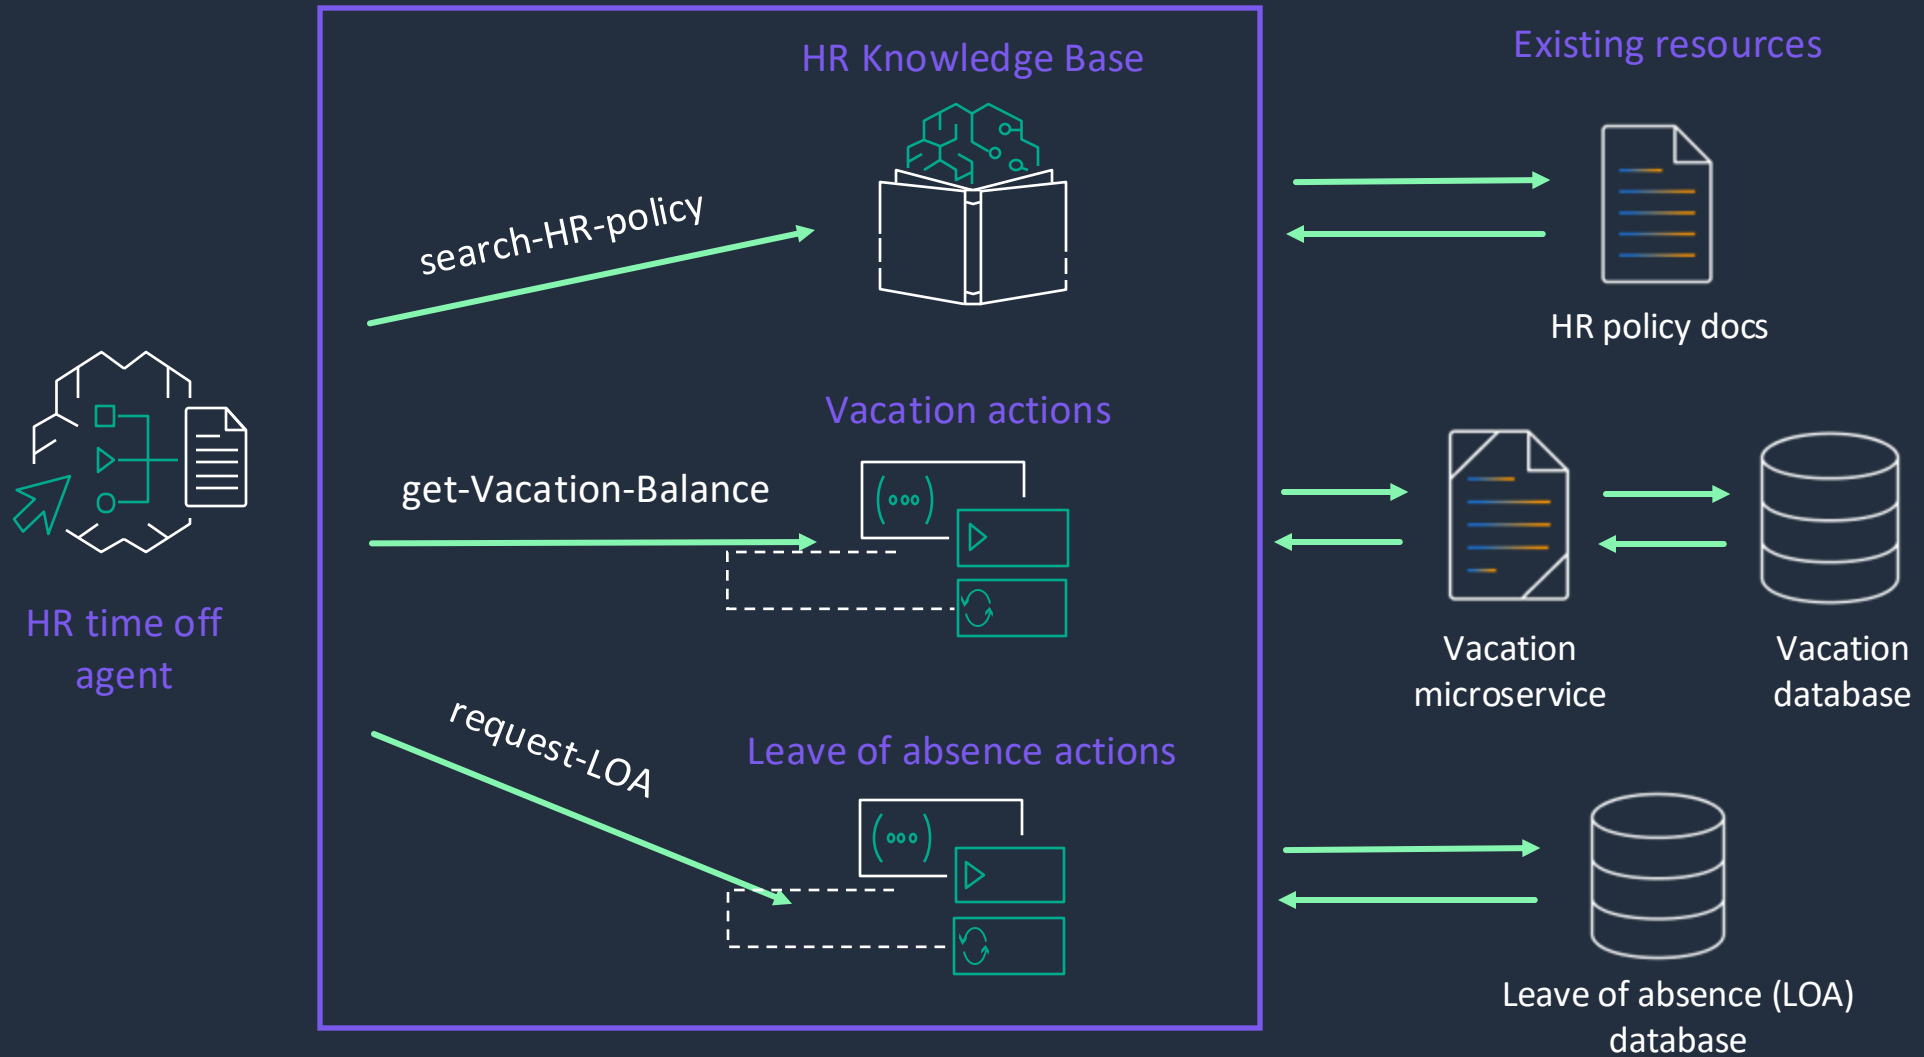
Remove personally identifiable information (PII) and sensitive information in gen AI apps

Define a set of words to detect and block in user inputs and model responses

Filter hallucinations by detecting grounded-ness and relevance of model responses based on context

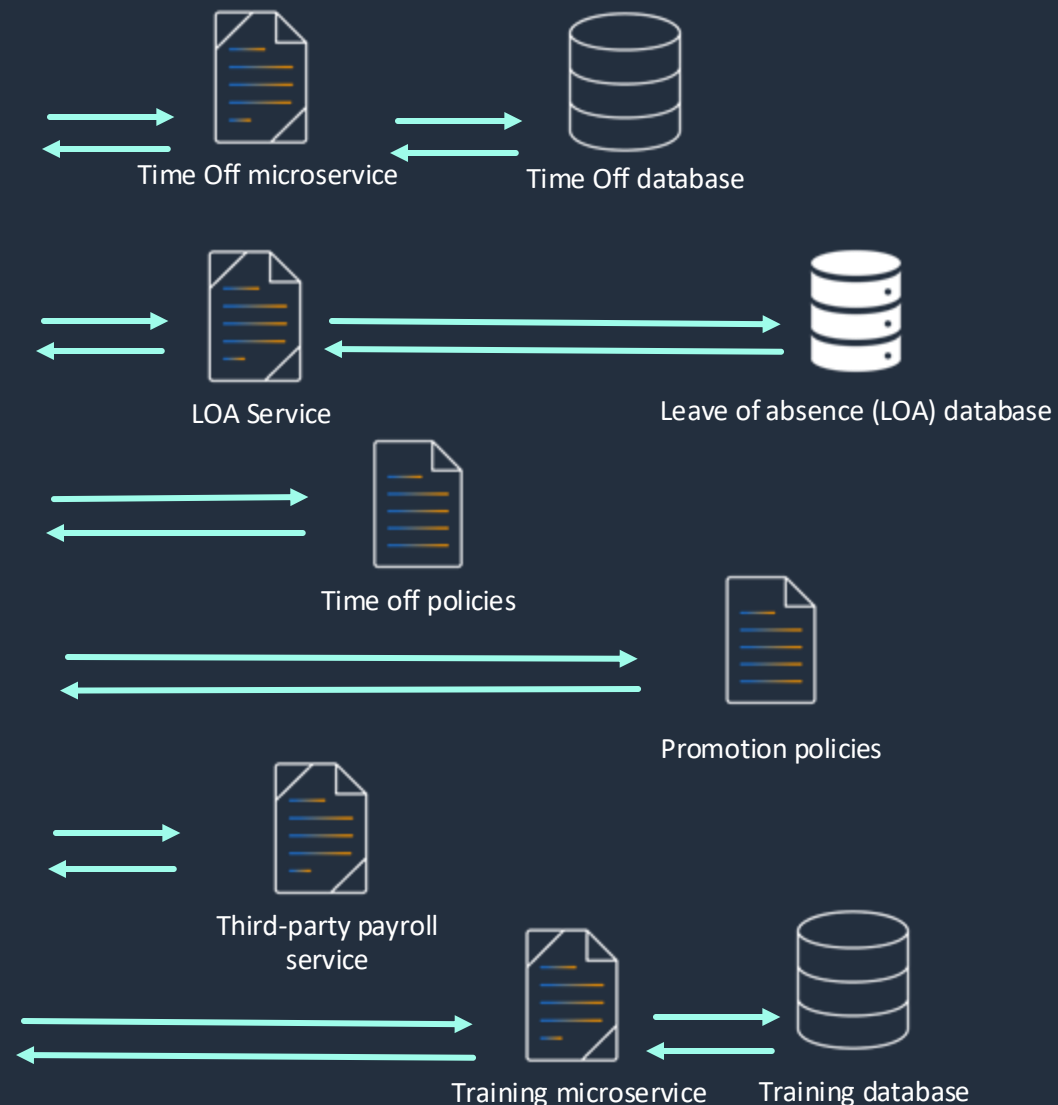
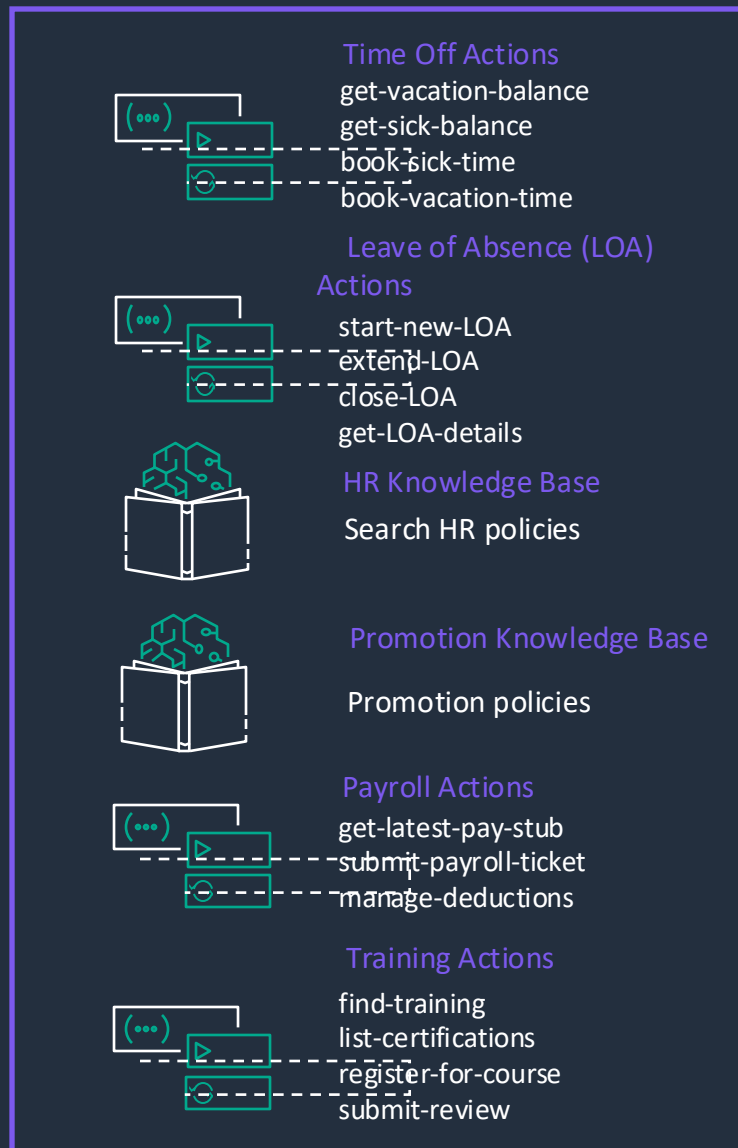


Example - HR time off agent



... but if you take a SINGLE agent too far

Version N



... it leads to challenges

Coding gets complicated



- Complex prompts to limit hallucinations
- Fragile, hard to maintain

Agent gets confused



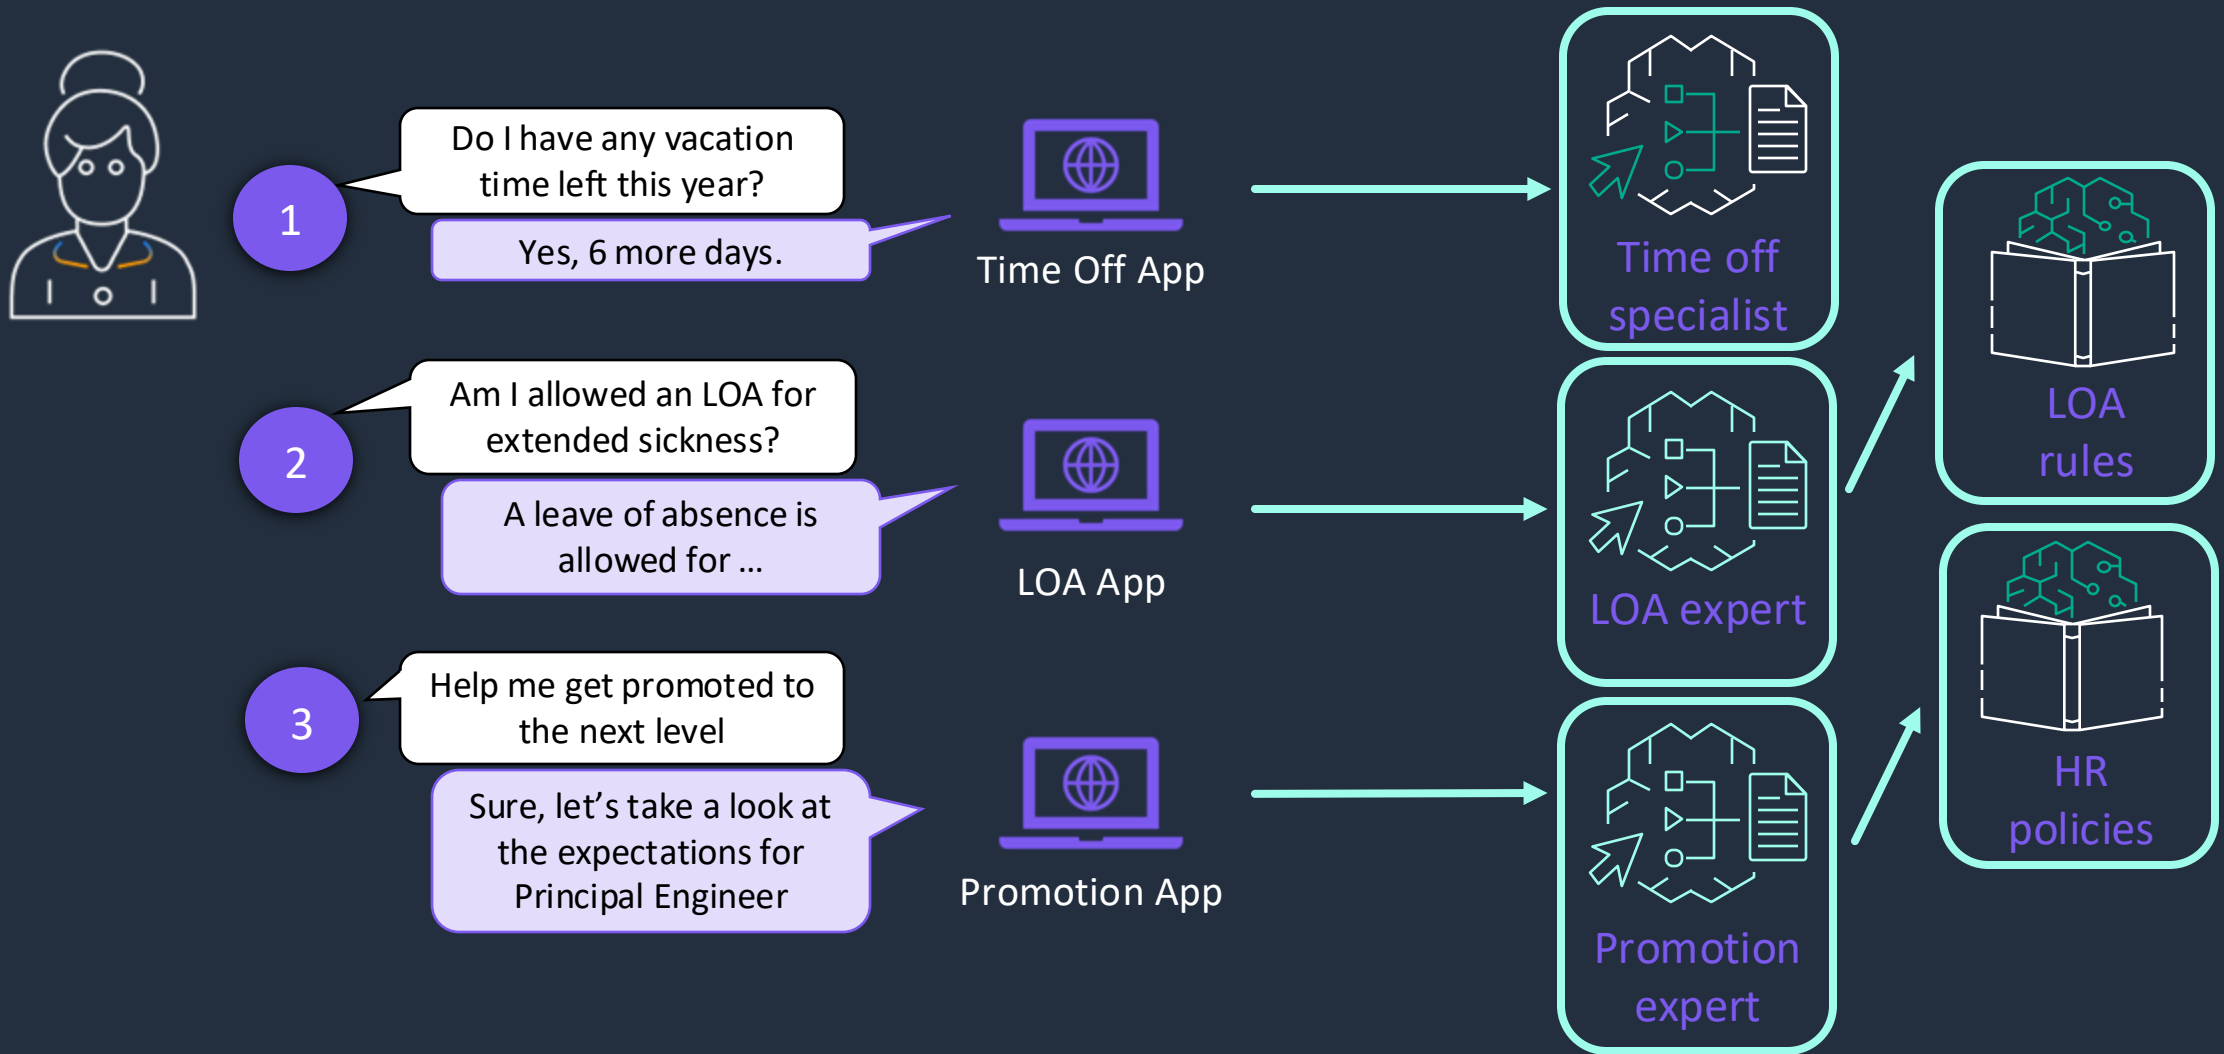
- Calling wrong tools
- Passing wrong arguments
- Inconsistent responses

Agent gets slower and more expensive



- Frontier models needed
- Prompt sizes grow
- Agents retry steps

Using multiple agents helps . . .



Amazon Bedrock Agents multi-agent collaboration

Scaling agentic experiences

Preview



Easily assemble agents and knowledge bases



Plan and execute complex tasks across agents



Unify conversations across agents with built-in intent classification



Observability across multi-agent flows



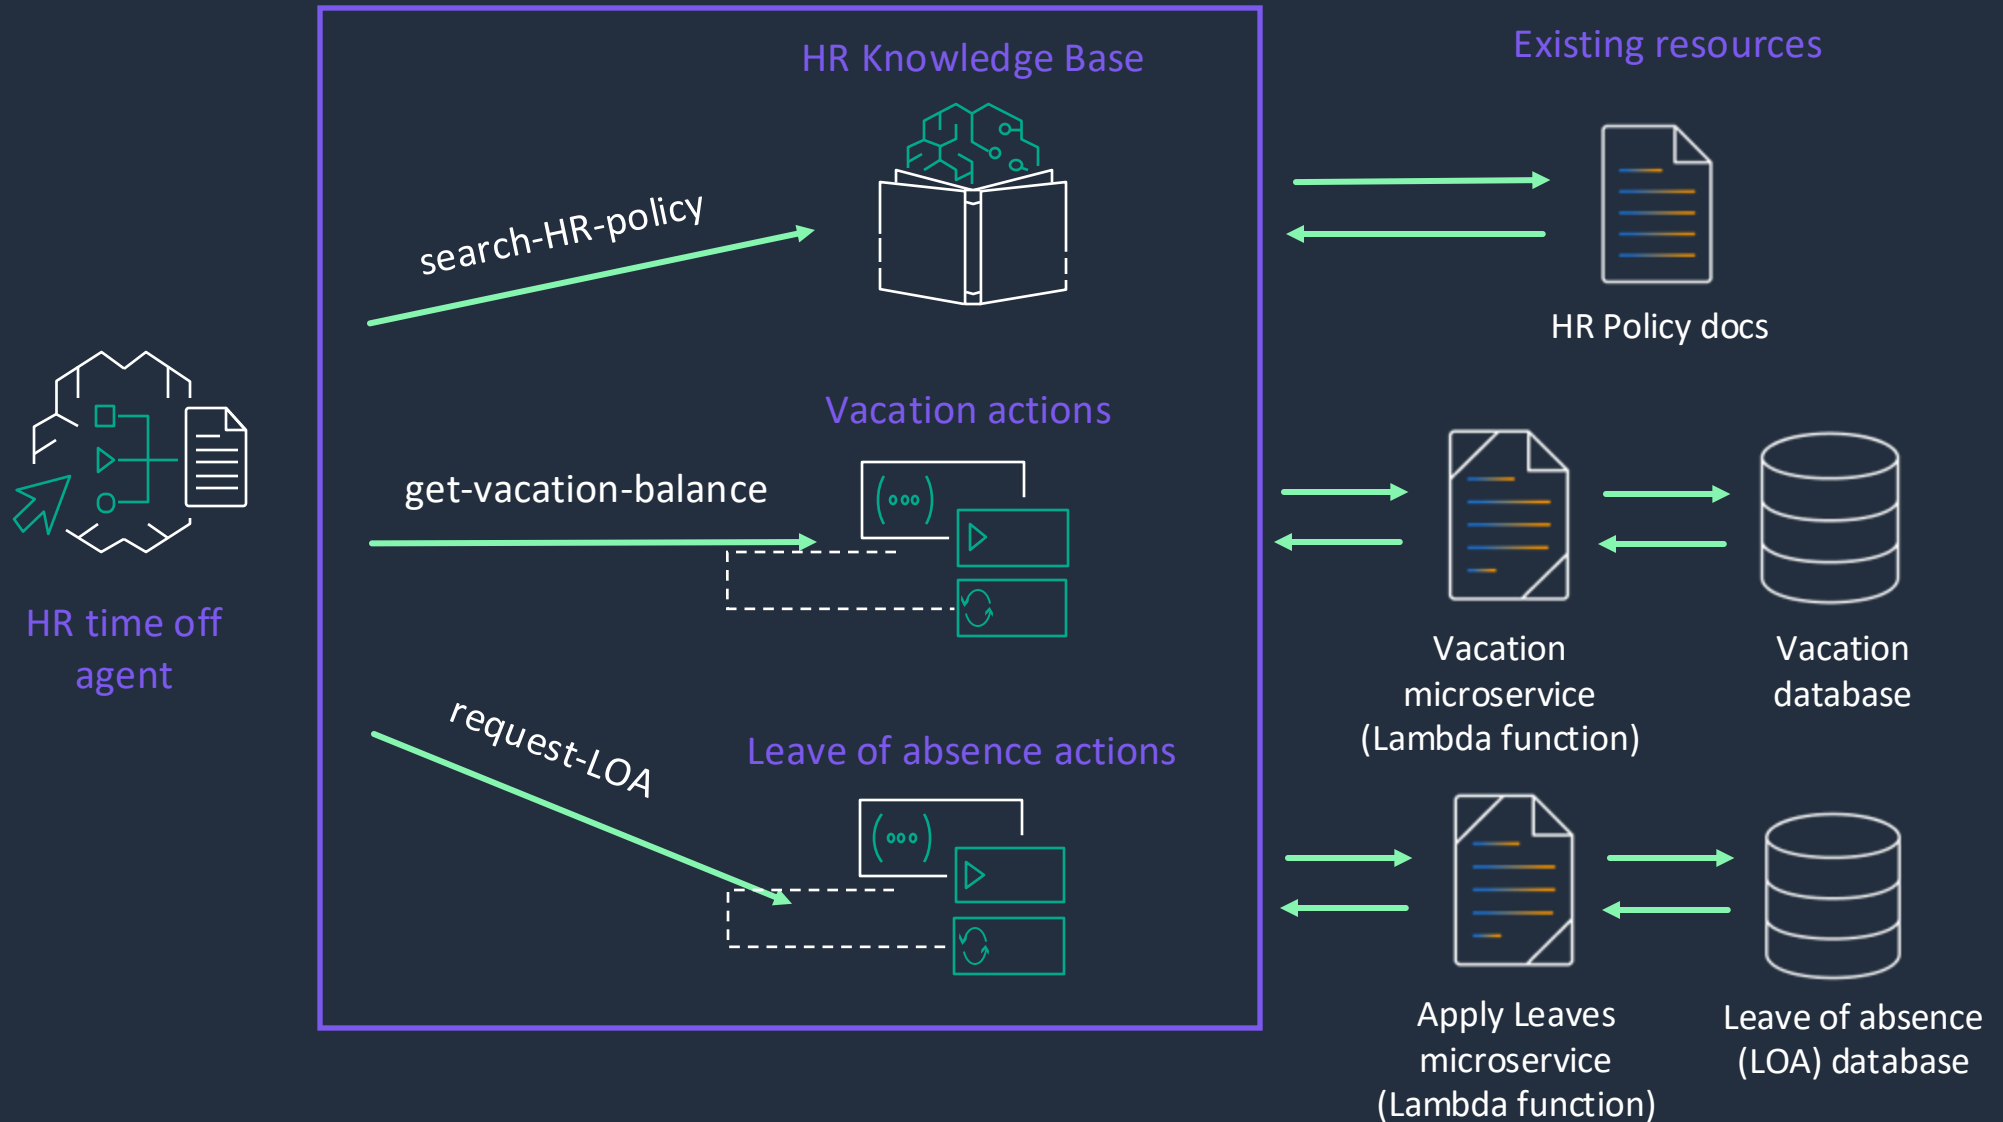
Guardrails, security, and privacy



Demo



Demo - HR time off agent



Recap and summary

AGENTS ON AMAZON BEDROCK

- Amazon Bedrock and agents
- Agent fundamentals
- Building Amazon Bedrock Agents
- Core capabilities: like custom orchestration, code interpretation, memory retention, guardrails, inline agents, etc.
- Multi-agent collaboration on Amazon Bedrock Agents
- Benefits on building agents with Amazon Bedrock



Resources



[Multi-agent collaboration
workshop](#)



[Amazon Bedrock Agents
training](#)



[Bedrock skill training](#)





Thank you!

Sahil Verma

Senior GTM Specialist Solutions Architect,
Generative AI
AWS India

